

```
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np

import keras
from keras.models import Sequential
from keras.layers import Dense
```

▼ Regresja softmax

```
s = [0.2, 0.1, 0.6, 0.1]
exps = [np.exp(i) for i in s]
sum_of_exps = sum(exps)
softmax = [j/sum_of_exps for j in exps]
```

▼ 3 gangi

Zbiór danych:

```
x_label0 = np.random.normal(1, 1.5, (1000, 1))
y_label0 = np.random.normal(1, 1.5, (1000, 1))
x_label1 = np.random.normal(5, 1.5, (1000, 1))
y_label1 = np.random.normal(4, 1.5, (1000, 1))
x_label2 = np.random.normal(8, 1.5, (1000, 1))
y_label2 = np.random.normal(0, 1.5, (1000, 1))

data_label0 = np.concatenate([x_label0, y_label0],axis=1)
data_label1 = np.concatenate([x_label1, y_label1],axis=1)
data_label2 = np.concatenate([x_label2, y_label2],axis=1)
points = np.concatenate([data_label0, data_label1, data_label2],axis=0)
```


Kodowanie one-hot

```
labels = np.array([[1., 0., 0.]] * len(data_label0) + [[0., 1., 0.]] * len(data_label1) + [[0.,0., 1.]] * len(data_label2))
```

```
points.shape, labels.shape
```

```
((3000, 2), (3000, 3))
```

```
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)  
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)  
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)  
plt.show()
```



```
model = Sequential()
```

```
model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))
```

```
#opt = tf.keras.optimizers.Adam(learning_rate=0.1)
```

```
opt = tf.keras.optimizers.Adam()
```

```
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)
```

```
model.compile(loss='binary_crossentropy', optimizer=opt)
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 3)	9
Total params: 9 (36.00 Byte)		
Trainable params: 9 (36.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

```
epochs = 1000
```

```
#h = model.fit(data_points,labels, verbose=1, epochs=epochs,validation_split=0.2)
```

```
h = model.fit(points,labels, verbose=1, epochs=epochs, batch_size= 70, validation_split=0.2)
```

```
Epoch 909/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2449
Epoch 990/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1567 - val_loss: 0.2451
Epoch 991/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2431
Epoch 992/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2419
Epoch 993/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2430
Epoch 994/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2433
Epoch 995/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2472
Epoch 996/1000
35/35 [=====] - 0s 4ms/step - loss: 0.1567 - val_loss: 0.2428
Epoch 997/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2434
Epoch 998/1000
35/35 [=====] - 0s 4ms/step - loss: 0.1567 - val_loss: 0.2433
Epoch 999/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1567 - val_loss: 0.2450
Epoch 1000/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1568 - val_loss: 0.2440
```

```
Loss = h.history['loss']
Loss
```

```
0.15074307942390442,
0.15672850608825684,
0.15678222477436066,
0.15671847760677338,
0.15672115981578827,
0.15673989057540894,
0.15675140917301178,
0.15671634674072266,
0.15681302547454834,
0.15670068562030792,
0.1567232459783554,
0.15674148499965668,
0.15671007335186005,
0.15672567486763,
0.15671215951442719,
0.15676255524158478,
0.15670548379421234,
0.15671847760677338,
0.15670041739940643,
0.1567092388868332,
0.15673935413360596,
0.15671317279338837,
0.15672899782657623,
0.15669238567352295,
0.15672625601291656,
0.15672796964645386,
0.15671248733997345,
0.1567256599664688,
0.15672150254249573,
0.15669448673725128,
0.15669465065002441,
0.1567067801952362,
0.1567171961069107,
0.15676064789295197]
```

```
weights = model.get_weights()
```

```
print(weights[0])
```


```
print(weights[1])    #bias
```

```
[[ -1.8542533   0.37775218  1.5208719 ]
 [ -0.9898408   1.4547888  -1.4014187 ]]
[ 7.970148  -5.1614227 -8.469914 ]
```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```

```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```


```
plt.show()
```



```
model.predict([[4,6]])
```

```
1/1 [=====] - 0s 65ms/step
array([[2.8538452e-05, 9.9997139e-01, 1.2773766e-07]], dtype=float32)
```

```
x=4.0
y=6.0
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)
plt.scatter(x,y,c='y', marker='s')
plt.show()
```



```
Model: "sequential_1"
```

```
epochs = 100
#h = model.fit(data_points,labels, verbose=1, epochs=epochs,validation_split=0.2)
h = model.fit(points,labels, verbose=1, epochs=epochs, batch_size= 70, validation_split=0.2)
```



```
35/35 [=====] - 0s 2ms/step - loss: 0.2490 - val_loss: 0.2688
Epoch 98/100
35/35 [=====] - 0s 2ms/step - loss: 0.2483 - val_loss: 0.2692
Epoch 99/100
35/35 [=====] - 0s 2ms/step - loss: 0.2470 - val_loss: 0.2691
Epoch 100/100
35/35 [=====] - 0s 2ms/step - loss: 0.2457 - val loss: 0.2656
```

```
Loss = h.history['loss']
Loss
```

```

0.2090559554019105,
0.26736921072006226,
0.265764057636261,
0.2641422152519226,
0.26258763670921326,
0.2610386908054352,
0.25954481959342957,
0.2580412030220032,
0.2566033601760864,
0.25513237714767456,
0.25371623039245605,
0.2523210048675537,
0.25096815824508667,
0.24962598085403442,
0.24828501045703888,
0.24699103832244873,
0.2457052618265152]

```

```
weights = model.get_weights()
```

```
print(weights[0])
```

```
print(weights[1])    #bias
```

```


[[-0.6192839   0.12776168  0.46283087]
 [-0.25504512  0.8719286   -0.91608995]]
[ 1.8876922 -2.4891806 -2.2122905]

```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```


```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```

```
plt.show()
```




```
model.predict([[4,6]])
```

1/1 [=====] - 0s 59ms/step
array([[4.6175518e-03, 9.9527246e-01, 1.0992816e-04]], dtype=float32)



```
x=4.0  
y=6.0  
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)  
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)  
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)  
plt.scatter(x,y,c='y', marker='s')  
plt.show()
```



▼ Number of epochs 2000

```
model = Sequential()

model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))

#opt = tf.keras.optimizers.Adam(learning_rate=0.1)
opt = tf.keras.optimizers.Adam()
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)

model.compile(loss='binary_crossentropy', optimizer=opt)

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 3)	9

=====
Total params: 9 (36.00 Byte)
Trainable params: 9 (36.00 Byte)
Non-trainable params: 0 (0.00 Byte)
=====

```
epochs = 2000
#h = model.fit(data_points, labels, verbose=1, epochs=epochs, validation_split=0.2)
h = model.fit(points, labels, verbose=1, epochs=epochs, batch_size= 70, validation_split=0.2)
```

```
Epoch 1980/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2443
Epoch 1981/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2444
Epoch 1982/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2419
Epoch 1983/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2434
Epoch 1984/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2454
Epoch 1985/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2434
Epoch 1986/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2448
Epoch 1987/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2445
Epoch 1988/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2450
Epoch 1989/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2447
Epoch 1990/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2470
Epoch 1991/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2474
Epoch 1992/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2466
Epoch 1993/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2487
Epoch 1994/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2475
Epoch 1995/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2439
Epoch 1996/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2469
Epoch 1997/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2461
Epoch 1998/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2453
Epoch 1999/2000
35/35 [=====] - 0s 3ms/step - loss: 0.1566 - val_loss: 0.2455
Epoch 2000/2000
35/35 [=====] - 0s 2ms/step - loss: 0.1566 - val_loss: 0.2444
```

```
Loss = h.history['loss']
Loss
```


...]

```
weights = model.get_weights()
```

```
print(weights[0])
```


```
print(weights[1])    #bias
```

```
[[ -1.9186313   0.37638855  1.6253421 ]
 [ -1.0266261   1.4578031  -1.4661033 ]
 [  8.280908  -5.164768  -9.07732  ]
```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```

```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```

```
plt.show()
```




```
model.predict([[4,6]])
```

```
1/1 [=====] - 0s 43ms/step
array([[2.3915041e-05, 9.9997592e-01, 7.1027500e-08]], dtype=float32)
```

```

x=4.0
y=6.0
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)
plt.scatter(x,y,c='y', marker='s')
plt.show()

```



▼ Learning rate 0.1

```

model = Sequential()

model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))

opt = tf.keras.optimizers.Adam(learning_rate=0.1)
#opt = tf.keras.optimizers.Adam()
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)

```



```
model.compile(loss='binary_crossentropy',optimizer=opt)
```

```
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 3)	9
Total params: 9 (36.00 Byte)		
Trainable params: 9 (36.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

```
epochs = 1000
```

```
#h = model.fit(data_points,labels, verbose=1, epochs=epochs,validation_split=0.2)
```

```
h = model.fit(points,labels, verbose=1, epochs=epochs, batch_size= 70, validation_split=0.2)
```

```

35/35 [=====] - 0s 2ms/step - loss: 0.1602 - val_loss: 0.2604
Epoch 987/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1636 - val_loss: 0.2565
Epoch 988/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1628 - val_loss: 0.2772
Epoch 989/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1630 - val_loss: 0.2362
Epoch 990/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1616 - val_loss: 0.2718
Epoch 991/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1605 - val_loss: 0.2094
Epoch 992/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1604 - val_loss: 0.2640
Epoch 993/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1603 - val_loss: 0.2967
Epoch 994/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1599 - val_loss: 0.2399
Epoch 995/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1610 - val_loss: 0.2613
Epoch 996/1000
35/35 [=====] - 0s 3ms/step - loss: 0.1598 - val_loss: 0.2741
Epoch 997/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1600 - val_loss: 0.3083
Epoch 998/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1602 - val_loss: 0.3103
Epoch 999/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1608 - val_loss: 0.2647
Epoch 1000/1000
35/35 [=====] - 0s 2ms/step - loss: 0.1598 - val_loss: 0.2681

```

```
Loss = h.history['loss']
```

```
Loss
```

```

0.10379012107453700,
0.16042982041835785,
0.16075921058654785,
0.16320320963859558,
0.16297510266304016,
0.16055388748645782,
0.1608719378709793,
0.16025833785533905,
0.1604868471622467,
0.1593562811613083,
0.16037946939468384,
0.16341416537761688,
0.15937060117721558,
0.16094814240932465,
0.15983568131923676,
0.1607755869626999,
0.1638484001159668,
0.16113193333148956,
0.16345638036727905,
0.17447996139526367,
0.16293780505657196,
0.16053014993667603,
0.16081421077251434,
0.1586494892835617,
0.16016899049282074,
0.16359543800354004,
0.16283239424228668,
0.16302351653575897,
0.16162802278995514,
0.16053351759910583,
0.16038376092910767,
0.16029150784015656,
0.15989957749843597,
0.1610257625579834,
0.1598190814256668,
0.15997277200222015,
0.1602095067501068,
0.16083262860774994,
0.1598118096590042]

```

```
weights = model.get_weights()
```

```
print(weights[0])
```


```
print(weights[1])    #bias
```

```

[[-1.9795942  0.4349273  1.706098 ]
 [-0.8989535  1.4161886 -1.5113696]]
[ 8.54315  -5.302328 -9.565231]

```


```
plt.scatter(np.arange(epochs),h.history['loss'])
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
plt.show()
```



```
model.predict([[4,6]])
```

```
1/1 [=====] - 0s 42ms/step
array([[6.1068109e-05, 9.9993896e-01, 5.3497580e-08]], dtype=float32)
```

```
x=4.0
y=6.0
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)
plt.scatter(x,y,c='y', marker='s')
plt.show()
```



▼ Learning rate 0.0001

```
model = Sequential()

model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))

opt = tf.keras.optimizers.Adam(learning_rate=0.0001)
#opt = tf.keras.optimizers.Adam()
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)

model.compile(loss='binary_crossentropy', optimizer=opt)

model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 3)	9

```
Total params: 9 (36.00 Byte)
Trainable params: 9 (36.00 Byte)
Non-trainable params: 0 (0.00 Byte)
```

```
epochs = 1000
```

```
#h = model.fit(data_points, labels, verbose=1, epochs=epochs, validation_split=0.2)
```

```
h = model.fit(points, labels, verbose=1, epochs=epochs, batch_size= 70, validation_split=0.2)
```

```
Epoch 994/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2106 - val_loss: 0.2484
Epoch 995/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2105 - val_loss: 0.2485
Epoch 996/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2104 - val_loss: 0.2480
Epoch 997/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2103 - val_loss: 0.2483
Epoch 998/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2102 - val_loss: 0.2483
Epoch 999/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2101 - val_loss: 0.2485
Epoch 1000/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2100 - val_loss: 0.2483
Epoch 1000/1000
35/35 [=====] - 0s 2ms/step - loss: 0.2099 - val_loss: 0.2481
```

```
Loss = h.history['loss']
Loss
```

```

0.2124300205742981,
0.21233676373958588,
0.21223615109920502,
0.21213240921497345,
0.21203115582466125,
0.21193036437034607,
0.21183323860168457,
0.21174253523349762,
0.21163320541381836,
0.21153868734836578,
0.2114371806383133,
0.21134229004383087,
0.21124067902565002,
0.21115081012248993,
0.21105551719665527,
0.21095231175422668,
0.21085263788700104,
0.21075329184532166,
0.21066242456436157,
0.2105596661567688,
0.21046391129493713,
0.21037010848522186,
0.2102724313735962,
0.21017520129680634,
0.21008704602718353,
0.20998099446296692,
0.20989182591438293]

```

```
weights = model.get_weights()
```

```
print(weights[0])
```

```
print(weights[1])    #bias
```

```


[[-0.86395454  0.15022355  0.5943968 ]
 [-0.38220945  0.91327024 -0.95378107]]
[ 3.0835235 -2.6992445 -2.9725344]

```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```

```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```


```
plt.show()
```

```
model.predict([[4,6]])
```

```
WARNING:tensorflow:5 out of the last 5 calls to <function Model.make_predict_function.<locals>.predict_function at 0x788676b92440> triggered tf.function
1/1 [=====] - 0s 48ms/step
array([[2.3596396e-03, 9.9757916e-01, 6.1202118e-05]], dtype=float32)
```

```
x=4.0
y=6.0
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)
plt.scatter(x,y,c='y', marker='s')
plt.show()
```



▼ Minibatch



▼ Minibatch-20



```
model = Sequential()
```



```
model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))
```



```
opt = tf.keras.optimizers.Adam(learning_rate=0.0001)
```

```
#opt = tf.keras.optimizers.Adam()
```

```
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)
```

```
model.compile(loss='binary_crossentropy', optimizer=opt)
```

```
model.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 3)	9
Total params: 9 (36.00 Byte)		
Trainable params: 9 (36.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

```
epochs = 1000
```

```
#h = model.fit(data_points, labels, verbose=1, epochs=epochs, validation_split=0.2)
```

```
h = model.fit(points, labels, verbose=1, epochs=epochs, validation_split=0.2, batch_size=20)
```



```
120/120 [=====] - 0s 2ms/step - loss: 0.1643 - val_loss: 0.2391
Epoch 998/1000
120/120 [=====] - 0s 2ms/step - loss: 0.1643 - val_loss: 0.2388
Epoch 999/1000
120/120 [=====] - 0s 2ms/step - loss: 0.1643 - val_loss: 0.2385
Epoch 1000/1000
120/120 [=====] - 0s 2ms/step - loss: 0.1643 - val_loss: 0.2391
```

```
Loss = h.history['loss']
Loss
```

```
0.16455519199371338,  
0.16453391313552856,  
0.1645217388868332,  
0.16450464725494385,  
0.1644873470067978,  
0.16446520388126373,  
0.16444945335388184,  
0.1644359529018402,  
0.1644258201122284,  
0.1644003540277481,  
0.16438671946525574,  
0.16437384486198425,  
0.16434895992279053,  
0.16434694826602936,  
0.16432400047779083,  
0.16430972516536713,  
0.1642874926328659]
```

```
weights = model.get_weights()
```

```
print(weights[0])
```

```
print(weights[1])    #bias
```

```
[[ -1.376739   0.3119349  1.0393476]  
 [ -0.7011174  1.2859708 -1.1403483]]  
[ 5.670027 -4.416901 -5.620191]
```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```

```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```


```
plt.show()
```



```
model.predict([[4,6]])
```

```
WARNING:tensorflow:6 out of the last 6 calls to <function Model.make_predict_function.<locals>.predict_function at 0x788710239ea0> triggered tf.funct
1/1 [=====] - 0s 64ms/step
array([[1.858657e-04, 9.998116e-01, 2.621559e-06]], dtype=float32)
```

```
x=4.0
y=6.0
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)
plt.scatter(x,y,c='y', marker='s')
plt.show()
```



▼ mini batch - 50

```
model = Sequential()

model.add(Dense(units = 3, use_bias=True, input_dim=2, activation = "softmax"))

opt = tf.keras.optimizers.Adam(learning_rate=0.0001)
#opt = tf.keras.optimizers.Adam()
#opt = tf.keras.optimizers.SGD(learning_rate=0.1)

model.compile(loss='binary_crossentropy',optimizer=opt)

model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 3)	9

=====
Total params: 9 (36.00 Byte)
Trainable params: 9 (36.00 Byte)
Non-trainable params: 0 (0.00 Byte)
=====

```
epochs = 1000
#h = model.fit(data_points,labels, verbose=1, epochs=epochs,validation_split=0.2)
h = model.fit(points,labels, verbose=1, epochs=epochs, batch_size= 50, validation_split=0.2)
```

```
Epoch 980/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1898 - val_loss: 0.2435
Epoch 981/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1897 - val_loss: 0.2442
Epoch 982/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1896 - val_loss: 0.2440
Epoch 983/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1896 - val_loss: 0.2440
Epoch 984/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1895 - val_loss: 0.2439
Epoch 985/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1894 - val_loss: 0.2440
Epoch 986/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1894 - val_loss: 0.2440
Epoch 987/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1893 - val_loss: 0.2439
Epoch 988/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1892 - val_loss: 0.2437
Epoch 989/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1892 - val_loss: 0.2441
Epoch 990/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1891 - val_loss: 0.2439
Epoch 991/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1890 - val_loss: 0.2436
Epoch 992/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1890 - val_loss: 0.2435
Epoch 993/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1889 - val_loss: 0.2437
Epoch 994/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1888 - val_loss: 0.2435
Epoch 995/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1888 - val_loss: 0.2440
Epoch 996/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1887 - val_loss: 0.2439
Epoch 997/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1886 - val_loss: 0.2436
Epoch 998/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1886 - val_loss: 0.2437
Epoch 999/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1885 - val_loss: 0.2438
Epoch 1000/1000
48/48 [=====] - 0s 2ms/step - loss: 0.1884 - val_loss: 0.2437
```

```
Loss = h.history['loss']
Loss
```



```
0.188443/0543950/5/1
```

```
weights = model.get_weights()
```

```
print(weights[0])
```


```
print(weights[1])    #bias
```

```
[[-0.98609793  0.21861346  0.7294145 ]
 [-0.46031675  1.0607656  -0.997672  ]]
 [ 3.722764  -3.4068813 -3.7914543]
```

```
plt.scatter(np.arange(epochs),h.history['loss'])
```

```
plt.scatter(np.arange(epochs),h.history['val_loss'],c='r')
```


```
plt.show()
```



```
model.predict([[4,6]])
```

```
1/1 [=====] - 0s 47ms/step
array([[1.0952058e-03, 9.9888200e-01, 2.2700757e-05]], dtype=float32)
```

```
x=4.0  
y=6.0  
plt.scatter(x_label0, y_label0, c='r', marker='x', s=20)  
plt.scatter(x_label1, y_label1, c='b', marker='x', s=20)  
plt.scatter(x_label2, y_label2, c='g', marker='x', s=20)  
plt.scatter(x,y,c='y', marker='s')  
plt.show()
```



Najlepsze wyniki otrzymałem dla współczynnika uczenia 0.1, liczby epok 2000, batcha równego 20, najgorsze dla współczynnika uczenia 0.0001, liczby epok 100, batcha równego 50.