

CNN na przykładzie MNIST

✓ Setup

Importujemy potrzebne biblioteki

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import keras
```

✓ 1 Przygotowanie danych

✓ 1.0 Pobranie zbioru danych

Pobieramy zbiór danych i sprawdzamy rozmiar 28 x 28 pixeli.

```
(x_train_data, y_train_data), (x_test_data, y_test_data) = tf.keras.datasets.mnist.load_data()
```

```
dataset_labels = ["0", # index 0
                  "1", # index 1
                  "2", # index 2
                  "3", # index 3
                  "4", # index 4
                  "5", # index 5
                  "6", # index 6
                  "7", # index 7
                  "8", # index 8
                  "9"] # index 9
```

```
print("x_train shape:", x_train_data.shape, "y_train shape:", y_train_data.shape)
print("x_test shape:", x_test_data.shape, "y_test shape:", y_test_data.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```


```
x_train shape: (60000, 28, 28) y_train shape: (60000,)  
x_test shape: (10000, 28, 28) y_test shape: (10000,)
```

✓ 1.1 Wizualizacja danych

Przykładowy obrazek ze bioru danych

```
def plot_image(img_index):  
    label_index = y_train_data[img_index]  
    plt.imshow(x_train_data[img_index]/255, cmap = 'gray')  
    plt.title("Image "+str(img_index)+":"+dataset_labels[label_index])
```

```
img_index = 60  
plot_image(img_index)
```



✓ 1.2 Normalizacja danych

Na początek sprawdzamy jakie są max i min wartości pixeli w obrazkach.

Wartości te powinny być zawarte w przedziale [0,1].

```
print("Wartości min:", np.min(x_train_data), " max:", np.max(x_train_data))

x_train_data = x_train_data.astype('float32') / 255
x_test_data = x_test_data.astype('float32') / 255

print("Wartości po przeskalowaniu min:", np.min(x_train_data), " max:", np.max(x_train_data))

Wartości min: 0 max: 255
Wartości po przeskalowaniu min: 0.0 max: 1.0
```

✓ 1.3 Podział zbioru danych na zbiór treningowy/walidacyjny/testowy

- **Zbiór treningowy** - wykorzystamy go do uczenia.
- **Zbiór walidacyjny** - wykorzystamy go do tuningu hiperparametrów.
- **Zbiór testowy** - wykorzystamy go do ostatecznego sprawdzenia modelu.

Zbiór walidacyjny stworzymy z 10% zbioru treningowego.

```
validation_fraction = .1

total_train_samples = len(x_train_data)
validation_samples = int(total_train_samples * validation_fraction)
train_samples = total_train_samples - validation_samples

(x_train, x_valid) = x_train_data[:train_samples], x_train_data[train_samples:]
(y_train, y_valid) = y_train_data[:train_samples], y_train_data[train_samples:]

x_test, y_test = x_test_data, y_test_data
print(train_samples, validation_samples, len(x_test))

54000 6000 10000
```

✓ 1.4 Dwa dodatkowe kroki

1. Większość zestawów danych obrazu składa się z obrazów rgb. Z tego powodu Keras oczekuje, że każdy obraz będzie miał 3 wymiary: [x_pixels, y_pixels, color_channels]. Ponieważ nasze obrazki są w skali szarości, wymiar koloru jest równy 1. Musimy zatem zmienić kształt obrazków.

2. W procesie uczenia naszego modelu będziemy wykorzystywali tzw. **kategoryczną entropię krzyżową** (<https://keras.io/losses/>). Musimy przekształcić wektory z etykietami (labelami) do **kodowania one-hot**. Wykorzystamy do tego funkcję `tf.keras.utils.to_categorical()`.

```
# Zmieniamy kształt z (28, 28) na (28, 28, 1)
w, h = 28, 28
x_train = x_train.reshape(x_train.shape[0], w, h, 1)
x_valid = x_valid.reshape(x_valid.shape[0], w, h, 1)
x_test = x_test.reshape(x_test.shape[0], w, h, 1)

# Kodowanie one-hot
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_valid = tf.keras.utils.to_categorical(y_valid, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)

print("x_train shape:", x_train.shape, "y_train shape:", y_train.shape)

# Ilość elementów w zbiorach
print(x_train.shape[0], 'train set')
print(x_valid.shape[0], 'validation set')
print(x_test.shape[0], 'test set')

x_train shape: (54000, 28, 28, 1) y_train shape: (54000, 10)
54000 train set
6000 validation set
10000 test set
```

✓ 2 Stworzenie modelu

Keras oferuje dwa API:

1. [Sequential model API](#)
2. [Functional API](#)

W naszym modelu wykorzystamy Sequential model API. Będziemy wykorzystywali następujące metody:

- `Dense()` [link text](#) - tworzy **warstwę gęstą**
- `Conv2D()` [link text](#) - tworzy **warstwę konwolucyjną**
- `Pooling()` [link text](#) - tworzy **warstwę pooling**
- `Dropout()` [link text](#) - zastosowanie **dropout**

✓ 2.0 Prosty model liniowy

Zacniemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przkształci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 10)	7850
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optimizer adam**
- Jako loss function użyjemy **'categorical_crossentropy'**
- Lista parametrów, tutaj zaczniemy od **'precyzji'**

Warto zerknąć: <https://keras.io/models/model/>

```
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???
```

✓ Uczenie modelu

Model uczymy wykorzystując fit().



```
history = model.fit(x_train, y_train, batch_size = 64, epochs = 5, validation_data = (x_valid,y_valid))
```

```
Epoch 1/5
844/844 [=====] - 9s 8ms/step - loss: 0.5838 - accuracy: 0.8506 - val_loss: 0.2990 - val_accuracy: 0.9217
Epoch 2/5
844/844 [=====] - 5s 6ms/step - loss: 0.3377 - accuracy: 0.9074 - val_loss: 0.2616 - val_accuracy: 0.9287
Epoch 3/5
844/844 [=====] - 6s 7ms/step - loss: 0.3056 - accuracy: 0.9153 - val_loss: 0.2440 - val_accuracy: 0.9333
Epoch 4/5
844/844 [=====] - 3s 3ms/step - loss: 0.2911 - accuracy: 0.9193 - val_loss: 0.2362 - val_accuracy: 0.9342
Epoch 5/5
844/844 [=====] - 2s 3ms/step - loss: 0.2823 - accuracy: 0.9207 - val_loss: 0.2326 - val_accuracy: 0.9367
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`.
saving_api.save_model(
```

Wczytanie modelu

```
#from keras.models import load_model
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.9244999885559082
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'


```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])


        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                   dataset_labels[true_index],
                                   color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 1s 1ms/step

Test wyniki:



✓ 1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.


```

for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))

visualize_filters(filters, 'Filtr')

```



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.


```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:, :, :, label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:, :, :, i] = avg_images[:, :, :, i] / class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.0 Prosty model liniowy v2

Zacznijemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przeksztalci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
flatten_6 (Flatten)	(None, 784)	0
dense_15 (Dense)	(None, 10)	7850
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optimizera adam**
- Jako loss function użyjemy '**categorical_crossentropy**'
- Lista parametrów, tutaj zaczniemy od '**precyzji**'

Warto zerknąć: <https://keras.io/models/model/>

```
opt = keras.optimizers.Adam(learning_rate=0.002)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???
```

✓ Uczenie modelu

Model uczy my wykorzystując fit().

```
history = model.fit(x_train, y_train, batch_size = 32, epochs = 5, validation_data = (x_valid,y_valid))
```



```
Epoch 1/5
1688/1688 [=====] - 6s 3ms/step - loss: 0.4105 - accuracy: 0.8882 - val_loss: 0.2464 - val_accuracy: 0.9313
Epoch 2/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.2977 - accuracy: 0.9161 - val_loss: 0.2368 - val_accuracy: 0.9327
```

```
Epoch 3/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.2817 - accuracy: 0.9210 - val_loss: 0.2271 - val_accuracy: 0.9377
Epoch 4/5
1688/1688 [=====] - 6s 3ms/step - loss: 0.2742 - accuracy: 0.9228 - val_loss: 0.2275 - val_accuracy: 0.9363
Epoch 5/5
1688/1688 [=====] - 4s 3ms/step - loss: 0.2673 - accuracy: 0.9247 - val_loss: 0.2318 - val_accuracy: 0.9335
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`.
saving_api.save_model(
```

Wczytanie modelu

```
#from keras.models import load_model
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.9204000234603882
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'

```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])

        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                   dataset_labels[true_index],
                                   color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 0s 1ms/step

Test wyniki:



✓ 1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.

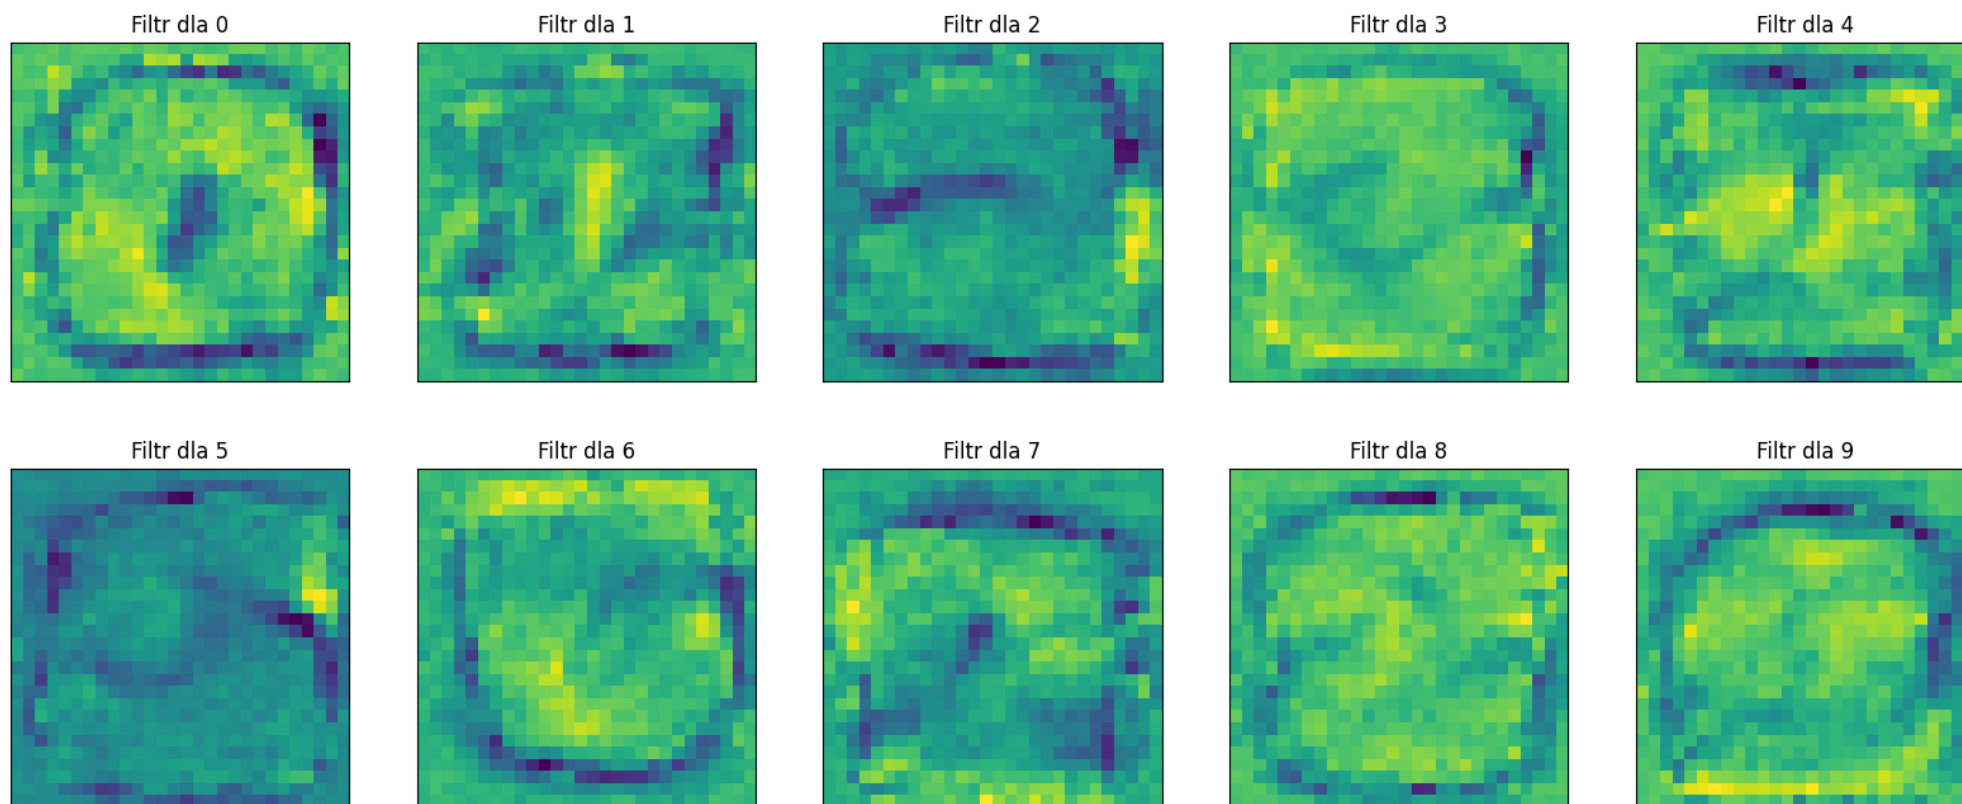
```

for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))

visualize_filters(filters, 'Filtr')

```



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.

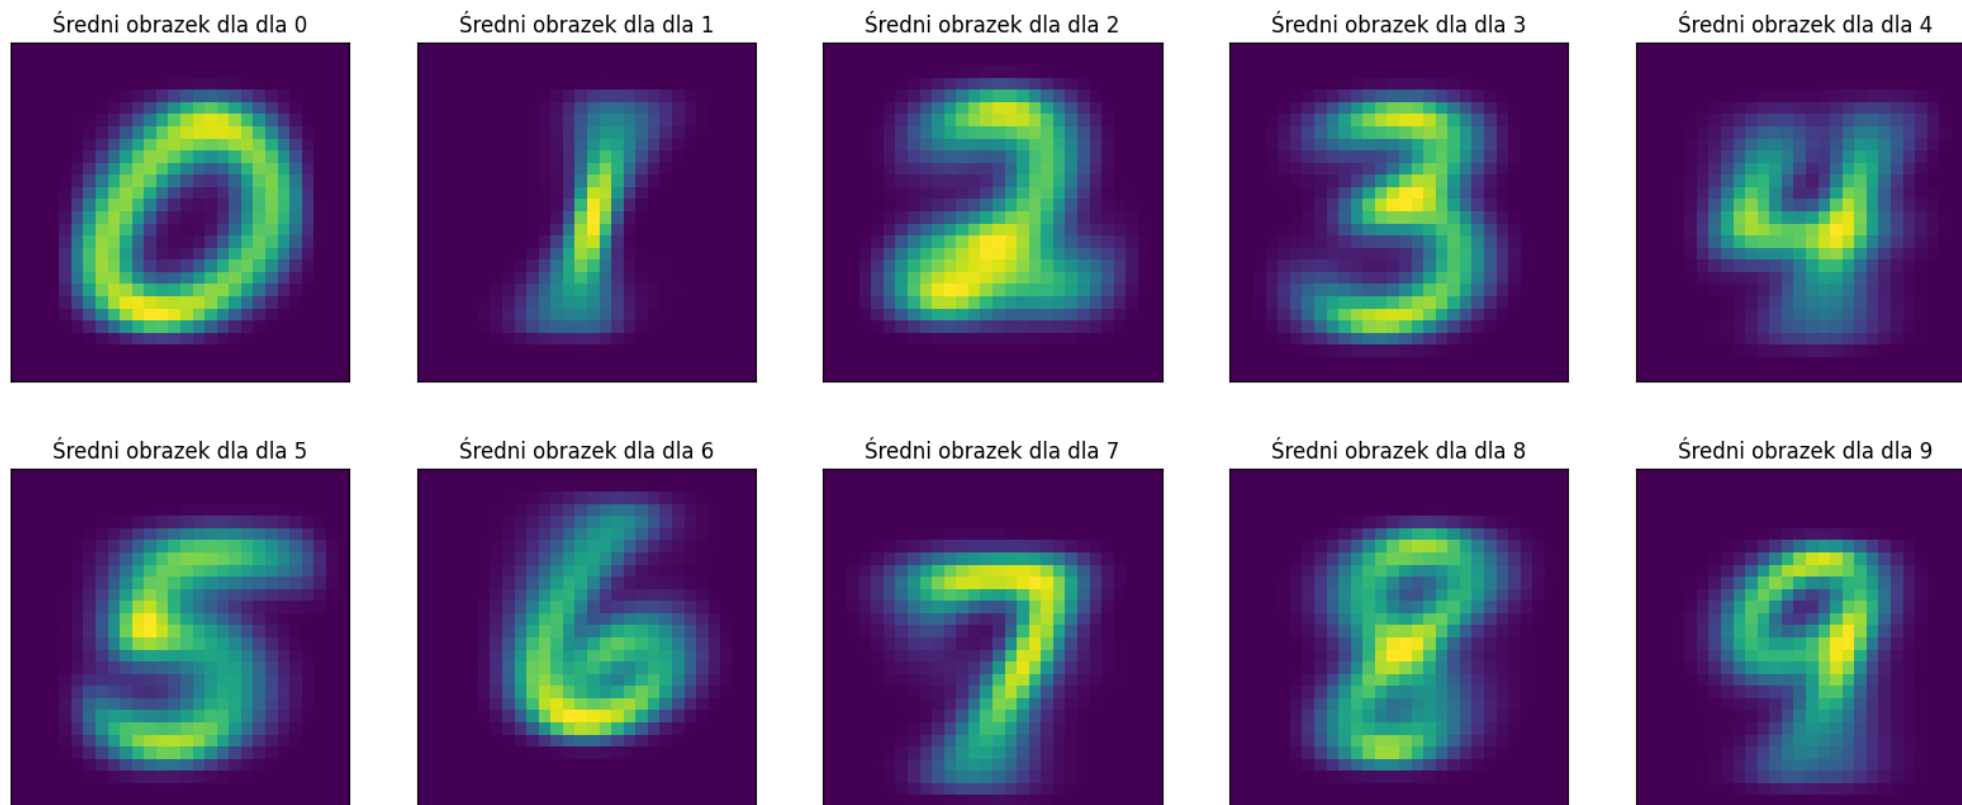
```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:, :, :, label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:, :, :, i] = avg_images[:, :, :, i]/class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.0 Prosty model liniowy v3

Zacniemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przeksztalci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential_17"

Layer (type)	Output Shape	Param #
flatten_17 (Flatten)	(None, 784)	0
dense_26 (Dense)	(None, 10)	7850
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optimizera** `sgd`
- Jako loss function użyjemy `'categorical_crossentropy'`
- Lista parametrów, tutaj zaczniemy od `'precyzji'`

Warto zerknąć: <https://keras.io/models/model/>

```
opt = keras.optimizers.SGD(learning_rate=0.04)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

✓ Uczenie modelu

Model uczymy wykorzystując `fit()`.



```
history = model.fit(x_train, y_train, batch_size = 64, epochs = 75, validation_data = (x_valid,y_valid))
```


844/844 [=====] - 25 imgs/step - loss: 0.2373 - accuracy: 0.9290 - val_loss: 0.2207 - val_accuracy: 0.9373

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

Wczytanie modelu

```
#from keras.models import load_model  
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)  
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.9243000149726868
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'

```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])

        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                   dataset_labels[true_index],
                                   color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 0s 1ms/step

Test wyniki:



1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.

```


for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))


visualize_filters(filters, 'Filtr')

```


Filtr dla 0




Filtr dla 1




Filtr dla 2




Filtr dla 3




Filtr dla 4




Filtr dla 5




Filtr dla 6




Filtr dla 7



Filtr dla 8



Filtr dla 9



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.

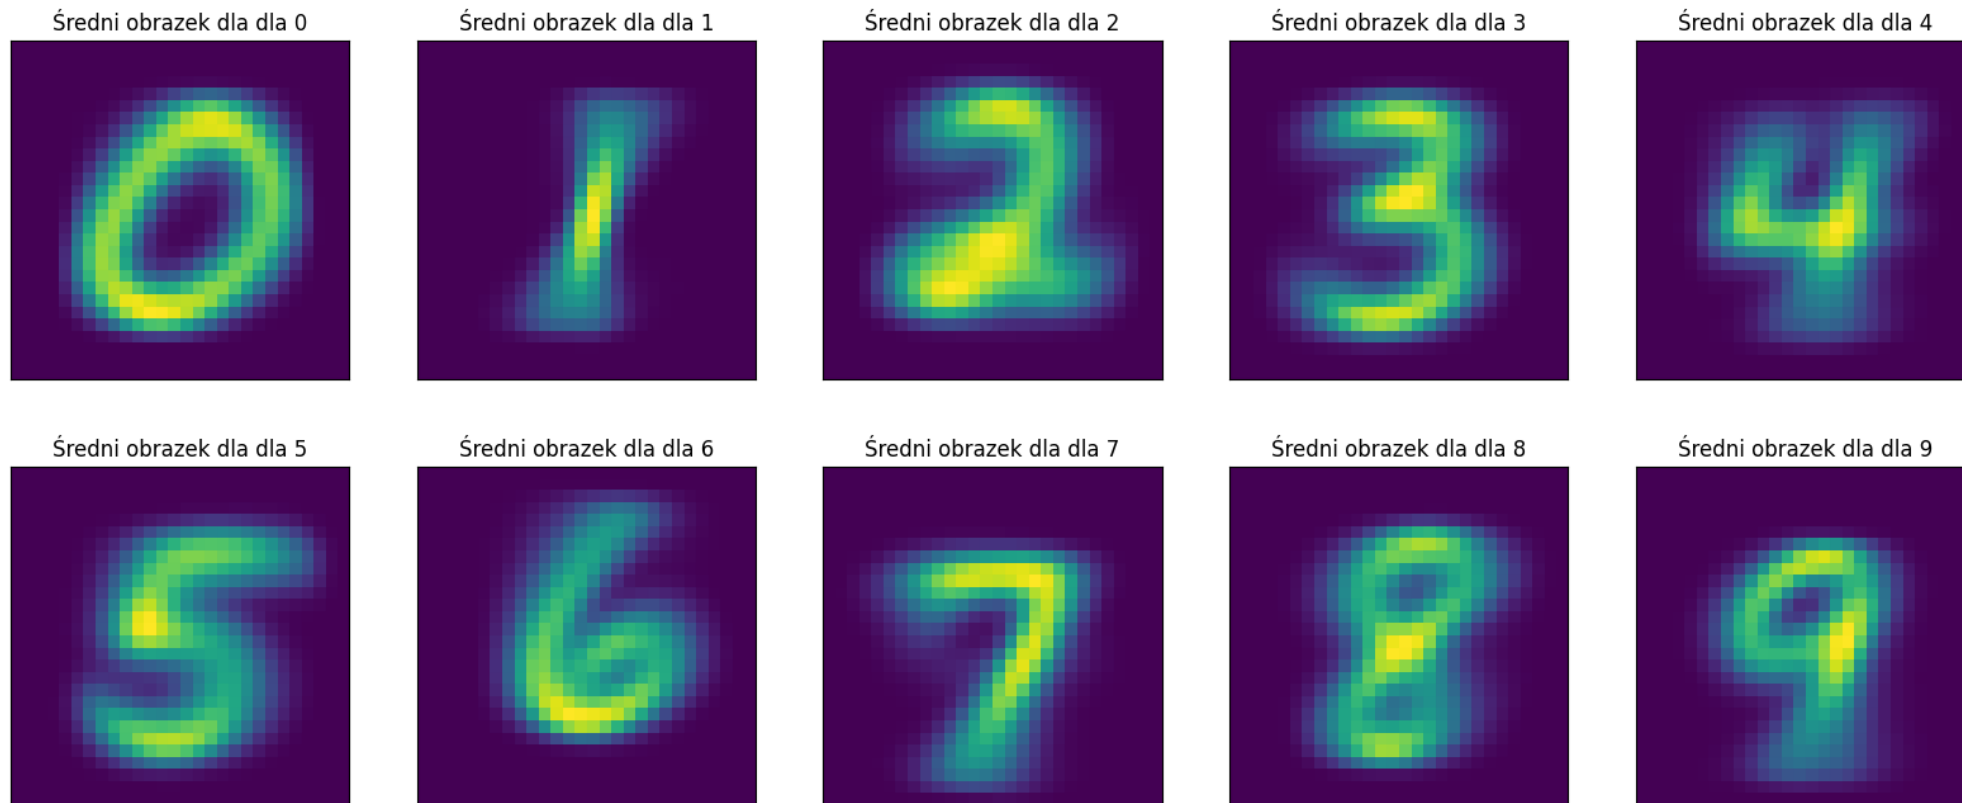
```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:,:,: ,label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:,:,: ,i] = avg_images[:,:,: ,i]/class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.1 A teraz sieć neuronowa

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(60,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()

model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))
```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 60)	47100
dense_2 (Dense)	(None, 10)	610

```

=====
Total params: 47710 (186.37 KB)
Trainable params: 47710 (186.37 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```

```

Epoch 1/10
844/844 [=====] - 4s 3ms/step - loss: 0.3693 - accuracy: 0.8986 - val_loss: 0.1734 - val_accuracy: 0.9537
Epoch 2/10
844/844 [=====] - 3s 4ms/step - loss: 0.1755 - accuracy: 0.9497 - val_loss: 0.1264 - val_accuracy: 0.9653
Epoch 3/10
844/844 [=====] - 3s 3ms/step - loss: 0.1296 - accuracy: 0.9629 - val_loss: 0.1046 - val_accuracy: 0.9707
Epoch 4/10
844/844 [=====] - 3s 3ms/step - loss: 0.1047 - accuracy: 0.9691 - val_loss: 0.0971 - val_accuracy: 0.9723
Epoch 5/10
844/844 [=====] - 3s 3ms/step - loss: 0.0858 - accuracy: 0.9745 - val_loss: 0.0955 - val_accuracy: 0.9720
Epoch 6/10
844/844 [=====] - 3s 4ms/step - loss: 0.0739 - accuracy: 0.9783 - val_loss: 0.0853 - val_accuracy: 0.9763
Epoch 7/10
844/844 [=====] - 3s 4ms/step - loss: 0.0627 - accuracy: 0.9816 - val_loss: 0.0924 - val_accuracy: 0.9735
Epoch 8/10
844/844 [=====] - 3s 3ms/step - loss: 0.0551 - accuracy: 0.9836 - val_loss: 0.0856 - val_accuracy: 0.9763
Epoch 9/10
844/844 [=====] - 3s 3ms/step - loss: 0.0483 - accuracy: 0.9854 - val_loss: 0.0906 - val_accuracy: 0.9750
Epoch 10/10
844/844 [=====] - 3s 3ms/step - loss: 0.0418 - accuracy: 0.9874 - val_loss: 0.0940 - val_accuracy: 0.9747

```

Precyzja

```



score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Precyzja: 0.9746000170707703

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```




▼ Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.1 Sieć neuronowa v2

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(64,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
opt = keras.optimizers.Adam(learning_rate=0.004)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???

history = model.fit(x_train, y_train, batch_size = 64, epochs = 13, validation_data = (x_valid,y_valid))
```

Model: "sequential_23"

Layer (type)	Output Shape	Param #
=====		
flatten_23 (Flatten)	(None, 784)	0
dense_37 (Dense)	(None, 64)	50240
dense_38 (Dense)	(None, 10)	650

=====

Total params: 50890 (198.79 KB)
 Trainable params: 50890 (198.79 KB)
 Non-trainable params: 0 (0.00 Byte)

Epoch 1/13
 844/844 [=====] - 3s 3ms/step - loss: 0.2624 - accuracy: 0.9215 - val_loss: 0.1202 - val_accuracy: 0.9658
 Epoch 2/13
 844/844 [=====] - 3s 3ms/step - loss: 0.1204 - accuracy: 0.9627 - val_loss: 0.1179 - val_accuracy: 0.9638
 Epoch 3/13
 844/844 [=====] - 4s 4ms/step - loss: 0.0912 - accuracy: 0.9722 - val_loss: 0.0917 - val_accuracy: 0.9738
 Epoch 4/13
 844/844 [=====] - 3s 3ms/step - loss: 0.0745 - accuracy: 0.9761 - val_loss: 0.1010 - val_accuracy: 0.9730
 Epoch 5/13
 844/844 [=====] - 3s 3ms/step - loss: 0.0638 - accuracy: 0.9795 - val_loss: 0.1001 - val_accuracy: 0.9728
 Epoch 6/13
 844/844 [=====] - 3s 3ms/step - loss: 0.0556 - accuracy: 0.9818 - val_loss: 0.0992 - val_accuracy: 0.9732
 Epoch 7/13
 844/844 [=====] - 3s 3ms/step - loss: 0.0452 - accuracy: 0.9851 - val_loss: 0.1033 - val_accuracy: 0.9720
 Epoch 8/13
 844/844 [=====] - 3s 4ms/step - loss: 0.0461 - accuracy: 0.9844 - val_loss: 0.1229 - val_accuracy: 0.9683
 Epoch 9/13
 844/844 [=====] - 3s 3ms/step - loss: 0.0395 - accuracy: 0.9866 - val_loss: 0.0974 - val_accuracy: 0.9768

```
Epoch 10/13
844/844 [=====] - 3s 3ms/step - loss: 0.0330 - accuracy: 0.9885 - val_loss: 0.1032 - val_accuracy: 0.9768
Epoch 11/13
844/844 [=====] - 3s 3ms/step - loss: 0.0317 - accuracy: 0.9893 - val_loss: 0.1274 - val_accuracy: 0.9725
Epoch 12/13
844/844 [=====] - 3s 4ms/step - loss: 0.0275 - accuracy: 0.9904 - val_loss: 0.1137 - val_accuracy: 0.9743
Epoch 13/13
844/844 [=====] - 3s 3ms/step - loss: 0.0234 - accuracy: 0.9920 - val_loss: 0.1196 - val_accuracy: 0.9765
```

Precyzja


```
score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])
```

```
Precyzja: 0.972100019454956
```


Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

Training and validation accuracy



Training and validation loss



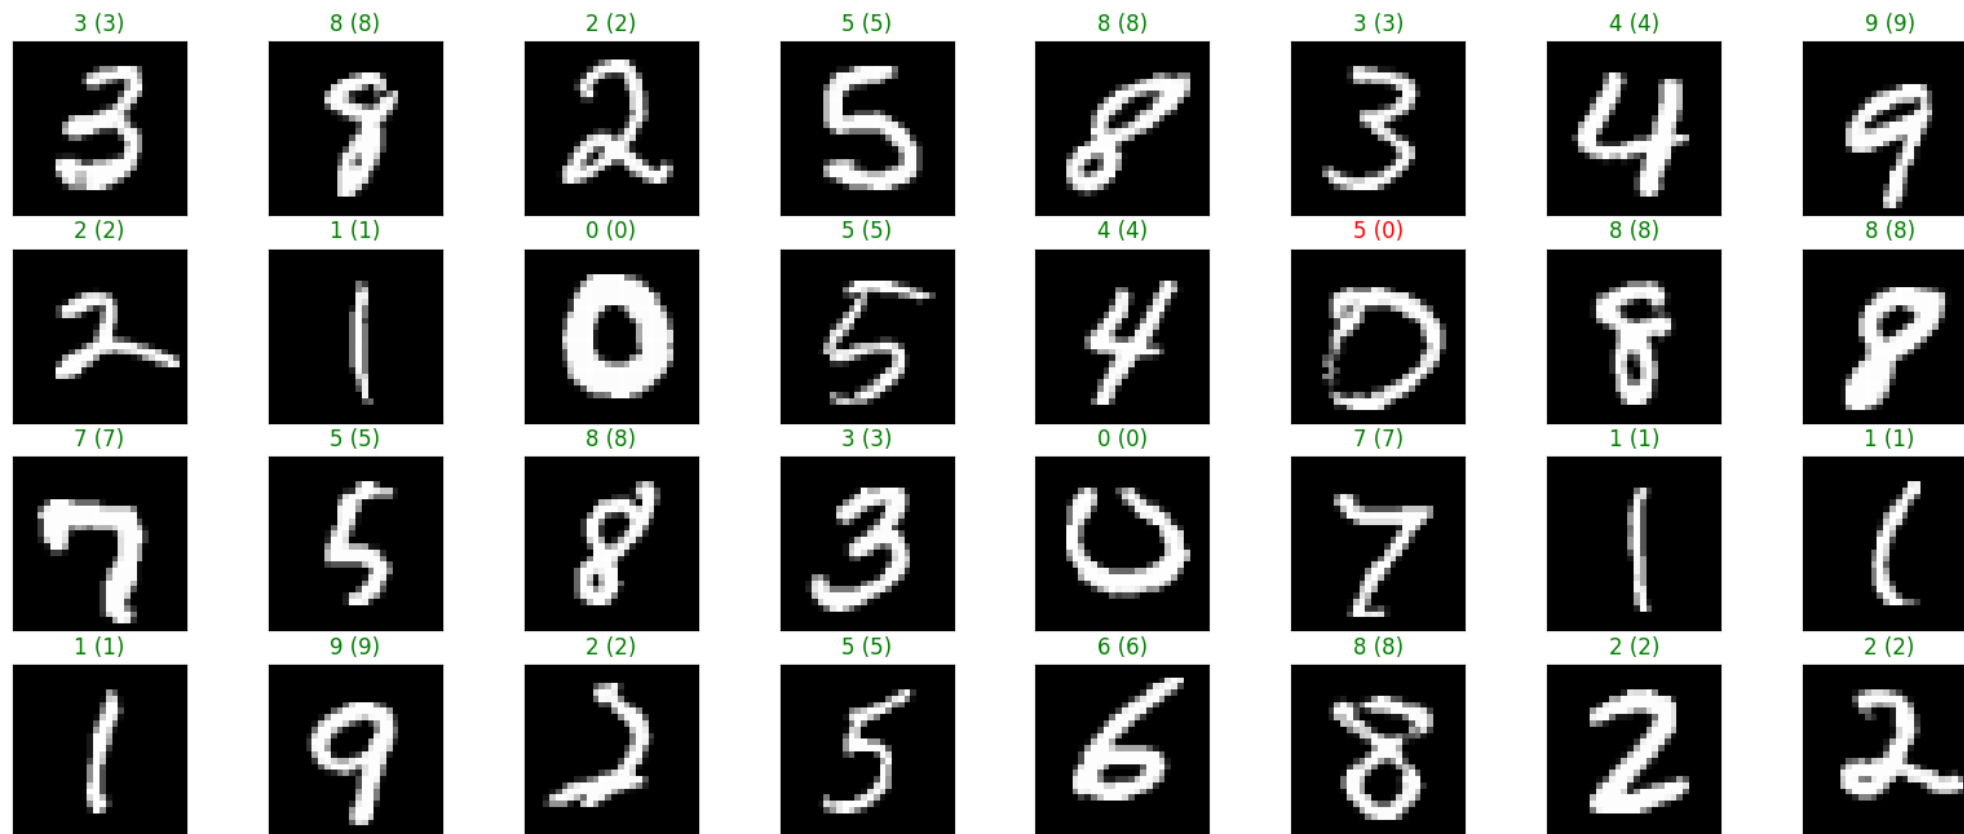
▼ Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 0s 1ms/step
```

test wyniki:



✓ 2.1 Sieć neuronowa v3

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(64,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
opt = keras.optimizers.SGD(learning_rate=0.08)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 128, epochs = 50, validation_data = (x_valid,y_valid))
```

```

422/422 [=====] - 1s 3ms/step - loss: 0.0447 - accuracy: 0.9882 - val_loss: 0.0841 - val_accuracy: 0.9700
Epoch 39/50
422/422 [=====] - 2s 4ms/step - loss: 0.0432 - accuracy: 0.9884 - val_loss: 0.0833 - val_accuracy: 0.9767
Epoch 40/50
422/422 [=====] - 2s 4ms/step - loss: 0.0420 - accuracy: 0.9893 - val_loss: 0.0826 - val_accuracy: 0.9770
Epoch 41/50
422/422 [=====] - 1s 3ms/step - loss: 0.0411 - accuracy: 0.9899 - val_loss: 0.0818 - val_accuracy: 0.9770
Epoch 42/50
422/422 [=====] - 1s 3ms/step - loss: 0.0399 - accuracy: 0.9898 - val_loss: 0.0821 - val_accuracy: 0.9767
Epoch 43/50
422/422 [=====] - 1s 3ms/step - loss: 0.0387 - accuracy: 0.9904 - val_loss: 0.0819 - val_accuracy: 0.9768
Epoch 44/50
422/422 [=====] - 1s 3ms/step - loss: 0.0376 - accuracy: 0.9905 - val_loss: 0.0819 - val_accuracy: 0.9777
Epoch 45/50
422/422 [=====] - 1s 3ms/step - loss: 0.0369 - accuracy: 0.9906 - val_loss: 0.0828 - val_accuracy: 0.9775
Epoch 46/50
422/422 [=====] - 1s 3ms/step - loss: 0.0359 - accuracy: 0.9909 - val_loss: 0.0825 - val_accuracy: 0.9782
Epoch 47/50
422/422 [=====] - 1s 3ms/step - loss: 0.0350 - accuracy: 0.9916 - val_loss: 0.0827 - val_accuracy: 0.9775
Epoch 48/50
422/422 [=====] - 2s 4ms/step - loss: 0.0337 - accuracy: 0.9917 - val_loss: 0.0836 - val_accuracy: 0.9773
Epoch 49/50
422/422 [=====] - 2s 4ms/step - loss: 0.0331 - accuracy: 0.9920 - val_loss: 0.0811 - val_accuracy: 0.9787
Epoch 50/50
422/422 [=====] - 1s 3ms/step - loss: 0.0322 - accuracy: 0.9924 - val_loss: 0.0821 - val_accuracy: 0.9773

```

Precyzja

```



score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

```
Precyzja: 0.9751999974250793
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```




✓ Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.2 Spróbujmy pogłębić nasz model!

Dodajmy 3 warstwy gęste.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 128)	100480
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 64)	4160
dense_6 (Dense)	(None, 32)	2080
dense_7 (Dense)	(None, 10)	330

```

=====
Total params: 115306 (450.41 KB)
Trainable params: 115306 (450.41 KB)
Non-trainable params: 0 (0.00 Byte)

```

```

Epoch 1/10
844/844 [=====] - 5s 4ms/step - loss: 0.3040 - accuracy: 0.9111 - val_loss: 0.1194 - val_accuracy: 0.9658
Epoch 2/10
844/844 [=====] - 4s 4ms/step - loss: 0.1197 - accuracy: 0.9641 - val_loss: 0.0991 - val_accuracy: 0.9698
Epoch 3/10
844/844 [=====] - 4s 4ms/step - loss: 0.0836 - accuracy: 0.9741 - val_loss: 0.0861 - val_accuracy: 0.9753
Epoch 4/10
844/844 [=====] - 3s 4ms/step - loss: 0.0677 - accuracy: 0.9789 - val_loss: 0.0818 - val_accuracy: 0.9770

```

```

Epoch 5/10
844/844 [=====] - 4s 5ms/step - loss: 0.0501 - accuracy: 0.9842 - val_loss: 0.0818 - val_accuracy: 0.9748
Epoch 6/10
844/844 [=====] - 4s 5ms/step - loss: 0.0426 - accuracy: 0.9868 - val_loss: 0.0935 - val_accuracy: 0.9730
Epoch 7/10
844/844 [=====] - 3s 4ms/step - loss: 0.0381 - accuracy: 0.9874 - val_loss: 0.0822 - val_accuracy: 0.9777
Epoch 8/10
844/844 [=====] - 3s 4ms/step - loss: 0.0318 - accuracy: 0.9898 - val_loss: 0.0943 - val_accuracy: 0.9745
Epoch 9/10
844/844 [=====] - 3s 4ms/step - loss: 0.0272 - accuracy: 0.9909 - val_loss: 0.0856 - val_accuracy: 0.9785
Epoch 10/10
844/844 [=====] - 4s 5ms/step - loss: 0.0283 - accuracy: 0.9908 - val_loss: 0.0873 - val_accuracy: 0.9783
Precyzja: 0.9760000109672546



```

Wykresy precyzji i błędu

```

import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()


```




```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.2 Pogłębinienie modelu v2

Dodajmy 2 warstwy gęste.

```
model = tf.keras.Sequential()
```

```

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_28"

Layer (type)	Output Shape	Param #
=====		
flatten_28 (Flatten)	(None, 784)	0
dense_47 (Dense)	(None, 128)	100480
dense_48 (Dense)	(None, 64)	8256
dense_49 (Dense)	(None, 32)	2080
dense_50 (Dense)	(None, 10)	330

=====

Total params: 111146 (434.16 KB)
 Trainable params: 111146 (434.16 KB)
 Non-trainable params: 0 (0.00 Byte)

```



Epoch 1/10
844/844 [=====] - 4s 4ms/step - loss: 0.2942 - accuracy: 0.9129 - val_loss: 0.1300 - val_accuracy: 0.9608
Epoch 2/10
844/844 [=====] - 4s 5ms/step - loss: 0.1168 - accuracy: 0.9649 - val_loss: 0.0922 - val_accuracy: 0.9728
Epoch 3/10
844/844 [=====] - 4s 4ms/step - loss: 0.0854 - accuracy: 0.9734 - val_loss: 0.0857 - val_accuracy: 0.9762
Epoch 4/10
844/844 [=====] - 3s 4ms/step - loss: 0.0623 - accuracy: 0.9807 - val_loss: 0.0902 - val_accuracy: 0.9745
Epoch 5/10
844/844 [=====] - 3s 4ms/step - loss: 0.0492 - accuracy: 0.9841 - val_loss: 0.1011 - val_accuracy: 0.9730
Epoch 6/10
844/844 [=====] - 3s 4ms/step - loss: 0.0413 - accuracy: 0.9867 - val_loss: 0.0784 - val_accuracy: 0.9792
Epoch 7/10
844/844 [=====] - 3s 4ms/step - loss: 0.0334 - accuracy: 0.9891 - val_loss: 0.0789 - val_accuracy: 0.9803
Epoch 8/10

```

```
844/844 [=====] - 3s 4ms/step - loss: 0.0277 - accuracy: 0.9909 - val_loss: 0.0969 - val_accuracy: 0.9768
Epoch 9/10
844/844 [=====] - 3s 4ms/step - loss: 0.0257 - accuracy: 0.9916 - val_loss: 0.0830 - val_accuracy: 0.9783
Epoch 10/10
844/844 [=====] - 4s 4ms/step - loss: 0.0213 - accuracy: 0.9923 - val_loss: 0.0925 - val_accuracy: 0.9780
Precyzja: 0.9775000214576721
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.2 Pogłębiamy model v3

Dodajmy 4 warstwy gęste.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
=====		
flatten_29 (Flatten)	(None, 784)	0
dense_51 (Dense)	(None, 128)	100480
dense_52 (Dense)	(None, 128)	16512
dense_53 (Dense)	(None, 64)	8256
dense_54 (Dense)	(None, 64)	4160
dense_55 (Dense)	(None, 32)	2080
dense_56 (Dense)	(None, 10)	330

=====



Total params: 131818 (514.91 KB)
Trainable params: 131818 (514.91 KB)
Non-trainable params: 0 (0.00 Byte)

Epoch 1/10
844/844 [=====] - 6s 5ms/step - loss: 0.3030 - accuracy: 0.9069 - val_loss: 0.1236 - val_accuracy: 0.9647
Epoch 2/10
844/844 [=====] - 3s 4ms/step - loss: 0.1200 - accuracy: 0.9635 - val_loss: 0.0996 - val_accuracy: 0.9737
Epoch 3/10
844/844 [=====] - 3s 4ms/step - loss: 0.0849 - accuracy: 0.9738 - val_loss: 0.1039 - val_accuracy: 0.9703
Epoch 4/10

```
844/844 [=====] - 3s 4ms/step - loss: 0.0672 - accuracy: 0.9789 - val_loss: 0.0725 - val_accuracy: 0.9778
Epoch 5/10
844/844 [=====] - 4s 5ms/step - loss: 0.0528 - accuracy: 0.9832 - val_loss: 0.0849 - val_accuracy: 0.9753
Epoch 6/10
844/844 [=====] - 3s 4ms/step - loss: 0.0474 - accuracy: 0.9846 - val_loss: 0.0802 - val_accuracy: 0.9768
Epoch 7/10
844/844 [=====] - 3s 4ms/step - loss: 0.0387 - accuracy: 0.9878 - val_loss: 0.0943 - val_accuracy: 0.9753
Epoch 8/10
844/844 [=====] - 4s 5ms/step - loss: 0.0339 - accuracy: 0.9890 - val_loss: 0.1035 - val_accuracy: 0.9742
Epoch 9/10
844/844 [=====] - 3s 4ms/step - loss: 0.0303 - accuracy: 0.9905 - val_loss: 0.0915 - val_accuracy: 0.9772
Epoch 10/10
844/844 [=====] - 3s 4ms/step - loss: 0.0286 - accuracy: 0.9909 - val_loss: 0.0797 - val_accuracy: 0.9810
Precyzja: 0.978600025177002
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```




```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.3 Konwolucja

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```
model = tf.keras.Sequential()
```

```
model.add(tf.keras.layers.Conv2D(filters=4, kernel_size=2, padding='same', activation='relu', input_shape=(28,28,1)))
```

```
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))
```

```
model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))
```

```
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))
```

```
model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 4)	20
max_pooling2d (MaxPooling2D)	(None, 14, 14, 4)	0
conv2d_1 (Conv2D)	(None, 14, 14, 2)	34
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 2)	0
flatten_3 (Flatten)	(None, 98)	0
dense_8 (Dense)	(None, 128)	12672
dense_9 (Dense)	(None, 64)	8256
dense_10 (Dense)	(None, 10)	650

```
=====  
Total params: 21632 (84.50 KB)  
Trainable params: 21632 (84.50 KB)  
Non-trainable params: 0 (0.00 Byte)
```

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```
Epoch 1/25  
422/422 [=====] - 6s 5ms/step - loss: 0.6456 - accuracy: 0.8120 - val_loss: 0.2564 - val_accuracy: 0.9202  
Epoch 2/25  
422/422 [=====] - 2s 4ms/step - loss: 0.2554 - accuracy: 0.9202 - val_loss: 0.1794 - val_accuracy: 0.9455  
Epoch 3/25  
422/422 [=====] - 2s 4ms/step - loss: 0.1903 - accuracy: 0.9398 - val_loss: 0.1564 - val_accuracy: 0.9518  
Epoch 4/25
```

```

422/422 [=====] - 2s 4ms/step - loss: 0.1560 - accuracy: 0.9503 - val_loss: 0.1370 - val_accuracy: 0.9587
Epoch 5/25
422/422 [=====] - 2s 4ms/step - loss: 0.1347 - accuracy: 0.9570 - val_loss: 0.1276 - val_accuracy: 0.9618
Epoch 6/25
422/422 [=====] - 2s 6ms/step - loss: 0.1176 - accuracy: 0.9624 - val_loss: 0.1130 - val_accuracy: 0.9675
Epoch 7/25
422/422 [=====] - 2s 4ms/step - loss: 0.1051 - accuracy: 0.9662 - val_loss: 0.1122 - val_accuracy: 0.9672
Epoch 8/25
422/422 [=====] - 2s 4ms/step - loss: 0.0954 - accuracy: 0.9688 - val_loss: 0.1011 - val_accuracy: 0.9723
Epoch 9/25
422/422 [=====] - 2s 4ms/step - loss: 0.0891 - accuracy: 0.9710 - val_loss: 0.1059 - val_accuracy: 0.9703
Epoch 10/25
422/422 [=====] - 2s 4ms/step - loss: 0.0811 - accuracy: 0.9736 - val_loss: 0.0932 - val_accuracy: 0.9740
Epoch 11/25
422/422 [=====] - 2s 4ms/step - loss: 0.0757 - accuracy: 0.9755 - val_loss: 0.1021 - val_accuracy: 0.9713
Epoch 12/25
422/422 [=====] - 2s 4ms/step - loss: 0.0702 - accuracy: 0.9768 - val_loss: 0.0919 - val_accuracy: 0.9745
Epoch 13/25
422/422 [=====] - 2s 6ms/step - loss: 0.0659 - accuracy: 0.9791 - val_loss: 0.1043 - val_accuracy: 0.9717
Epoch 14/25
422/422 [=====] - 2s 4ms/step - loss: 0.0615 - accuracy: 0.9796 - val_loss: 0.0942 - val_accuracy: 0.9735
Epoch 15/25
422/422 [=====] - 2s 4ms/step - loss: 0.0579 - accuracy: 0.9803 - val_loss: 0.0941 - val_accuracy: 0.9752
Epoch 16/25
422/422 [=====] - 2s 4ms/step - loss: 0.0523 - accuracy: 0.9828 - val_loss: 0.0899 - val_accuracy: 0.9760
Epoch 17/25
422/422 [=====] - 2s 4ms/step - loss: 0.0494 - accuracy: 0.9837 - val_loss: 0.0909 - val_accuracy: 0.9752
Epoch 18/25
422/422 [=====] - 2s 4ms/step - loss: 0.0469 - accuracy: 0.9847 - val_loss: 0.0942 - val_accuracy: 0.9740
Epoch 19/25
422/422 [=====] - 2s 5ms/step - loss: 0.0436 - accuracy: 0.9854 - val_loss: 0.0986 - val_accuracy: 0.9753
Epoch 20/25
422/422 [=====] - 2s 6ms/step - loss: 0.0408 - accuracy: 0.9864 - val_loss: 0.0929 - val_accuracy: 0.9763
Epoch 21/25
422/422 [=====] - 2s 4ms/step - loss: 0.0389 - accuracy: 0.9869 - val_loss: 0.1009 - val_accuracy: 0.9770
Epoch 22/25
422/422 [=====] - 2s 4ms/step - loss: 0.0362 - accuracy: 0.9876 - val_loss: 0.1025 - val_accuracy: 0.9748
Epoch 23/25
422/422 [=====] - 2s 4ms/step - loss: 0.0345 - accuracy: 0.9885 - val_loss: 0.1083 - val_accuracy: 0.9752
Epoch 24/25
422/422 [=====] - 2s 4ms/step - loss: 0.0320 - accuracy: 0.9896 - val_loss: 0.1095 - val_accuracy: 0.9757
Epoch 25/25
422/422 [=====] - 2s 4ms/step - loss: 0.0312 - accuracy: 0.9893 - val_loss: 0.1088 - val_accuracy: 0.9732

```

▼ Ewaluacja modelu

```



score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ', score[1])

```

Precyzja: 0.9750000238418579

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test,"convnet")
```

313/313 [=====] - 1s 3ms/step

convnet wyniki:



✓ 2.3 Konwolucja model v2

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=4, kernel_size=2, padding='same', activation='relu', input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_33"

Layer (type)	Output Shape	Param #
=====		
conv2d_14 (Conv2D)	(None, 28, 28, 4)	20
max_pooling2d_11 (MaxPooling2D)	(None, 14, 14, 4)	0
conv2d_15 (Conv2D)	(None, 14, 14, 2)	34
max_pooling2d_12 (MaxPooling2D)	(None, 7, 7, 2)	0
conv2d_16 (Conv2D)	(None, 7, 7, 2)	18
max_pooling2d_13 (MaxPooling2D)	(None, 3, 3, 2)	0
flatten_33 (Flatten)	(None, 18)	0
dense_66 (Dense)	(None, 128)	2432
dense_67 (Dense)	(None, 64)	8256
dense_68 (Dense)	(None, 10)	650

```
=====
Total params: 11410 (44.57 KB)
Trainable params: 11410 (44.57 KB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 4s 6ms/step - loss: 0.8775 - accuracy: 0.7366 - val_loss: 0.3990 - val_accuracy: 0.8743
Epoch 2/25
422/422 [=====] - 3s 6ms/step - loss: 0.4237 - accuracy: 0.8690 - val_loss: 0.3221 - val_accuracy: 0.8957
Epoch 3/25
422/422 [=====] - 2s 4ms/step - loss: 0.3533 - accuracy: 0.8911 - val_loss: 0.2699 - val_accuracy: 0.9163
Epoch 4/25
422/422 [=====] - 2s 5ms/step - loss: 0.3125 - accuracy: 0.9053 - val_loss: 0.2398 - val_accuracy: 0.9233
Epoch 5/25
422/422 [=====] - 2s 5ms/step - loss: 0.2787 - accuracy: 0.9141 - val_loss: 0.2199 - val_accuracy: 0.9333
Epoch 6/25
422/422 [=====] - 2s 5ms/step - loss: 0.2535 - accuracy: 0.9216 - val_loss: 0.2120 - val_accuracy: 0.9320
Epoch 7/25
422/422 [=====] - 2s 4ms/step - loss: 0.2363 - accuracy: 0.9268 - val_loss: 0.1925 - val_accuracy: 0.9370
Epoch 8/25
422/422 [=====] - 3s 6ms/step - loss: 0.2213 - accuracy: 0.9307 - val_loss: 0.1857 - val_accuracy: 0.9400
Epoch 9/25
422/422 [=====] - 2s 5ms/step - loss: 0.2109 - accuracy: 0.9335 - val_loss: 0.1840 - val_accuracy: 0.9407
Epoch 10/25
422/422 [=====] - 2s 5ms/step - loss: 0.1997 - accuracy: 0.9374 - val_loss: 0.1729 - val_accuracy: 0.9435
Epoch 11/25
422/422 [=====] - 2s 5ms/step - loss: 0.1901 - accuracy: 0.9395 - val_loss: 0.1750 - val_accuracy: 0.9435
Epoch 12/25
422/422 [=====] - 2s 4ms/step - loss: 0.1848 - accuracy: 0.9411 - val_loss: 0.1595 - val_accuracy: 0.9478
Epoch 13/25
422/422 [=====] - 2s 4ms/step - loss: 0.1783 - accuracy: 0.9439 - val_loss: 0.1611 - val_accuracy: 0.9467
Epoch 14/25
422/422 [=====] - 3s 8ms/step - loss: 0.1728 - accuracy: 0.9449 - val_loss: 0.1623 - val_accuracy: 0.9458
Epoch 15/25
422/422 [=====] - 2s 6ms/step - loss: 0.1694 - accuracy: 0.9457 - val_loss: 0.1513 - val_accuracy: 0.9522
Epoch 16/25
422/422 [=====] - 2s 6ms/step - loss: 0.1626 - accuracy: 0.9479 - val_loss: 0.1536 - val_accuracy: 0.9505
Epoch 17/25
422/422 [=====] - 2s 5ms/step - loss: 0.1605 - accuracy: 0.9482 - val_loss: 0.1481 - val_accuracy: 0.9528
Epoch 18/25
422/422 [=====] - 3s 7ms/step - loss: 0.1578 - accuracy: 0.9501 - val_loss: 0.1495 - val_accuracy: 0.9508
Epoch 19/25
422/422 [=====] - 4s 8ms/step - loss: 0.1535 - accuracy: 0.9511 - val_loss: 0.1503 - val_accuracy: 0.9513
Epoch 20/25
422/422 [=====] - 2s 6ms/step - loss: 0.1507 - accuracy: 0.9519 - val_loss: 0.1469 - val_accuracy: 0.9523
Epoch 21/25
422/422 [=====] - 2s 5ms/step - loss: 0.1488 - accuracy: 0.9530 - val_loss: 0.1425 - val_accuracy: 0.9537
```



```
Epoch 22/25
422/422 [=====] - 2s 5ms/step - loss: 0.1447 - accuracy: 0.9536 - val_loss: 0.1412 - val_accuracy: 0.9553
Epoch 23/25
422/422 [=====] - 3s 6ms/step - loss: 0.1428 - accuracy: 0.9538 - val_loss: 0.1436 - val_accuracy: 0.9560
Epoch 24/25
422/422 [=====] - 3s 8ms/step - loss: 0.1421 - accuracy: 0.9536 - val_loss: 0.1441 - val_accuracy: 0.9540
Epoch 25/25
422/422 [=====] - 3s 6ms/step - loss: 0.1373 - accuracy: 0.9554 - val_loss: 0.1466 - val_accuracy: 0.9518
```



✓ Ewaluacja modelu

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ', score[1])
```

Precyzja: 0.9480000138282776

Wykresy precyzji i błędu


```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

```
313/313 [======] - 1s 2ms/step
```

convnet wyniki:



✓ 2.3 Konwolucja model v3

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2, padding='same', activation='relu', input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

Model: "sequential_34"

Layer (type)	Output Shape	Param #
=====		
conv2d_17 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_14 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_18 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_15 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_19 (Conv2D)	(None, 7, 7, 16)	2064
max_pooling2d_16 (MaxPooling2D)	(None, 3, 3, 16)	0
flatten_34 (Flatten)	(None, 144)	0
dense_69 (Dense)	(None, 128)	18560
dense_70 (Dense)	(None, 64)	8256
dense_71 (Dense)	(None, 10)	650

=====

Total params: 38074 (148.73 KB)
Trainable params: 38074 (148.73 KB)

Non-trainable params: 0 (0.00 Byte)

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 6s 8ms/step - loss: 0.3922 - accuracy: 0.8789 - val_loss: 0.1028 - val_accuracy: 0.9685
Epoch 2/25
422/422 [=====] - 3s 6ms/step - loss: 0.1002 - accuracy: 0.9681 - val_loss: 0.0702 - val_accuracy: 0.9792
Epoch 3/25
422/422 [=====] - 3s 7ms/step - loss: 0.0714 - accuracy: 0.9776 - val_loss: 0.0678 - val_accuracy: 0.9785
Epoch 4/25
422/422 [=====] - 3s 8ms/step - loss: 0.0596 - accuracy: 0.9811 - val_loss: 0.0638 - val_accuracy: 0.9812
Epoch 5/25
422/422 [=====] - 4s 8ms/step - loss: 0.0479 - accuracy: 0.9846 - val_loss: 0.0465 - val_accuracy: 0.9845
Epoch 6/25
422/422 [=====] - 3s 7ms/step - loss: 0.0433 - accuracy: 0.9863 - val_loss: 0.0469 - val_accuracy: 0.9850
Epoch 7/25
422/422 [=====] - 3s 7ms/step - loss: 0.0377 - accuracy: 0.9880 - val_loss: 0.0464 - val_accuracy: 0.9853
Epoch 8/25
422/422 [=====] - 3s 8ms/step - loss: 0.0327 - accuracy: 0.9897 - val_loss: 0.0532 - val_accuracy: 0.9843
Epoch 9/25
422/422 [=====] - 4s 9ms/step - loss: 0.0311 - accuracy: 0.9898 - val_loss: 0.0478 - val_accuracy: 0.9865
Epoch 10/25
422/422 [=====] - 2s 6ms/step - loss: 0.0280 - accuracy: 0.9911 - val_loss: 0.0412 - val_accuracy: 0.9878
Epoch 11/25
422/422 [=====] - 2s 6ms/step - loss: 0.0262 - accuracy: 0.9917 - val_loss: 0.0480 - val_accuracy: 0.9845
Epoch 12/25
422/422 [=====] - 2s 6ms/step - loss: 0.0230 - accuracy: 0.9922 - val_loss: 0.0362 - val_accuracy: 0.9902
Epoch 13/25
422/422 [=====] - 3s 7ms/step - loss: 0.0209 - accuracy: 0.9932 - val_loss: 0.0338 - val_accuracy: 0.9893
Epoch 14/25
422/422 [=====] - 3s 6ms/step - loss: 0.0201 - accuracy: 0.9928 - val_loss: 0.0513 - val_accuracy: 0.9872
Epoch 15/25
422/422 [=====] - 2s 6ms/step - loss: 0.0180 - accuracy: 0.9938 - val_loss: 0.0433 - val_accuracy: 0.9898
Epoch 16/25
422/422 [=====] - 2s 6ms/step - loss: 0.0158 - accuracy: 0.9947 - val_loss: 0.0382 - val_accuracy: 0.9888
Epoch 17/25
422/422 [=====] - 2s 6ms/step - loss: 0.0141 - accuracy: 0.9951 - val_loss: 0.0341 - val_accuracy: 0.9910
Epoch 18/25
422/422 [=====] - 3s 7ms/step - loss: 0.0144 - accuracy: 0.9950 - val_loss: 0.0321 - val_accuracy: 0.9915
Epoch 19/25
422/422 [=====] - 3s 6ms/step - loss: 0.0116 - accuracy: 0.9960 - val_loss: 0.0480 - val_accuracy: 0.9872
Epoch 20/25
422/422 [=====] - 2s 6ms/step - loss: 0.0122 - accuracy: 0.9958 - val_loss: 0.0483 - val_accuracy: 0.9877
Epoch 21/25
422/422 [=====] - 2s 6ms/step - loss: 0.0130 - accuracy: 0.9955 - val_loss: 0.0352 - val_accuracy: 0.9908
Epoch 22/25
422/422 [=====] - 3s 6ms/step - loss: 0.0113 - accuracy: 0.9960 - val_loss: 0.0423 - val_accuracy: 0.9888
Epoch 23/25
422/422 [=====] - 3s 7ms/step - loss: 0.0110 - accuracy: 0.9966 - val_loss: 0.0361 - val_accuracy: 0.9910
```

Epoch 24/25

422/422 [=====] - 3s 6ms/step - loss: 0.0087 - accuracy: 0.9970 - val_loss: 0.0417 - val_accuracy: 0.9892

Epoch 25/25

422/422 [=====] - 2s 6ms/step - loss: 0.0104 - accuracy: 0.9965 - val_loss: 0.0380 - val_accuracy: 0.9905



✓ Ewaluacja modelu

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ', score[1])
```

Precyzja: 0.9902999997138977

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

```
313/313 [=====] - 1s 2ms/step
```

convnet wyniki:



3 Regularyzacja

Nasz model ma obecnie dużo stopni swobody (ma DUŻO parametrów i dlatego może dopasować się do niemal każdej funkcji, jeśli tylko będziemy trenować wystarczająco długo). Oznacza to, że nasza sieć jest również podatna na przeuczenie.

W tej sekcji dodajmy warstwy dropout pomiędzy głównymi warstwami naszej sieci, aby uniknąć przeuczenia.


```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2, padding='same', activation='relu', input_shape=(28, 28, 1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2, padding='same', activation='relu'))

model.add(tf.keras.layers.Flatten())

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_30"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_4 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_3 (Dropout)	(None, 14, 14, 64)	0
conv2d_6 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_5 (MaxPooling2D)	(None, 7, 7, 32)	0
dropout_4 (Dropout)	(None, 7, 7, 32)	0
conv2d_7 (Conv2D)	(None, 7, 7, 16)	2064
flatten_30 (Flatten)	(None, 784)	0
dropout_5 (Dropout)	(None, 784)	0

dense_57 (Dense)	(None, 128)	100480
dense_58 (Dense)	(None, 64)	8256
dense_59 (Dense)	(None, 10)	650

```
=====
Total params: 119994 (468.73 KB)
Trainable params: 119994 (468.73 KB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
model.fit(x_train,
          y_train,
          batch_size=128,
          epochs=25,
          validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 6s 9ms/step - loss: 0.6871 - accuracy: 0.7697 - val_loss: 0.1493 - val_accuracy: 0.9603
Epoch 2/25
422/422 [=====] - 4s 10ms/step - loss: 0.2656 - accuracy: 0.9146 - val_loss: 0.0969 - val_accuracy: 0.9705
Epoch 3/25
422/422 [=====] - 3s 7ms/step - loss: 0.2066 - accuracy: 0.9339 - val_loss: 0.0831 - val_accuracy: 0.9742
Epoch 4/25
422/422 [=====] - 3s 7ms/step - loss: 0.1753 - accuracy: 0.9437 - val_loss: 0.0647 - val_accuracy: 0.9808
Epoch 5/25
422/422 [=====] - 3s 8ms/step - loss: 0.1546 - accuracy: 0.9506 - val_loss: 0.0589 - val_accuracy: 0.9847
Epoch 6/25
422/422 [=====] - 3s 7ms/step - loss: 0.1381 - accuracy: 0.9556 - val_loss: 0.0532 - val_accuracy: 0.9835
Epoch 7/25
422/422 [=====] - 3s 7ms/step - loss: 0.1296 - accuracy: 0.9588 - val_loss: 0.0505 - val_accuracy: 0.9863
Epoch 8/25
422/422 [=====] - 3s 7ms/step - loss: 0.1240 - accuracy: 0.9605 - val_loss: 0.0492 - val_accuracy: 0.9860
Epoch 9/25
422/422 [=====] - 3s 8ms/step - loss: 0.1163 - accuracy: 0.9628 - val_loss: 0.0448 - val_accuracy: 0.9877
Epoch 10/25
422/422 [=====] - 3s 7ms/step - loss: 0.1094 - accuracy: 0.9654 - val_loss: 0.0418 - val_accuracy: 0.9873
Epoch 11/25
422/422 [=====] - 3s 7ms/step - loss: 0.1035 - accuracy: 0.9666 - val_loss: 0.0421 - val_accuracy: 0.9882
Epoch 12/25
422/422 [=====] - 3s 7ms/step - loss: 0.1044 - accuracy: 0.9664 - val_loss: 0.0404 - val_accuracy: 0.9888
Epoch 13/25
422/422 [=====] - 3s 8ms/step - loss: 0.1016 - accuracy: 0.9683 - val_loss: 0.0393 - val_accuracy: 0.9888
Epoch 14/25
422/422 [=====] - 3s 7ms/step - loss: 0.0960 - accuracy: 0.9684 - val_loss: 0.0390 - val_accuracy: 0.9883
Epoch 15/25
422/422 [=====] - 3s 7ms/step - loss: 0.0930 - accuracy: 0.9702 - val_loss: 0.0351 - val_accuracy: 0.9905
Epoch 16/25
422/422 [=====] - 3s 7ms/step - loss: 0.0900 - accuracy: 0.9710 - val_loss: 0.0358 - val_accuracy: 0.9898
```

```

Epoch 17/25
422/422 [=====] - 3s 8ms/step - loss: 0.0872 - accuracy: 0.9719 - val_loss: 0.0354 - val_accuracy: 0.9890
Epoch 18/25
422/422 [=====] - 3s 8ms/step - loss: 0.0874 - accuracy: 0.9716 - val_loss: 0.0334 - val_accuracy: 0.9912
Epoch 19/25
422/422 [=====] - 3s 7ms/step - loss: 0.0854 - accuracy: 0.9728 - val_loss: 0.0352 - val_accuracy: 0.9895
Epoch 20/25
422/422 [=====] - 3s 7ms/step - loss: 0.0838 - accuracy: 0.9731 - val_loss: 0.0358 - val_accuracy: 0.9895
Epoch 21/25
422/422 [=====] - 3s 8ms/step - loss: 0.0793 - accuracy: 0.9743 - val_loss: 0.0313 - val_accuracy: 0.9907
Epoch 22/25
422/422 [=====] - 3s 8ms/step - loss: 0.0810 - accuracy: 0.9736 - val_loss: 0.0330 - val_accuracy: 0.9898
Epoch 23/25
422/422 [=====] - 3s 7ms/step - loss: 0.0761 - accuracy: 0.9749 - val_loss: 0.0343 - val_accuracy: 0.9903
Epoch 24/25
422/422 [=====] - 3s 7ms/step - loss: 0.0764 - accuracy: 0.9749 - val_loss: 0.0303 - val_accuracy: 0.9905
Epoch 25/25
422/422 [=====] - 3s 8ms/step - loss: 0.0739 - accuracy: 0.9760 - val_loss: 0.0310 - val_accuracy: 0.9915
<keras.src.callbacks.History at 0x7fdb787b2e60>

```

✓ Evaluate model:

```

test_score = model.evaluate(x_test, y_test, verbose=0)
train_score = model.evaluate(x_train, y_train, verbose=0)



print('Train accuracy: ',train_score[1], ' Test accuracy: ',test_score[1])

Train accuracy:  0.9929259419441223  Test accuracy:  0.9897000193595886

```

Wykresy precyzji i błędu

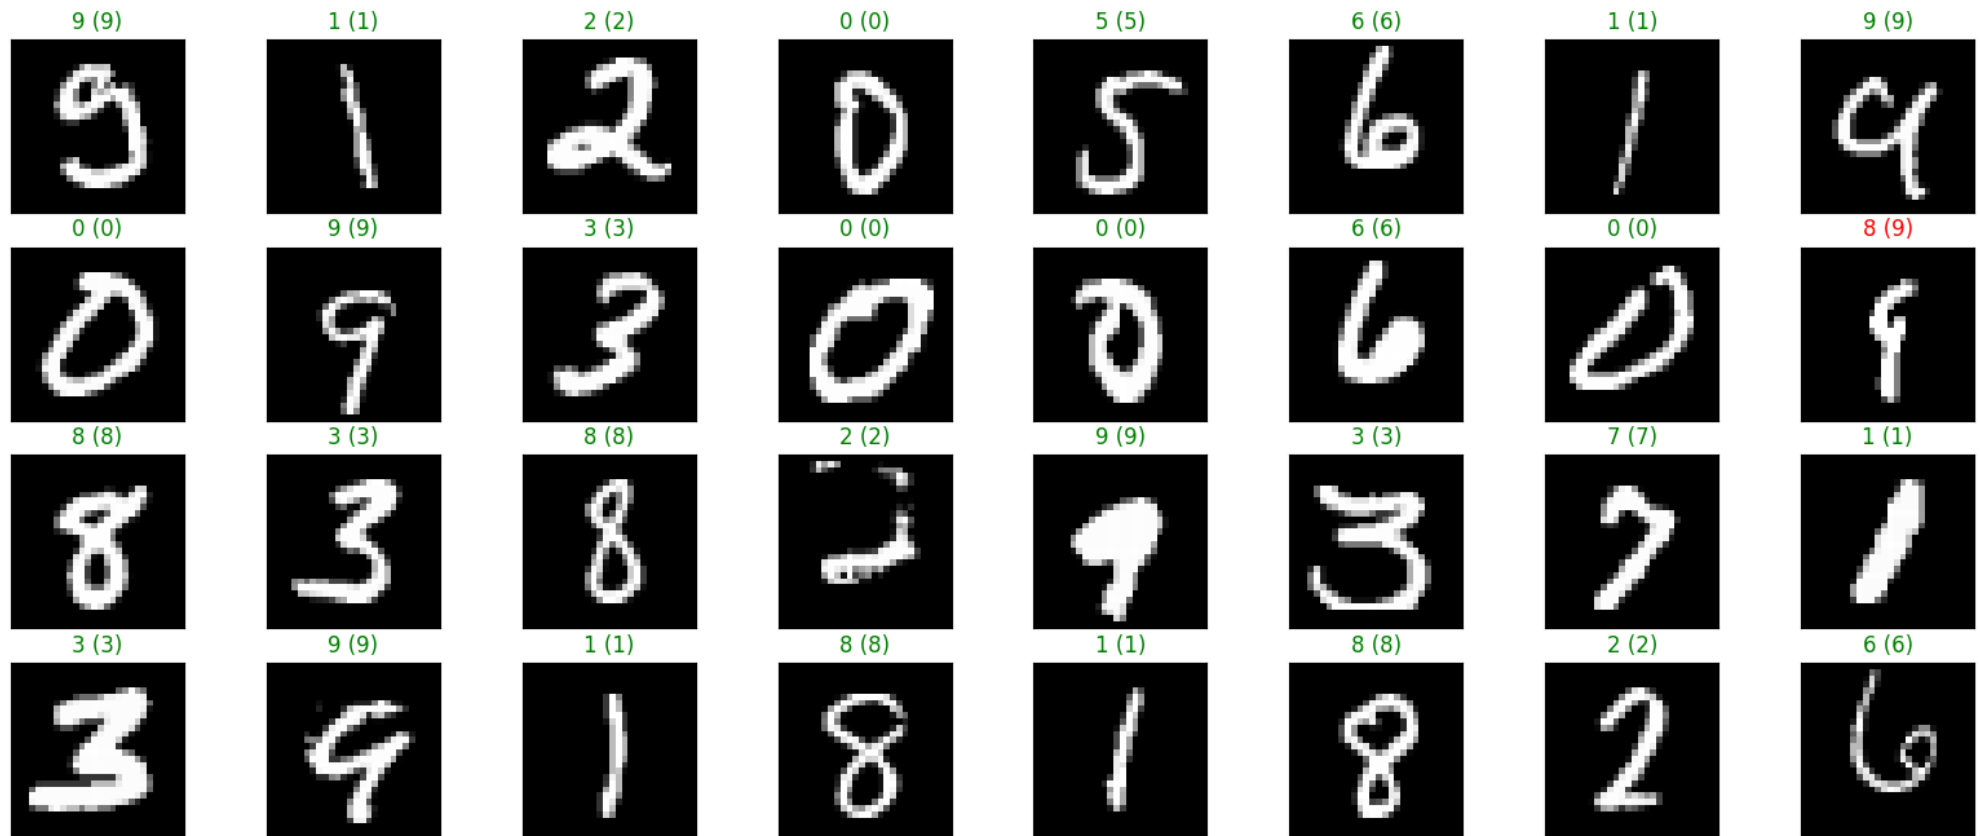
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

313/313 [=====] - 1s 2ms/step

convnet wyniki:



3 Regularyzacja v2

Nasz model ma obecnie dużo stopni swobody (ma DUŻO parametrów i dlatego może dopasować się do niemal każdej funkcji, jeśli tylko będziemy trenować wystarczająco długo). Oznacza to, że nasza sieć jest również podatna na przeuczenie.

W tej sekcji dodajmy warstwy dropout pomiędzy warstwami naszej sieci, aby uniknąć przeuczenia.

```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2, padding='same', activation='relu', input_shape=(28, 28, 1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2, padding='same', activation='relu'))

model.add(tf.keras.layers.Flatten())

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dropout(0.1))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_31"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_6 (Dropout)	(None, 14, 14, 64)	0
conv2d_9 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 32)	0