

CNN na przykładzie MNIST

✓ Setup

Importujemy potrzebne biblioteki

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import keras
```

✓ 1 Przygotowanie danych

✓ 1.0 Pobranie zbioru danych

Pobieramy zbiór danych i sprawdzamy rozmiar 28 x 28 pixeli.

```
(x_train_data, y_train_data), (x_test_data, y_test_data) = tf.keras.datasets.fashion_mnist.load_data()
```

```
dataset_labels = ["0", # index 0
                  "1", # index 1
                  "2", # index 2
                  "3", # index 3
                  "4", # index 4
                  "5", # index 5
                  "6", # index 6
                  "7", # index 7
                  "8", # index 8
                  "9"] # index 9
```

```
print("x_train shape:", x_train_data.shape, "y_train shape:", y_train_data.shape)
print("x_test shape:", x_test_data.shape, "y_test shape:", y_test_data.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 [=====] - 0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
```


```
26421880/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 [=====] - 0s 0us/step
x_train shape: (60000, 28, 28) y_train shape: (60000,)
x_test shape: (10000, 28, 28) y_test shape: (10000,)
```

✓ 1.1 Wizualizacja danych

Przykładowy obrazek ze bioru danych

```
def plot_image(img_index):
    label_index = y_train_data[img_index]
    plt.imshow(x_train_data[img_index]/255, cmap = 'gray')
    plt.title("Image "+str(img_index)+": "+dataset_labels[label_index])
```

```
img_index = 60
plot_image(img_index)
```



✓ 1.2 Normalizacja danych

Na początek sprawdzamy jakie są max i min wartości pixeli w obrazkach.

Wartości te powinny być zawarte w przedziale [0,1].

```
print("Wartości min:", np.min(x_train_data), " max:", np.max(x_train_data))

x_train_data = x_train_data.astype('float32') / 255
x_test_data = x_test_data.astype('float32') / 255

print("Wartości po przeskalowaniu min:", np.min(x_train_data), " max:", np.max(x_train_data))

Wartości min: 0 max: 255
Wartości po przeskalowaniu min: 0.0 max: 1.0
```

✓ 1.3 Podział zbioru danych na zbiór treningowy/walidacyjny/testowy

- **Zbiór treningowy** - wykorzystamy go do uczenia.
- **Zbiór walidacyjny** - wykorzystamy go do tuningu hiperparametrów.
- **Zbiór testowy** - wykorzystamy go do ostatecznego sprawdzenia modelu.

Zbiór walidacyjny stworzymy z 10% zbioru treningowego.

```
validation_fraction = .1

total_train_samples = len(x_train_data)
validation_samples = int(total_train_samples * validation_fraction)
train_samples = total_train_samples - validation_samples

(x_train, x_valid) = x_train_data[:train_samples], x_train_data[train_samples:]
(y_train, y_valid) = y_train_data[:train_samples], y_train_data[train_samples:]

x_test, y_test = x_test_data, y_test_data
print(train_samples, validation_samples, len(x_test))

54000 6000 10000
```

✓ 1.4 Dwa dodatkowe kroki

- Większość zestawów danych obrazu składa się z obrazów rgb. Z tego powodu Keras oczekuje, że każdy obraz będzie miał 3 wymiary: [x_pixels, y_pixels, color_channels]. Ponieważ nasze obrazki są w skali szarości, wymiar koloru jest równy 1. Musimy zatem zmienić kształt obrazków.
- W procesie uczenia naszego modelu będziemy wykorzystywali tzw. **kategoryczną entropię krzyżową** (<https://keras.io/losses/>). Musimy przekształcić wektory z etykietami (labelami) do **kodowania one-hot**. Wykorzystamy do tego funkcję `tf.keras.utils.to_categorical()`.

```
# Zmieniamy kształt z (28, 28) na (28, 28, 1)
w, h = 28, 28
x_train = x_train.reshape(x_train.shape[0], w, h, 1)
x_valid = x_valid.reshape(x_valid.shape[0], w, h, 1)
x_test = x_test.reshape(x_test.shape[0], w, h, 1)

# Kodowanie one-hot
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_valid = tf.keras.utils.to_categorical(y_valid, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)

print("x_train shape:", x_train.shape, "y_train shape:", y_train.shape)

# Ilość elementów w zbiorach
print(x_train.shape[0], 'train set')
print(x_valid.shape[0], 'validation set')
print(x_test.shape[0], 'test set')

x_train shape: (54000, 28, 28, 1) y_train shape: (54000, 10)
54000 train set
6000 validation set
10000 test set
```

✓ 2 Stworzenie modelu

Keras oferuje dwa API:

1. [Sequential model API](#)
2. [Functional API](#)

W naszym modelu wykorzystamy Sequential model API. Będziemy wykorzystywali następujące metody:

- `Dense()` [link text](#) - tworzy **warstwę gęstą**
- `Conv2D()` [link text](#) - tworzy **warstwę konwolucyjną**
- `Pooling()` [link text](#) - tworzy **warstwę pooling**

- Dropout() [link text](#) - zastosowanie **dropout**

✓ 2.0 Prosty model liniowy

Zacniemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przształci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 10)	7850
=====		
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optimizer adam**
- Jako loss function użyjemy '**categorical_crossentropy**'
- Lista parametrów, tutaj zaczniemy od '**precyzji**'

Warto zerknąć: <https://keras.io/models/model/>

```
model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???
```

✓ Uczenie modelu

Model uczymy wykorzystując fit().



```
history = model.fit(x_train, y_train, batch_size = 64, epochs = 5, validation_data = (x_valid,y_valid))
```

```
Epoch 1/5
844/844 [=====] - 11s 10ms/step - loss: 0.6659 - accuracy: 0.7815 - val_loss: 0.5227 - val_accuracy: 0.8168
Epoch 2/5
844/844 [=====] - 5s 6ms/step - loss: 0.4892 - accuracy: 0.8341 - val_loss: 0.4733 - val_accuracy: 0.8378
Epoch 3/5
844/844 [=====] - 3s 3ms/step - loss: 0.4558 - accuracy: 0.8447 - val_loss: 0.4471 - val_accuracy: 0.8448
Epoch 4/5
844/844 [=====] - 3s 4ms/step - loss: 0.4362 - accuracy: 0.8511 - val_loss: 0.4455 - val_accuracy: 0.8442
Epoch 5/5
844/844 [=====] - 2s 3ms/step - loss: 0.4250 - accuracy: 0.8549 - val_loss: 0.4297 - val_accuracy: 0.8505
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`.
saving_api.save_model(
```

Wczytanie modelu

```
#from keras.models import load_model
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.8393999934196472
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'


```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])

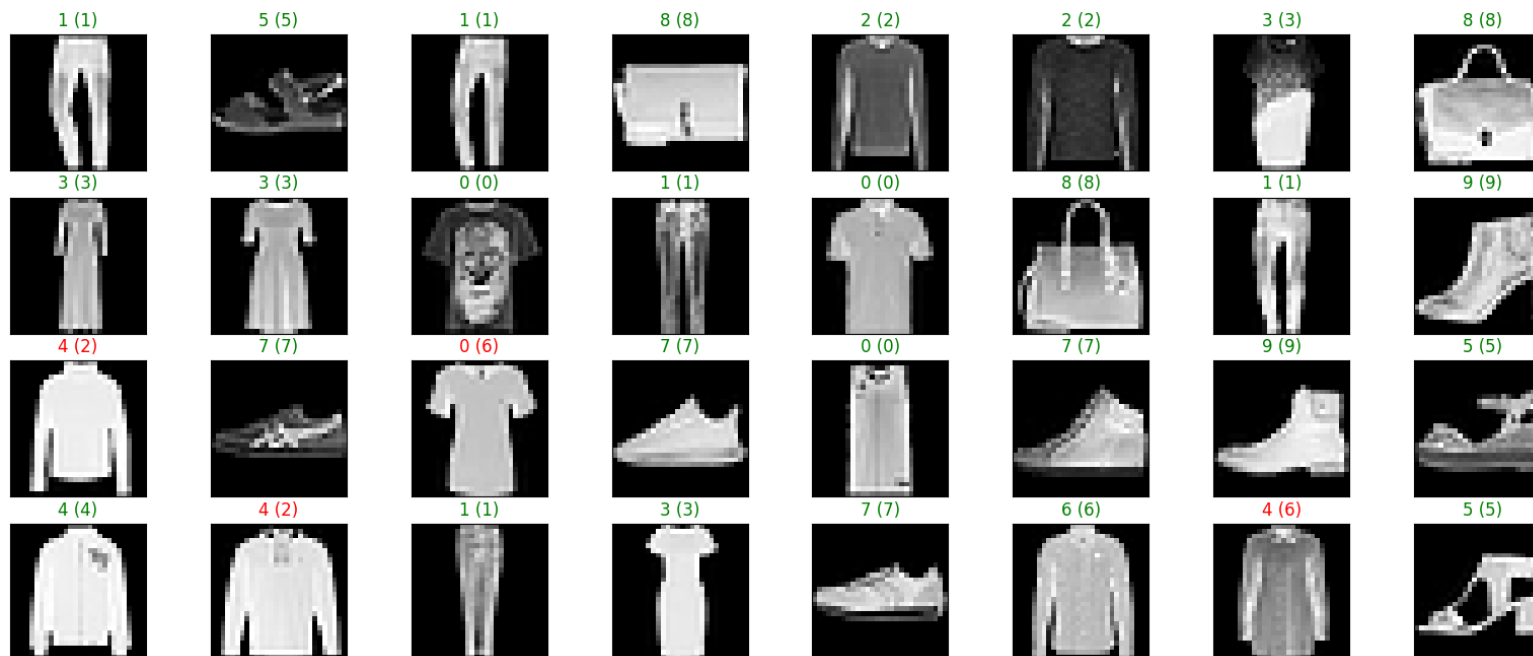
        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                     dataset_labels[true_index],
                                     color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 1s 1ms/step

Test wyniki:



✓ 1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.


```

for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))

visualize_filters(filters, 'Filtr')

```



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.

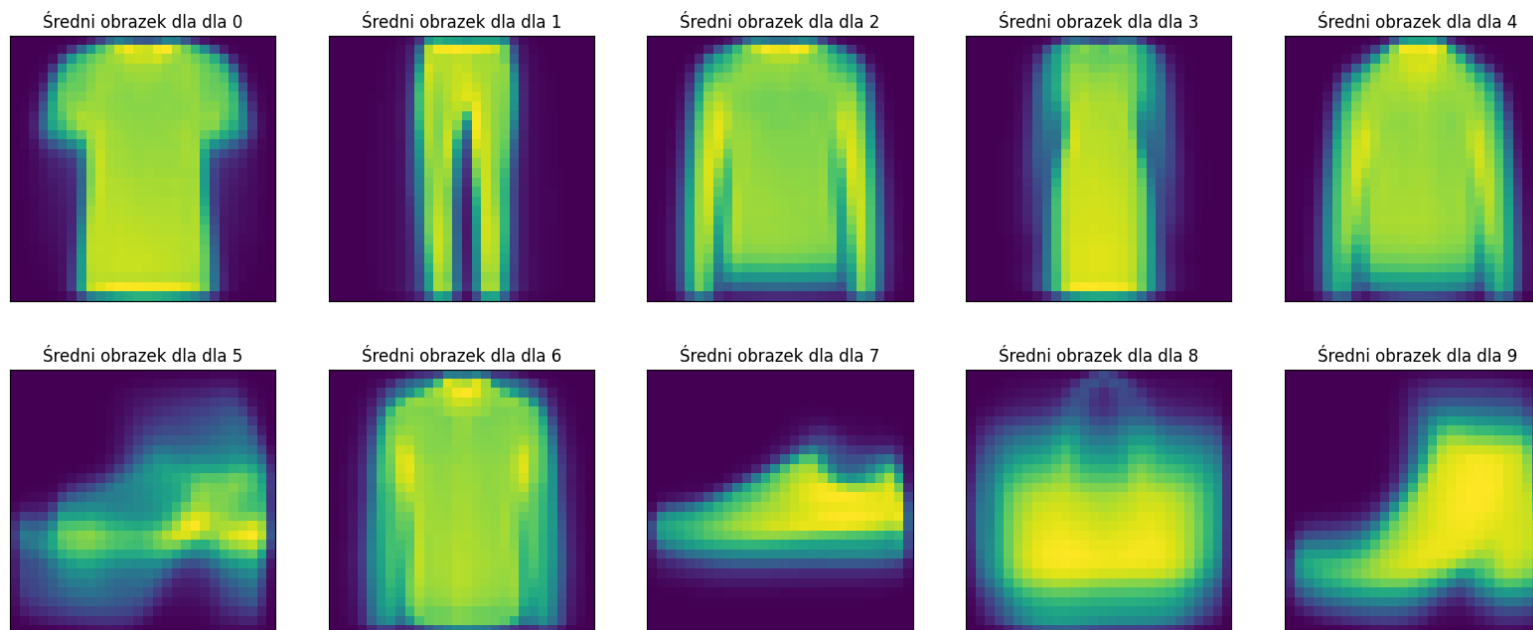
```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:, :, :, label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:, :, :, i] = avg_images[:, :, :, i] / class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.0 Prosty model liniowy v2

Zacniemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przeksztalci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 10)	7850
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optimizera adam**
- Jako loss function użyjemy '**categorical_crossentropy**'
- Lista parametrów, tutaj zaczniemy od '**precyzji**'

Warto zerknąć: <https://keras.io/models/model/>

```
opt = keras.optimizers.Adam(learning_rate=0.002)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???
```

✓ Uczenie modelu

Model uczy my wykorzystując fit().

```
history = model.fit(x_train, y_train, batch_size = 32, epochs = 5, validation_data = (x_valid,y_valid))
```



```
Epoch 1/5
1688/1688 [=====] - 6s 3ms/step - loss: 0.5608 - accuracy: 0.8092 - val_loss: 0.4680 - val_accuracy: 0.8358
Epoch 2/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.4523 - accuracy: 0.8437 - val_loss: 0.4662 - val_accuracy: 0.8410
```

```
Epoch 3/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.4311 - accuracy: 0.8516 - val_loss: 0.4286 - val_accuracy: 0.8487
Epoch 4/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.4212 - accuracy: 0.8538 - val_loss: 0.4370 - val_accuracy: 0.8490
Epoch 5/5
1688/1688 [=====] - 5s 3ms/step - loss: 0.4146 - accuracy: 0.8562 - val_loss: 0.4180 - val_accuracy: 0.8540
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

Wczytanie modelu

```
#from keras.models import load_model  
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)  
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.8410999774932861
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'

```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])

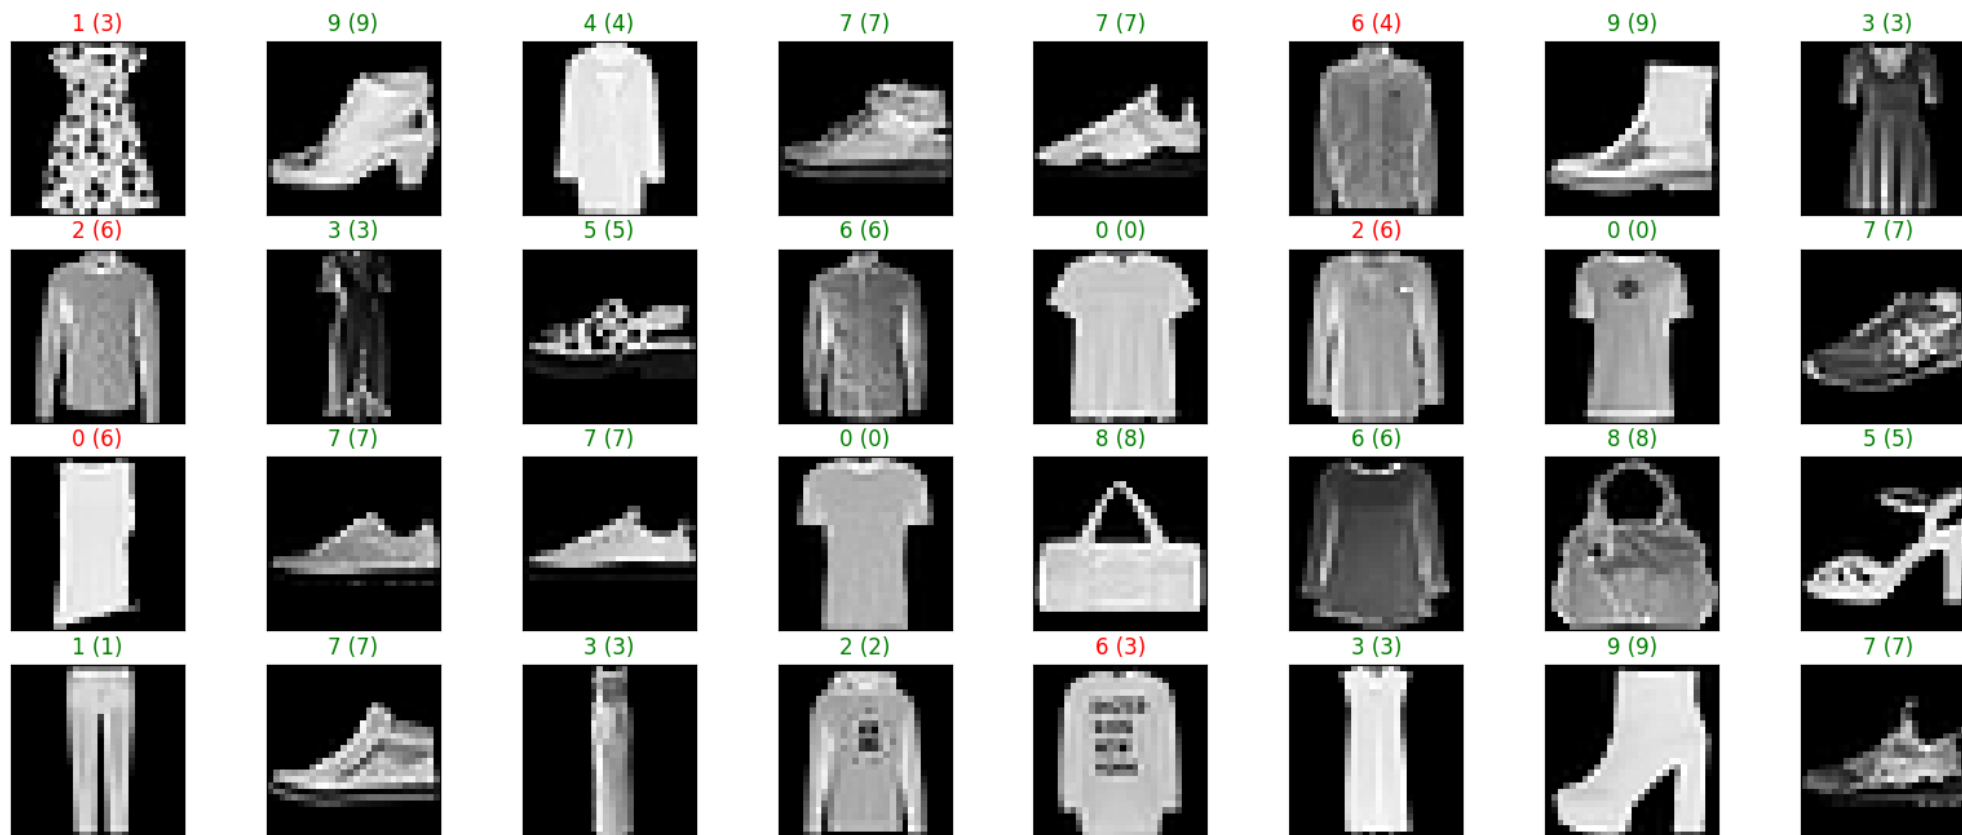
        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                   dataset_labels[true_index],
                                   color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 0s 1ms/step

Test wyniki:



✓ 1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.

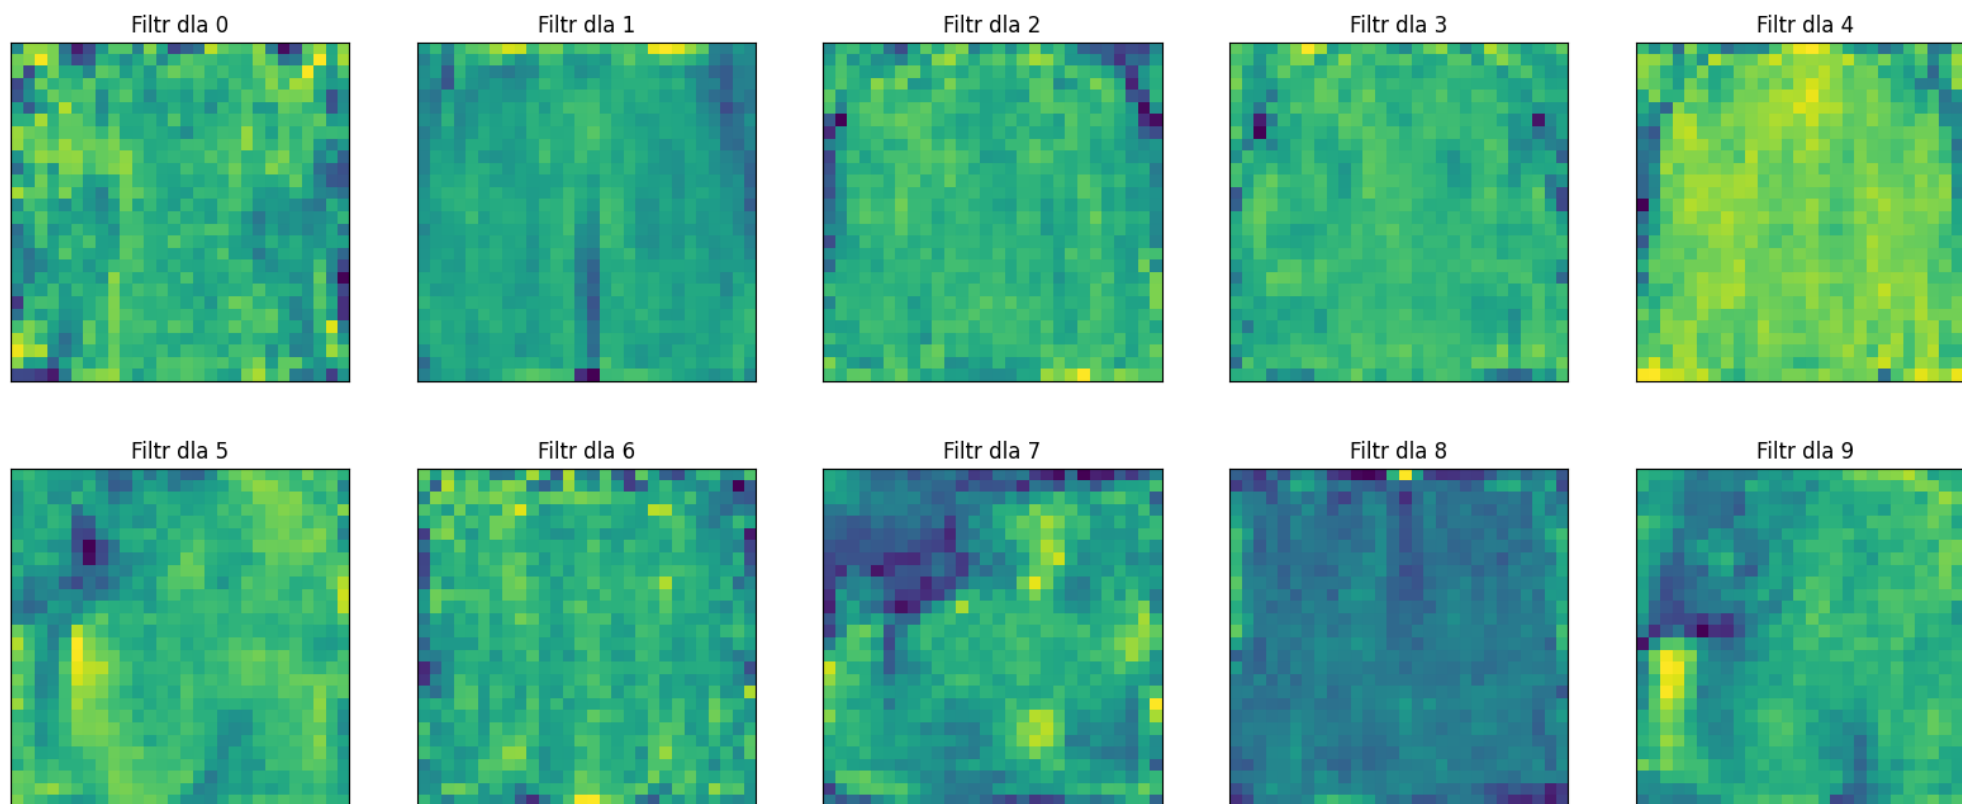
```

for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))

visualize_filters(filters, 'Filtr')

```



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.


```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:, :, :, label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:, :, :, i] = avg_images[:, :, :, i] / class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.0 Prosty model liniowy v3

Zacznijemy od prostego modelu składającego się z jednej transformacji liniowej.

- Model stworzymy za pomocą `tf.keras.Sequential()` (https://www.tensorflow.org/api_docs/python/tf/keras/models/Sequential). Ponieważ nie zastosujemy jeszcze konwolucji zatem możemy spłaszczyć obrazki do wektorów zawierających 28x28 wartości.
- Następnie dodamy jedną warstwę liniową, która przeksztalci wejściowe piksele w 10 klas. Ponieważ wyniki reprezentują prawdopodobieństwa możemy użyć funkcji aktywacji softmax.
- Szczegóły modelu uzyskamy z pomocą `model.summary()`

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(10,activation="softmax"))
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 10)	7850
Total params: 7850 (30.66 KB)		
Trainable params: 7850 (30.66 KB)		
Non-trainable params: 0 (0.00 Byte)		

✓ Kompilacja modelu

Uwagi:

- Użyjemy **optymizera sgd**
- Jako loss function użyjemy '**categorical_crossentropy**'
- Lista parametrów, tutaj zaczniemy od '**precyzji**'

Warto zerknąć: <https://keras.io/models/model/>

```
opt = keras.optimizers.SGD(learning_rate=0.04)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

✓ Uczenie modelu

Model uczymy wykorzystując fit().



```
history = model.fit(x_train, y_train, batch_size = 64, epochs = 75, validation_data = (x_valid,y_valid))
```


844/844 [=====] - 35 imgs/step - loss: 0.3781 - accuracy: 0.8098 - val_loss: 0.4073 - val_accuracy: 0.8508

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



✓ Zapisanie i wczytanie modelu

Zapisanie modelu

```
model.save("mnist_simple.h5")
```

Wczytanie modelu

```
#from keras.models import load_model  
#model = load_model("mnist_simple.h5")
```

✓ Precyzja

Wykorzystamy funkcję evaluate()

```
score = model.evaluate(x_test,y_test,verbose=0)  
print('Test accuracy:',score[1])
```

```
Test accuracy: 0.8434000015258789
```

✓ Przewidywania modelu

Przetestujmy przewidywania naszego modelu. Sprawdzimy go na danych testowych. W tym celu wykorzystamy poniższą funkcję 'visualize_model_predictions(model, x, y)'

```
def visualize_model_predictions(model, x_test, y_test, title_string):
    y_hat = model.predict(x_test)

    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=32, replace=False)):
        ax = figure.add_subplot(4, 8, i + 1, xticks=[], yticks=[])

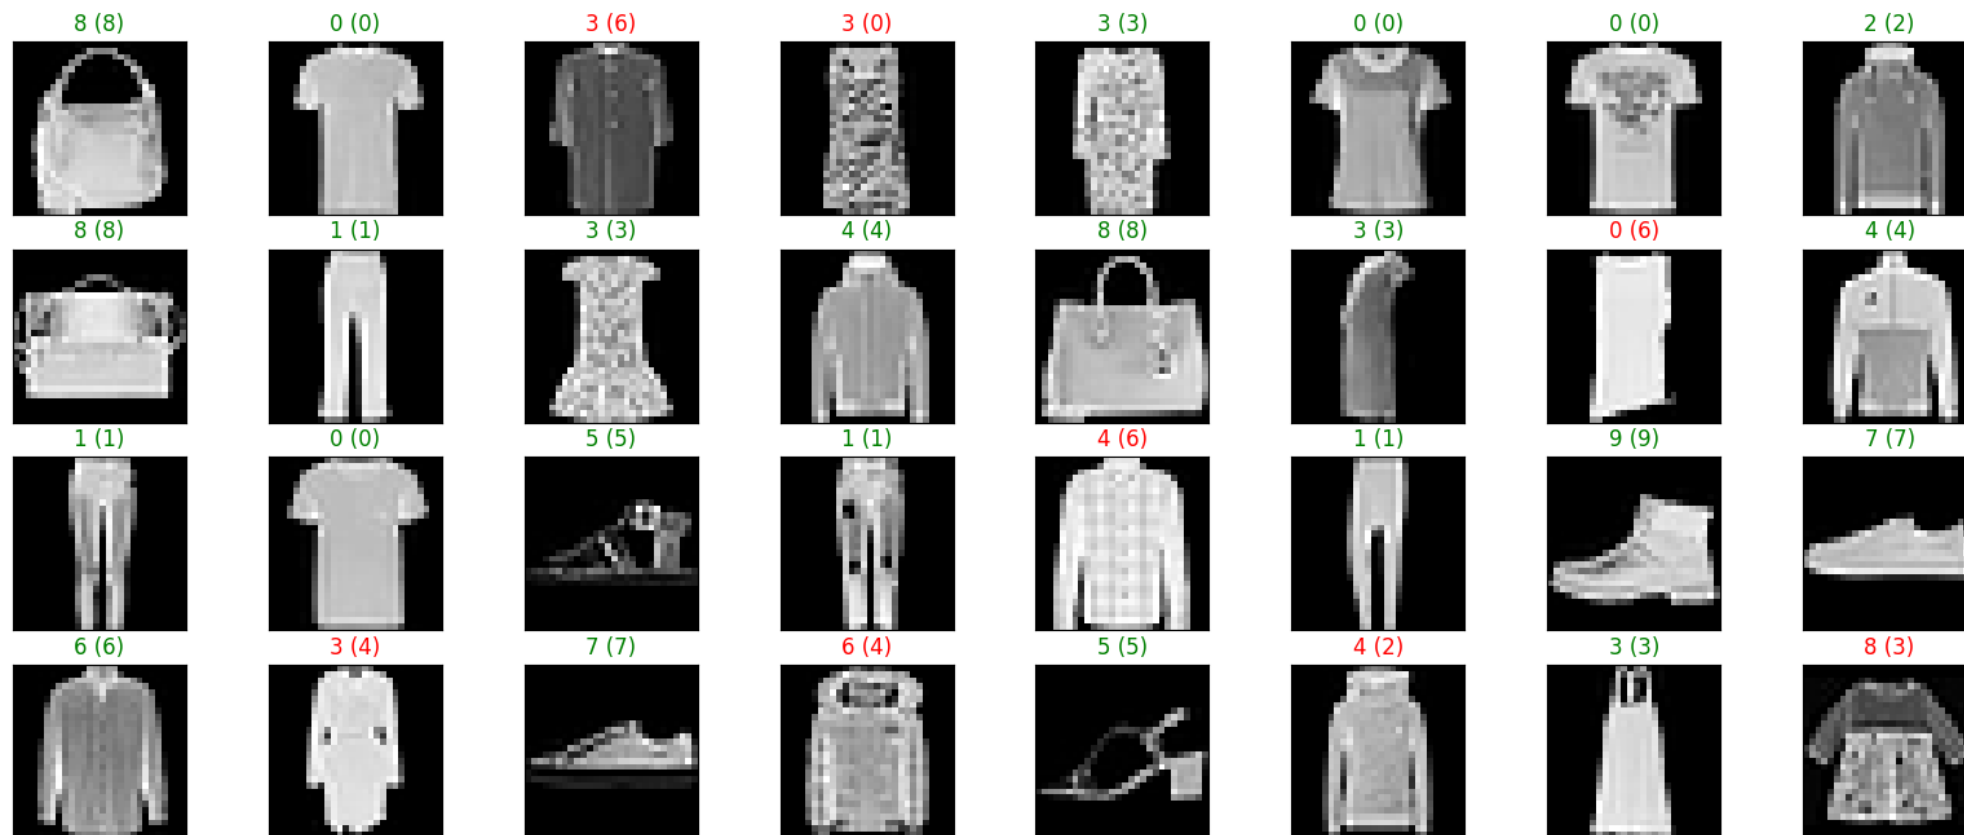
        ax.imshow(np.squeeze(x_test[index]), cmap = 'gray')
        predict_index = np.argmax(y_hat[index])
        true_index = np.argmax(y_test[index])

        ax.set_title("{} ({}).format(dataset_labels[predict_index],
                                   dataset_labels[true_index],
                                   color=("green" if predict_index == true_index else "red"))
    figure.suptitle("%s wyniki:" %title_string, fontsize=25)

visualize_model_predictions(model, x_test, y_test, 'Test')
```

313/313 [=====] - 0s 1ms/step

Test wyniki:



✓ 1. Wizualizacja wag dla każdej klasy

Warstwę transformacyjną naszego modelu można przedstawić za pomocą macierzy wag [28x28, 10]. Spróbujmy narysować każdy z 10 filtrów. W celu uzyskania wag modelu wykorzystamy funkcje `model.layers` i `get_weights()`.


```

for layer in model.layers:
    weights = layer.get_weights()
    if len(weights) > 0:
        w,b = weights
        filters = np.reshape(w, (28,28,10))

def visualize_filters(filters, title_string):
    figure = plt.figure(figsize=(20, 8))
    for i, index in enumerate(np.random.choice(x_test.shape[0], size=10, replace=False)):
        ax = figure.add_subplot(2, 5, i + 1, xticks=[], yticks=[])
        ax.imshow(filters[:, :, i], cmap = 'viridis')
        ax.set_title("%s dla %s" %(title_string, dataset_labels[i]))

visualize_filters(filters, 'Filtr')

```



✓ 2 I jeszcze jedna wizualizacja

Porównajmy powyższe filtry, ze średnim zdjęciem dla każdej klasy.

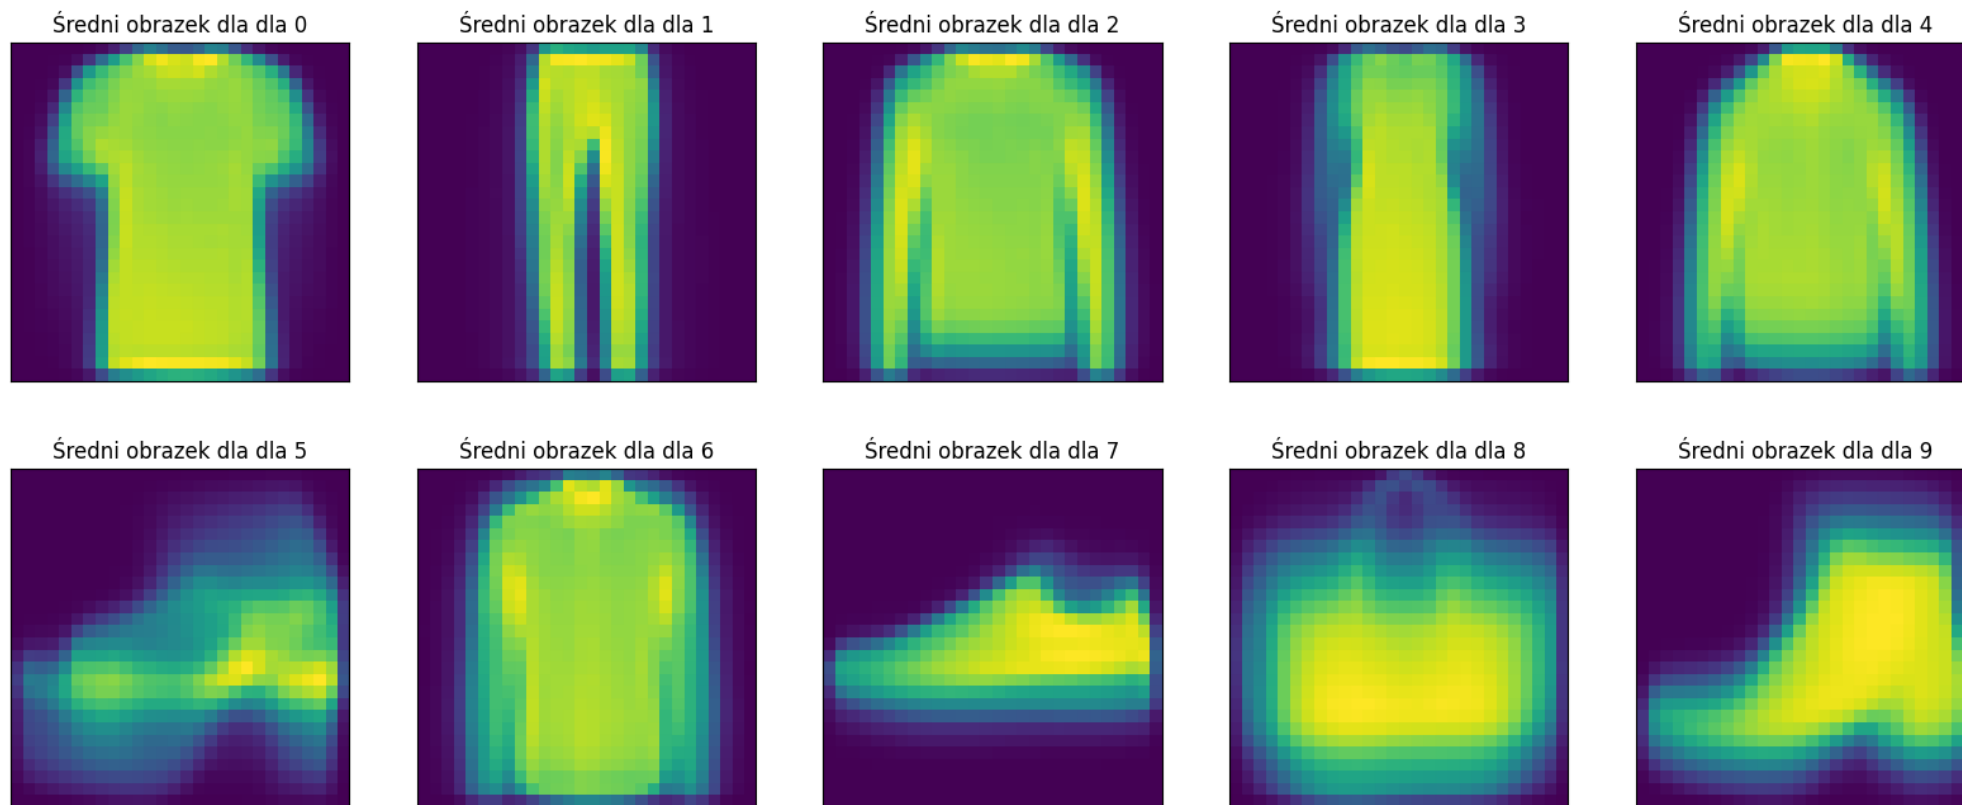
```
avg_images = np.zeros((28,28,1,10))
class_images = [0]*10

for i in range(len(x_train)):
    img = x_train[i]
    label = np.argmax(y_train[i])

    avg_images[:, :, :, label] += img
    class_images[label] += 1

for i in range(10):
    avg_images[:, :, :, i] = avg_images[:, :, :, i]/class_images[i]

avg_images = np.squeeze(avg_images)
visualize_filters(avg_images, 'Średni obrazek dla')
```



✓ 2.1 A teraz sieć neuronowa

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(60,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()

model.compile(optimizer="adam", loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))
```


Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 60)	47100
dense_4 (Dense)	(None, 10)	610

```

=====
Total params: 47710 (186.37 KB)
Trainable params: 47710 (186.37 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```

```

Epoch 1/10
844/844 [=====] - 4s 3ms/step - loss: 0.5653 - accuracy: 0.8077 - val_loss: 0.4265 - val_accuracy: 0.8442
Epoch 2/10
844/844 [=====] - 4s 4ms/step - loss: 0.4121 - accuracy: 0.8532 - val_loss: 0.4095 - val_accuracy: 0.8583
Epoch 3/10
844/844 [=====] - 3s 3ms/step - loss: 0.3755 - accuracy: 0.8665 - val_loss: 0.3804 - val_accuracy: 0.8623
Epoch 4/10
844/844 [=====] - 3s 3ms/step - loss: 0.3510 - accuracy: 0.8751 - val_loss: 0.3873 - val_accuracy: 0.8627
Epoch 5/10
844/844 [=====] - 3s 3ms/step - loss: 0.3331 - accuracy: 0.8804 - val_loss: 0.3767 - val_accuracy: 0.8678
Epoch 6/10
844/844 [=====] - 3s 4ms/step - loss: 0.3205 - accuracy: 0.8843 - val_loss: 0.3475 - val_accuracy: 0.8717
Epoch 7/10
844/844 [=====] - 3s 4ms/step - loss: 0.3064 - accuracy: 0.8886 - val_loss: 0.3397 - val_accuracy: 0.8783
Epoch 8/10
844/844 [=====] - 3s 3ms/step - loss: 0.2953 - accuracy: 0.8921 - val_loss: 0.3408 - val_accuracy: 0.8763
Epoch 9/10
844/844 [=====] - 3s 3ms/step - loss: 0.2856 - accuracy: 0.8970 - val_loss: 0.3514 - val_accuracy: 0.8767
Epoch 10/10
844/844 [=====] - 3s 3ms/step - loss: 0.2788 - accuracy: 0.8984 - val_loss: 0.3277 - val_accuracy: 0.8805

```

Precyzja

```



score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Precyzja: 0.8726999759674072

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



- Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

313/313 [=====] - 1s 2ms/step

test wyniki:



✓ 2.1 Sieć neuronowa v2

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(64,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
opt = keras.optimizers.Adam(learning_rate=0.004)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy']) #learnig rate???

history = model.fit(x_train, y_train, batch_size = 64, epochs = 13, validation_data = (x_valid,y_valid))
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
flatten_4 (Flatten)	(None, 784)	0
dense_5 (Dense)	(None, 64)	50240
dense_6 (Dense)	(None, 10)	650

=====

Total params: 50890 (198.79 KB)
 Trainable params: 50890 (198.79 KB)
 Non-trainable params: 0 (0.00 Byte)

Epoch 1/13
 844/844 [=====] - 4s 4ms/step - loss: 0.4945 - accuracy: 0.8241 - val_loss: 0.3975 - val_accuracy: 0.8548
 Epoch 2/13
 844/844 [=====] - 3s 3ms/step - loss: 0.3838 - accuracy: 0.8597 - val_loss: 0.3851 - val_accuracy: 0.8600
 Epoch 3/13
 844/844 [=====] - 4s 4ms/step - loss: 0.3518 - accuracy: 0.8710 - val_loss: 0.3633 - val_accuracy: 0.8670
 Epoch 4/13
 844/844 [=====] - 3s 3ms/step - loss: 0.3352 - accuracy: 0.8764 - val_loss: 0.3787 - val_accuracy: 0.8648
 Epoch 5/13
 844/844 [=====] - 3s 3ms/step - loss: 0.3216 - accuracy: 0.8816 - val_loss: 0.3957 - val_accuracy: 0.8547
 Epoch 6/13
 844/844 [=====] - 3s 3ms/step - loss: 0.3127 - accuracy: 0.8836 - val_loss: 0.3613 - val_accuracy: 0.8735
 Epoch 7/13
 844/844 [=====] - 3s 3ms/step - loss: 0.2995 - accuracy: 0.8890 - val_loss: 0.3703 - val_accuracy: 0.8655
 Epoch 8/13
 844/844 [=====] - 3s 4ms/step - loss: 0.2933 - accuracy: 0.8914 - val_loss: 0.3580 - val_accuracy: 0.8722
 Epoch 9/13
 844/844 [=====] - 3s 3ms/step - loss: 0.2811 - accuracy: 0.8966 - val_loss: 0.3680 - val_accuracy: 0.8728

```

Epoch 10/13
844/844 [=====] - 3s 3ms/step - loss: 0.2784 - accuracy: 0.8960 - val_loss: 0.3735 - val_accuracy: 0.8690
Epoch 11/13
844/844 [=====] - 3s 3ms/step - loss: 0.2719 - accuracy: 0.8977 - val_loss: 0.3533 - val_accuracy: 0.8778
Epoch 12/13
844/844 [=====] - 4s 4ms/step - loss: 0.2643 - accuracy: 0.9018 - val_loss: 0.3645 - val_accuracy: 0.8797
Epoch 13/13
844/844 [=====] - 3s 3ms/step - loss: 0.2613 - accuracy: 0.9029 - val_loss: 0.3778 - val_accuracy: 0.8732

```

Precyzja

```

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

```
Precyzja: 0.8708999752998352
```


Wykresy precyzji i błędu

```


import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()

```

Training and validation accuracy



Training and validation loss



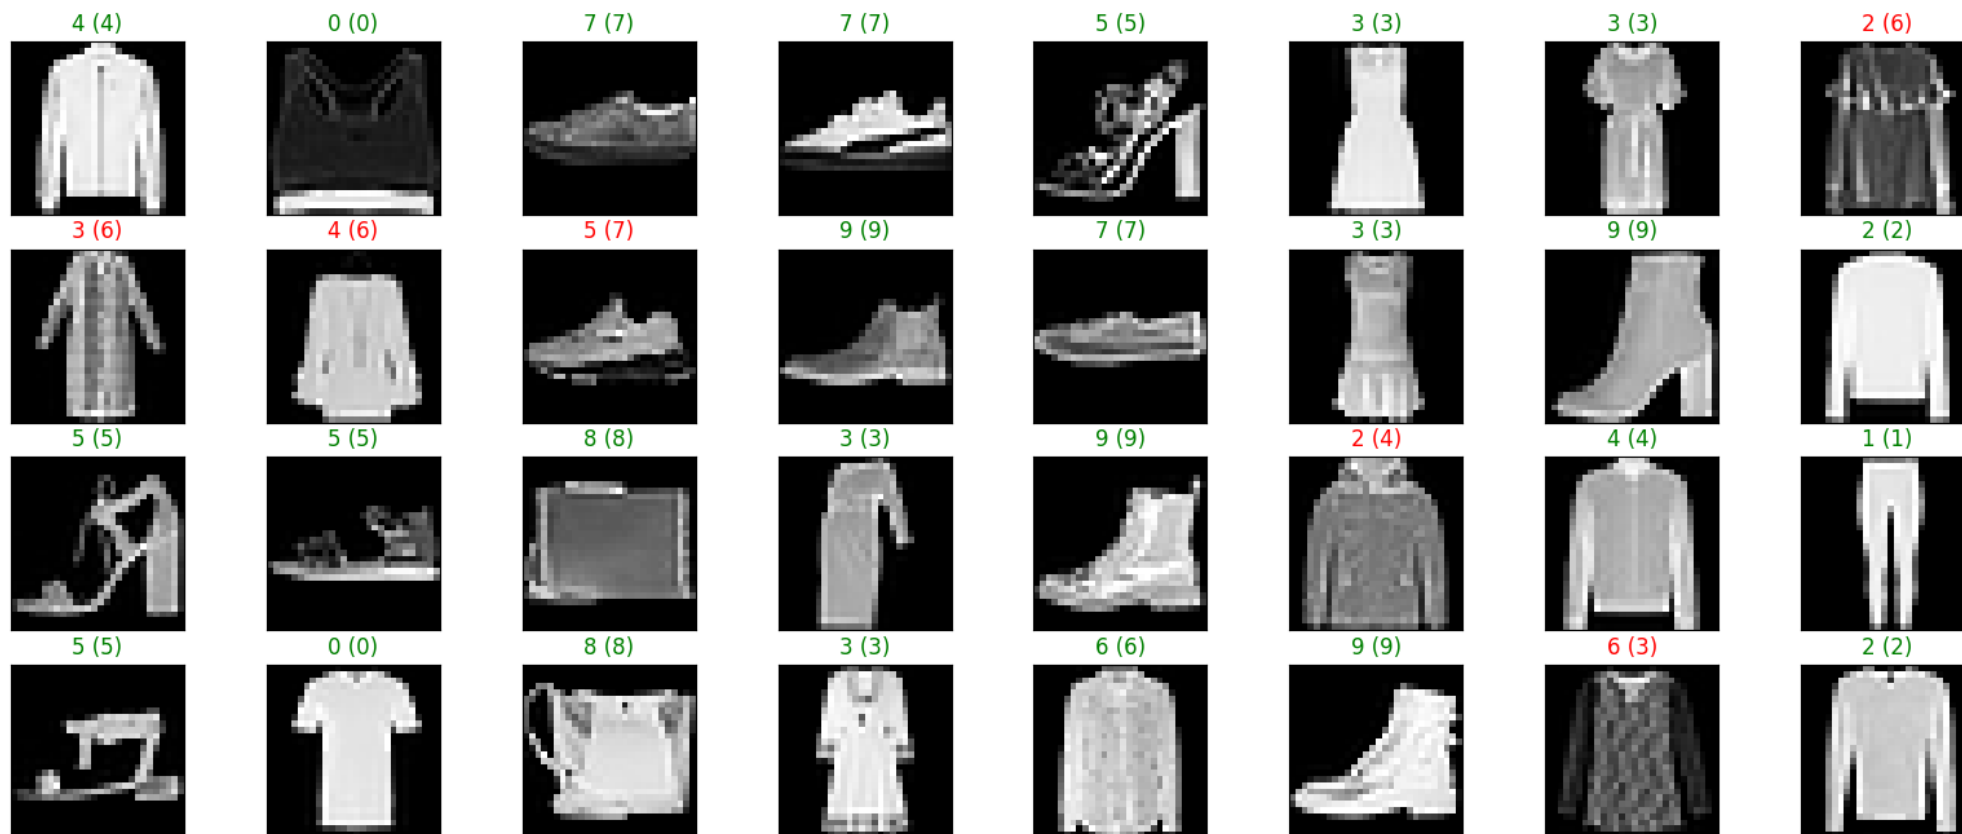
▼ Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 1ms/step
```

test wyniki:



✓ 2.1 Sieć neuronowa v3

Dodajmy teraz warstwy wewnętrzne w naszej sieci. Funkcja aktywacji w takich warstwach to zwykle relu.

```
model = tf.keras.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(64,activation="relu"))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
opt = keras.optimizers.SGD(learning_rate=0.08)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 128, epochs = 50, validation_data = (x_valid,y_valid))
```

```

422/422 [=====] - 1s 4ms/step - loss: 0.2412 - accuracy: 0.9128 - val_loss: 0.3240 - val_accuracy: 0.8840
Epoch 39/50
422/422 [=====] - 1s 3ms/step - loss: 0.2390 - accuracy: 0.9137 - val_loss: 0.3658 - val_accuracy: 0.8697
Epoch 40/50
422/422 [=====] - 1s 3ms/step - loss: 0.2376 - accuracy: 0.9130 - val_loss: 0.3291 - val_accuracy: 0.8818
Epoch 41/50
422/422 [=====] - 1s 3ms/step - loss: 0.2337 - accuracy: 0.9150 - val_loss: 0.3156 - val_accuracy: 0.8872
Epoch 42/50
422/422 [=====] - 1s 3ms/step - loss: 0.2321 - accuracy: 0.9164 - val_loss: 0.3318 - val_accuracy: 0.8818
Epoch 43/50
422/422 [=====] - 1s 3ms/step - loss: 0.2320 - accuracy: 0.9160 - val_loss: 0.3172 - val_accuracy: 0.8838
Epoch 44/50
422/422 [=====] - 1s 3ms/step - loss: 0.2293 - accuracy: 0.9178 - val_loss: 0.3276 - val_accuracy: 0.8843
Epoch 45/50
422/422 [=====] - 2s 4ms/step - loss: 0.2277 - accuracy: 0.9174 - val_loss: 0.3395 - val_accuracy: 0.8842
Epoch 46/50
422/422 [=====] - 2s 5ms/step - loss: 0.2249 - accuracy: 0.9197 - val_loss: 0.3272 - val_accuracy: 0.8855
Epoch 47/50
422/422 [=====] - 1s 3ms/step - loss: 0.2251 - accuracy: 0.9189 - val_loss: 0.3270 - val_accuracy: 0.8860
Epoch 48/50
422/422 [=====] - 1s 3ms/step - loss: 0.2217 - accuracy: 0.9207 - val_loss: 0.3282 - val_accuracy: 0.8855
Epoch 49/50
422/422 [=====] - 1s 3ms/step - loss: 0.2199 - accuracy: 0.9207 - val_loss: 0.3722 - val_accuracy: 0.8703
Epoch 50/50
422/422 [=====] - 1s 3ms/step - loss: 0.2191 - accuracy: 0.9205 - val_loss: 0.3217 - val_accuracy: 0.8875

```

Precyzja

```



score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

```
Precyzja: 0.8774999976158142
```

Wykresy precyzji i błędu

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



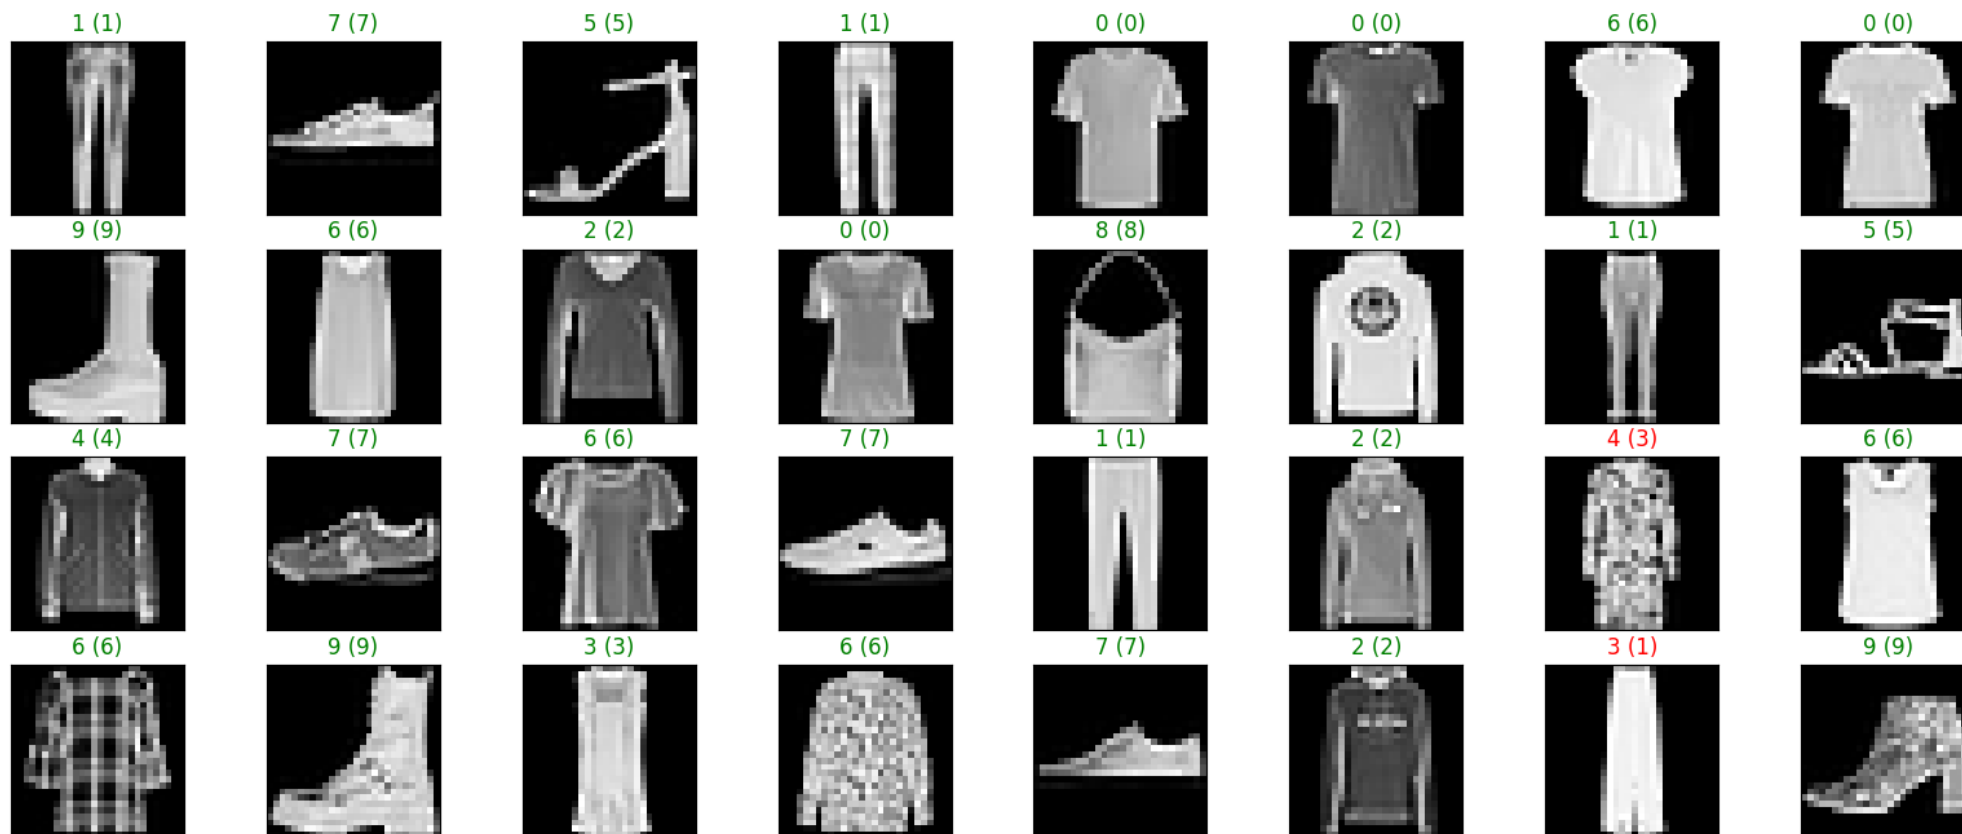
✓ Przewidywania modelu

Sprawdźmy jakie są przewidywania naszego modelu

```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.2 Spróbujmy pogłębić nasz model!

Dodajmy 3 warstwy gęste.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
flatten_6 (Flatten)	(None, 784)	0
dense_9 (Dense)	(None, 128)	100480
dense_10 (Dense)	(None, 64)	8256
dense_11 (Dense)	(None, 64)	4160
dense_12 (Dense)	(None, 32)	2080
dense_13 (Dense)	(None, 10)	330

```

=====
Total params: 115306 (450.41 KB)
Trainable params: 115306 (450.41 KB)
Non-trainable params: 0 (0.00 Byte)

```

```

Epoch 1/10
844/844 [=====] - 6s 5ms/step - loss: 0.5682 - accuracy: 0.7956 - val_loss: 0.4015 - val_accuracy: 0.8535
Epoch 2/10
844/844 [=====] - 3s 4ms/step - loss: 0.3907 - accuracy: 0.8580 - val_loss: 0.3665 - val_accuracy: 0.8657
Epoch 3/10
844/844 [=====] - 3s 4ms/step - loss: 0.3477 - accuracy: 0.8731 - val_loss: 0.3860 - val_accuracy: 0.8558
Epoch 4/10
844/844 [=====] - 4s 5ms/step - loss: 0.3252 - accuracy: 0.8798 - val_loss: 0.3581 - val_accuracy: 0.8722

```

```

Epoch 5/10
844/844 [=====] - 3s 4ms/step - loss: 0.3068 - accuracy: 0.8869 - val_loss: 0.3473 - val_accuracy: 0.8673
Epoch 6/10
844/844 [=====] - 3s 4ms/step - loss: 0.2932 - accuracy: 0.8919 - val_loss: 0.3538 - val_accuracy: 0.8710
Epoch 7/10
844/844 [=====] - 4s 5ms/step - loss: 0.2794 - accuracy: 0.8959 - val_loss: 0.3819 - val_accuracy: 0.8737
Epoch 8/10
844/844 [=====] - 4s 4ms/step - loss: 0.2680 - accuracy: 0.9000 - val_loss: 0.3414 - val_accuracy: 0.8782
Epoch 9/10
844/844 [=====] - 3s 4ms/step - loss: 0.2589 - accuracy: 0.9034 - val_loss: 0.3230 - val_accuracy: 0.8850
Epoch 10/10
844/844 [=====] - 3s 4ms/step - loss: 0.2510 - accuracy: 0.9039 - val_loss: 0.3588 - val_accuracy: 0.8740
Precyzja: 0.8686000108718872



```

Wykresy precyzji i błędu

```

import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()

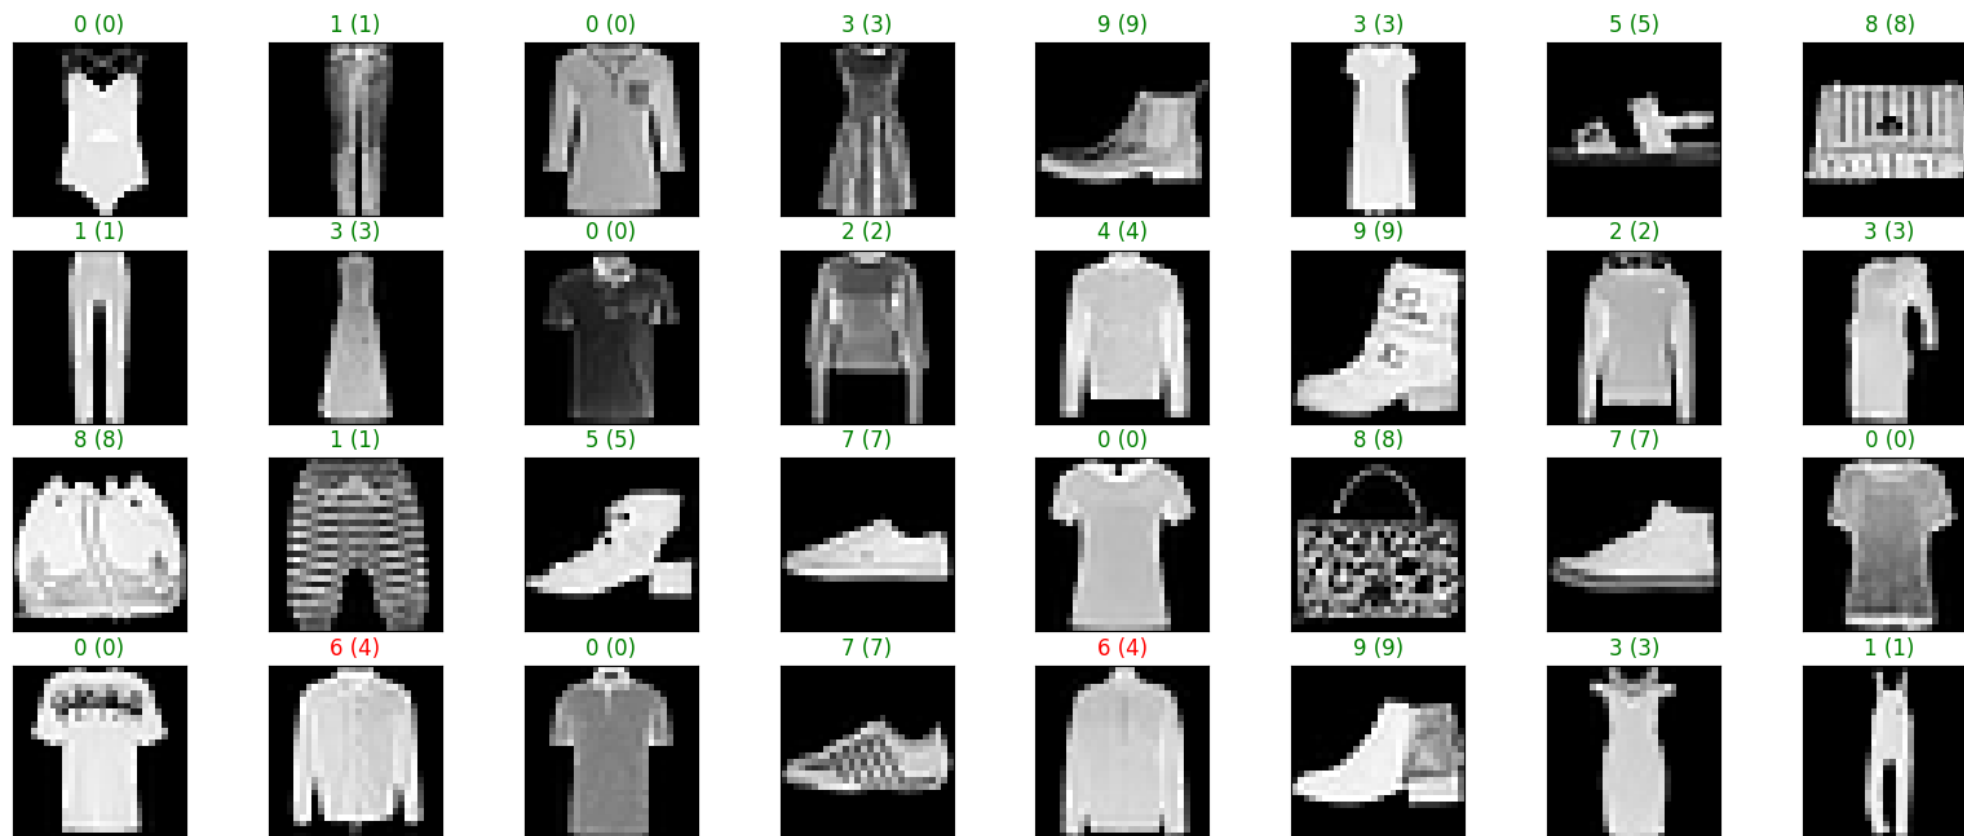
```




```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.2 Pogłębinienie modelu v2

Dodajmy 2 warstwy gęste.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
flatten_7 (Flatten)	(None, 784)	0
dense_14 (Dense)	(None, 128)	100480
dense_15 (Dense)	(None, 64)	8256
dense_16 (Dense)	(None, 32)	2080
dense_17 (Dense)	(None, 10)	330

Total params: 111146 (434.16 KB)
 Trainable params: 111146 (434.16 KB)
 Non-trainable params: 0 (0.00 Byte)

```



Epoch 1/10
844/844 [=====] - 5s 4ms/step - loss: 0.5522 - accuracy: 0.8073 - val_loss: 0.4940 - val_accuracy: 0.8123
Epoch 2/10
844/844 [=====] - 4s 5ms/step - loss: 0.3888 - accuracy: 0.8599 - val_loss: 0.3804 - val_accuracy: 0.8632
Epoch 3/10
844/844 [=====] - 3s 4ms/step - loss: 0.3449 - accuracy: 0.8739 - val_loss: 0.3514 - val_accuracy: 0.8690
Epoch 4/10
844/844 [=====] - 3s 4ms/step - loss: 0.3200 - accuracy: 0.8824 - val_loss: 0.3328 - val_accuracy: 0.8738
Epoch 5/10
844/844 [=====] - 4s 4ms/step - loss: 0.3033 - accuracy: 0.8869 - val_loss: 0.3304 - val_accuracy: 0.8800
Epoch 6/10
844/844 [=====] - 4s 4ms/step - loss: 0.2872 - accuracy: 0.8934 - val_loss: 0.3324 - val_accuracy: 0.8820
Epoch 7/10

```

```
844/844 [=====] - 3s 4ms/step - loss: 0.2787 - accuracy: 0.8966 - val_loss: 0.3252 - val_accuracy: 0.8810
Epoch 8/10
844/844 [=====] - 3s 4ms/step - loss: 0.2666 - accuracy: 0.9005 - val_loss: 0.3142 - val_accuracy: 0.8813
Epoch 9/10
844/844 [=====] - 4s 5ms/step - loss: 0.2534 - accuracy: 0.9055 - val_loss: 0.3081 - val_accuracy: 0.8860
Epoch 10/10
844/844 [=====] - 4s 5ms/step - loss: 0.2471 - accuracy: 0.9078 - val_loss: 0.3248 - val_accuracy: 0.8822
Precyzja: 0.8751999735832214
```

Wykresy precyzji i błędu

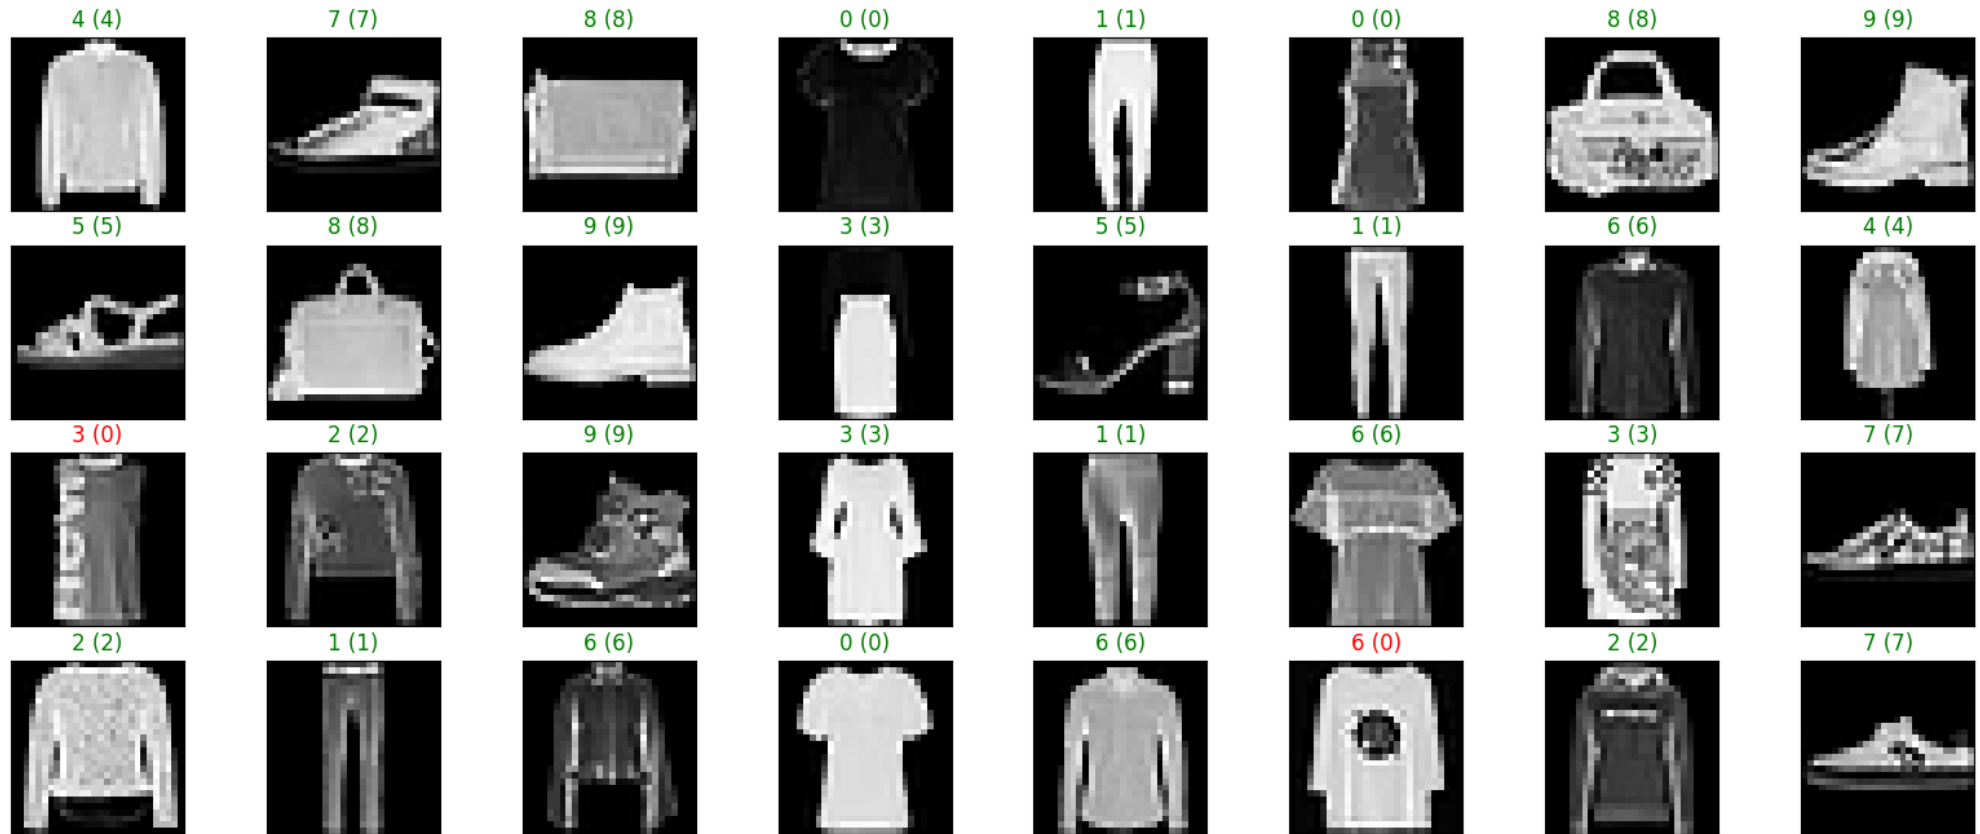
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 2s 4ms/step
```

test wyniki:



✓ 2.2 Pogłębiamy model v3

Dodajmy 4 warstwy gęste.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Flatten(input_shape=(28,28,1)))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(x_train, y_train, batch_size = 64, epochs = 10, validation_data = (x_valid,y_valid))

score = model.evaluate(x_test,y_test,verbose=0)
print('Precyzja: ',score[1])

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		
flatten_8 (Flatten)	(None, 784)	0
dense_18 (Dense)	(None, 128)	100480
dense_19 (Dense)	(None, 128)	16512
dense_20 (Dense)	(None, 64)	8256
dense_21 (Dense)	(None, 64)	4160
dense_22 (Dense)	(None, 32)	2080
dense_23 (Dense)	(None, 10)	330

=====

Total params: 131818 (514.91 KB)
Trainable params: 131818 (514.91 KB)
Non-trainable params: 0 (0.00 Byte)

Epoch 1/10
844/844 [=====] - 5s 4ms/step - loss: 0.5538 - accuracy: 0.7996 - val_loss: 0.4585 - val_accuracy: 0.8338
Epoch 2/10
844/844 [=====] - 4s 5ms/step - loss: 0.3875 - accuracy: 0.8581 - val_loss: 0.3851 - val_accuracy: 0.8600
Epoch 3/10
844/844 [=====] - 3s 4ms/step - loss: 0.3475 - accuracy: 0.8713 - val_loss: 0.3718 - val_accuracy: 0.8613
Epoch 4/10

```

844/844 [=====] - 3s 4ms/step - loss: 0.3237 - accuracy: 0.8804 - val_loss: 0.3308 - val_accuracy: 0.8780
Epoch 5/10
844/844 [=====] - 4s 5ms/step - loss: 0.3020 - accuracy: 0.8889 - val_loss: 0.3603 - val_accuracy: 0.8690
Epoch 6/10
844/844 [=====] - 4s 4ms/step - loss: 0.2914 - accuracy: 0.8919 - val_loss: 0.3694 - val_accuracy: 0.8612
Epoch 7/10
844/844 [=====] - 4s 4ms/step - loss: 0.2771 - accuracy: 0.8968 - val_loss: 0.3299 - val_accuracy: 0.8808
Epoch 8/10
844/844 [=====] - 4s 4ms/step - loss: 0.2637 - accuracy: 0.9004 - val_loss: 0.3125 - val_accuracy: 0.8850
Epoch 9/10
844/844 [=====] - 4s 5ms/step - loss: 0.2561 - accuracy: 0.9039 - val_loss: 0.3477 - val_accuracy: 0.8812
Epoch 10/10
844/844 [=====] - 3s 4ms/step - loss: 0.2487 - accuracy: 0.9070 - val_loss: 0.3360 - val_accuracy: 0.8828
Precyzja: 0.8799999952316284

```


Wykresy precyzji i błędu

```


import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()

```

Training and validation accuracy



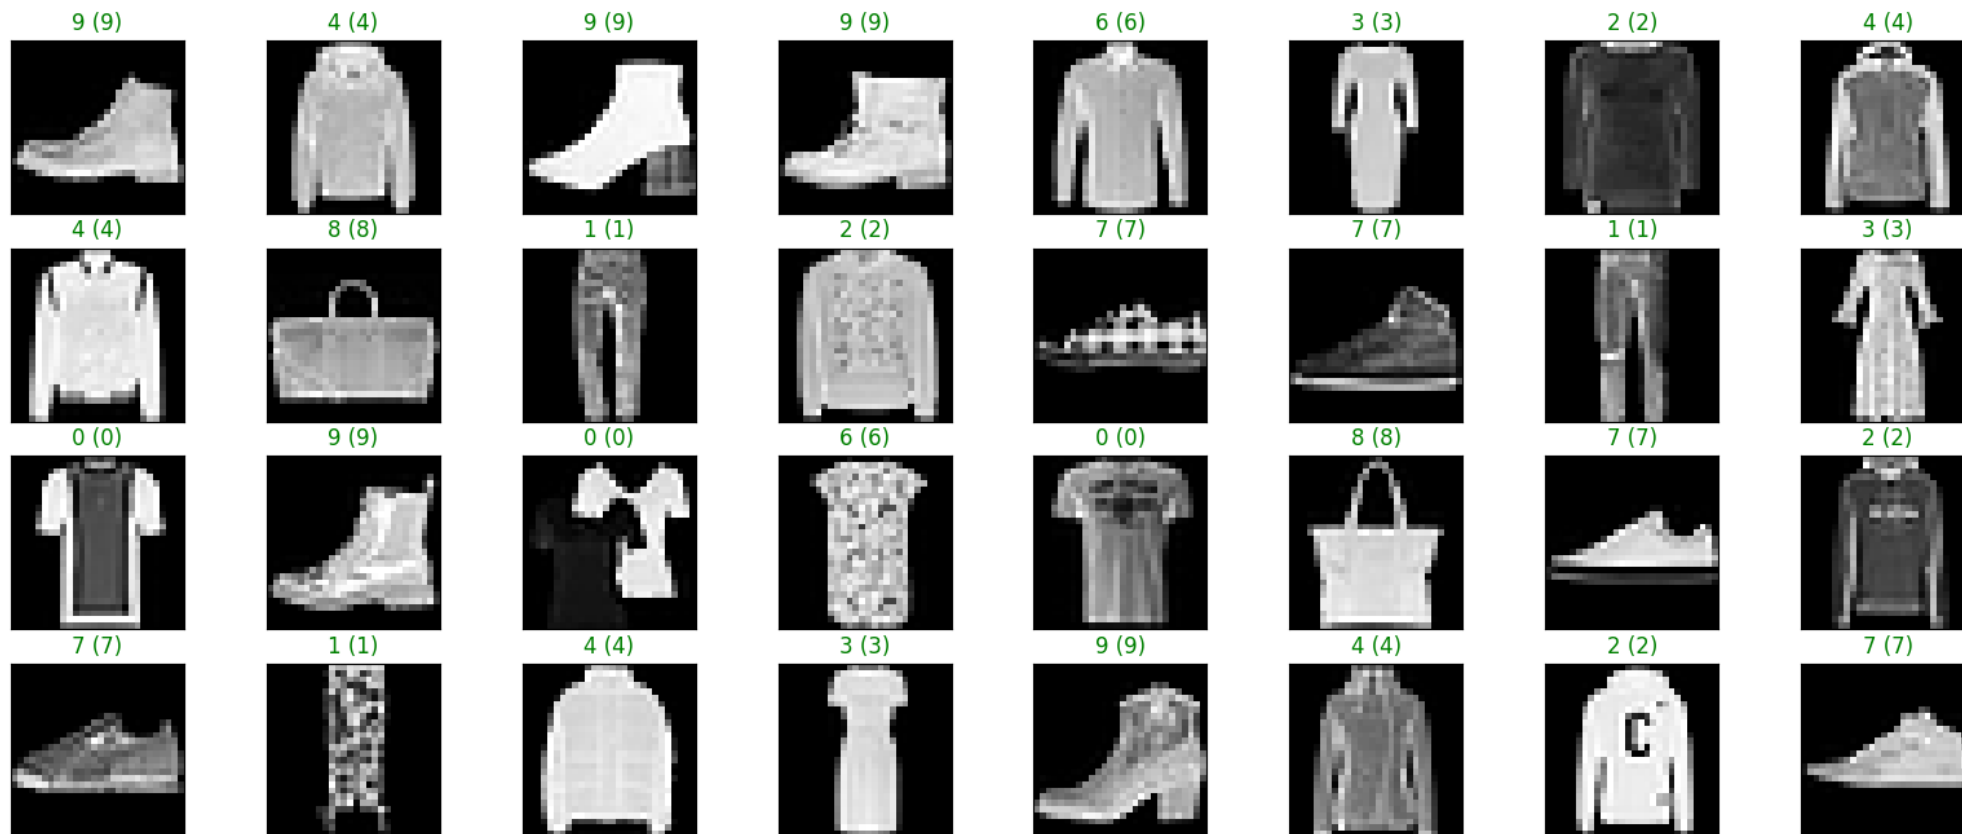
Training and validation loss




```
visualize_model_predictions(model, x_test, y_test, "test" )
```

```
313/313 [=====] - 1s 2ms/step
```

test wyniki:



✓ 2.3 Konwolucja

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=4, kernel_size=2,padding='same',activation='relu',input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2,padding='same',activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 4)	20
max_pooling2d (MaxPooling2D)	(None, 14, 14, 4)	0
conv2d_1 (Conv2D)	(None, 14, 14, 2)	34
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 2)	0
flatten_9 (Flatten)	(None, 98)	0
dense_24 (Dense)	(None, 128)	12672
dense_25 (Dense)	(None, 64)	8256
dense_26 (Dense)	(None, 10)	650
Total params: 21632 (84.50 KB)		
Trainable params: 21632 (84.50 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```

Epoch 1/25
422/422 [=====] - 6s 5ms/step - loss: 0.8460 - accuracy: 0.6909 - val_loss: 0.5659 - val_accuracy: 0.7833
Epoch 2/25

```

```
422/422 [=====] - 2s 4ms/step - loss: 0.5310 - accuracy: 0.8006 - val_loss: 0.4907 - val_accuracy: 0.8155
Epoch 3/25
422/422 [=====] - 2s 4ms/step - loss: 0.4643 - accuracy: 0.8292 - val_loss: 0.4464 - val_accuracy: 0.8360
Epoch 4/25
422/422 [=====] - 2s 4ms/step - loss: 0.4304 - accuracy: 0.8422 - val_loss: 0.4248 - val_accuracy: 0.8415
Epoch 5/25
422/422 [=====] - 2s 4ms/step - loss: 0.4082 - accuracy: 0.8496 - val_loss: 0.4117 - val_accuracy: 0.8502
Epoch 6/25
422/422 [=====] - 2s 5ms/step - loss: 0.3921 - accuracy: 0.8572 - val_loss: 0.3825 - val_accuracy: 0.8593
Epoch 7/25
422/422 [=====] - 2s 6ms/step - loss: 0.3769 - accuracy: 0.8616 - val_loss: 0.3723 - val_accuracy: 0.8607
Epoch 8/25
422/422 [=====] - 2s 4ms/step - loss: 0.3629 - accuracy: 0.8662 - val_loss: 0.3712 - val_accuracy: 0.8628
Epoch 9/25
422/422 [=====] - 2s 4ms/step - loss: 0.3506 - accuracy: 0.8703 - val_loss: 0.3608 - val_accuracy: 0.8635
Epoch 10/25
422/422 [=====] - 2s 4ms/step - loss: 0.3402 - accuracy: 0.8727 - val_loss: 0.3554 - val_accuracy: 0.8702
Epoch 11/25
422/422 [=====] - 2s 4ms/step - loss: 0.3330 - accuracy: 0.8768 - val_loss: 0.3616 - val_accuracy: 0.8672
Epoch 12/25
422/422 [=====] - 2s 4ms/step - loss: 0.3254 - accuracy: 0.8794 - val_loss: 0.3630 - val_accuracy: 0.8658
Epoch 13/25
422/422 [=====] - 3s 7ms/step - loss: 0.3177 - accuracy: 0.8810 - val_loss: 0.3492 - val_accuracy: 0.8723
Epoch 14/25
422/422 [=====] - 3s 6ms/step - loss: 0.3125 - accuracy: 0.8829 - val_loss: 0.3430 - val_accuracy: 0.8733
Epoch 15/25
422/422 [=====] - 2s 4ms/step - loss: 0.3064 - accuracy: 0.8849 - val_loss: 0.3284 - val_accuracy: 0.8758
Epoch 16/25
422/422 [=====] - 2s 4ms/step - loss: 0.2992 - accuracy: 0.8893 - val_loss: 0.3327 - val_accuracy: 0.8757
Epoch 17/25
422/422 [=====] - 2s 4ms/step - loss: 0.2923 - accuracy: 0.8909 - val_loss: 0.3256 - val_accuracy: 0.8778
Epoch 18/25
422/422 [=====] - 2s 4ms/step - loss: 0.2900 - accuracy: 0.8908 - val_loss: 0.3242 - val_accuracy: 0.8797
Epoch 19/25
422/422 [=====] - 2s 5ms/step - loss: 0.2844 - accuracy: 0.8936 - val_loss: 0.3248 - val_accuracy: 0.8775
Epoch 20/25
422/422 [=====] - 2s 5ms/step - loss: 0.2782 - accuracy: 0.8963 - val_loss: 0.3130 - val_accuracy: 0.8817
Epoch 21/25
422/422 [=====] - 2s 4ms/step - loss: 0.2759 - accuracy: 0.8971 - val_loss: 0.3142 - val_accuracy: 0.8853
Epoch 22/25
422/422 [=====] - 2s 4ms/step - loss: 0.2712 - accuracy: 0.8978 - val_loss: 0.3178 - val_accuracy: 0.8813
Epoch 23/25
422/422 [=====] - 2s 4ms/step - loss: 0.2672 - accuracy: 0.8993 - val_loss: 0.3177 - val_accuracy: 0.8837
Epoch 24/25
422/422 [=====] - 2s 4ms/step - loss: 0.2667 - accuracy: 0.8993 - val_loss: 0.3158 - val_accuracy: 0.8863
Epoch 25/25
422/422 [=====] - 2s 4ms/step - loss: 0.2621 - accuracy: 0.9018 - val_loss: 0.3237 - val_accuracy: 0.8765
```



✓ Ewaluacja modelu

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ',score[1])
```

Precyzja: 0.8812000155448914

Wykresy precyzji i błędu

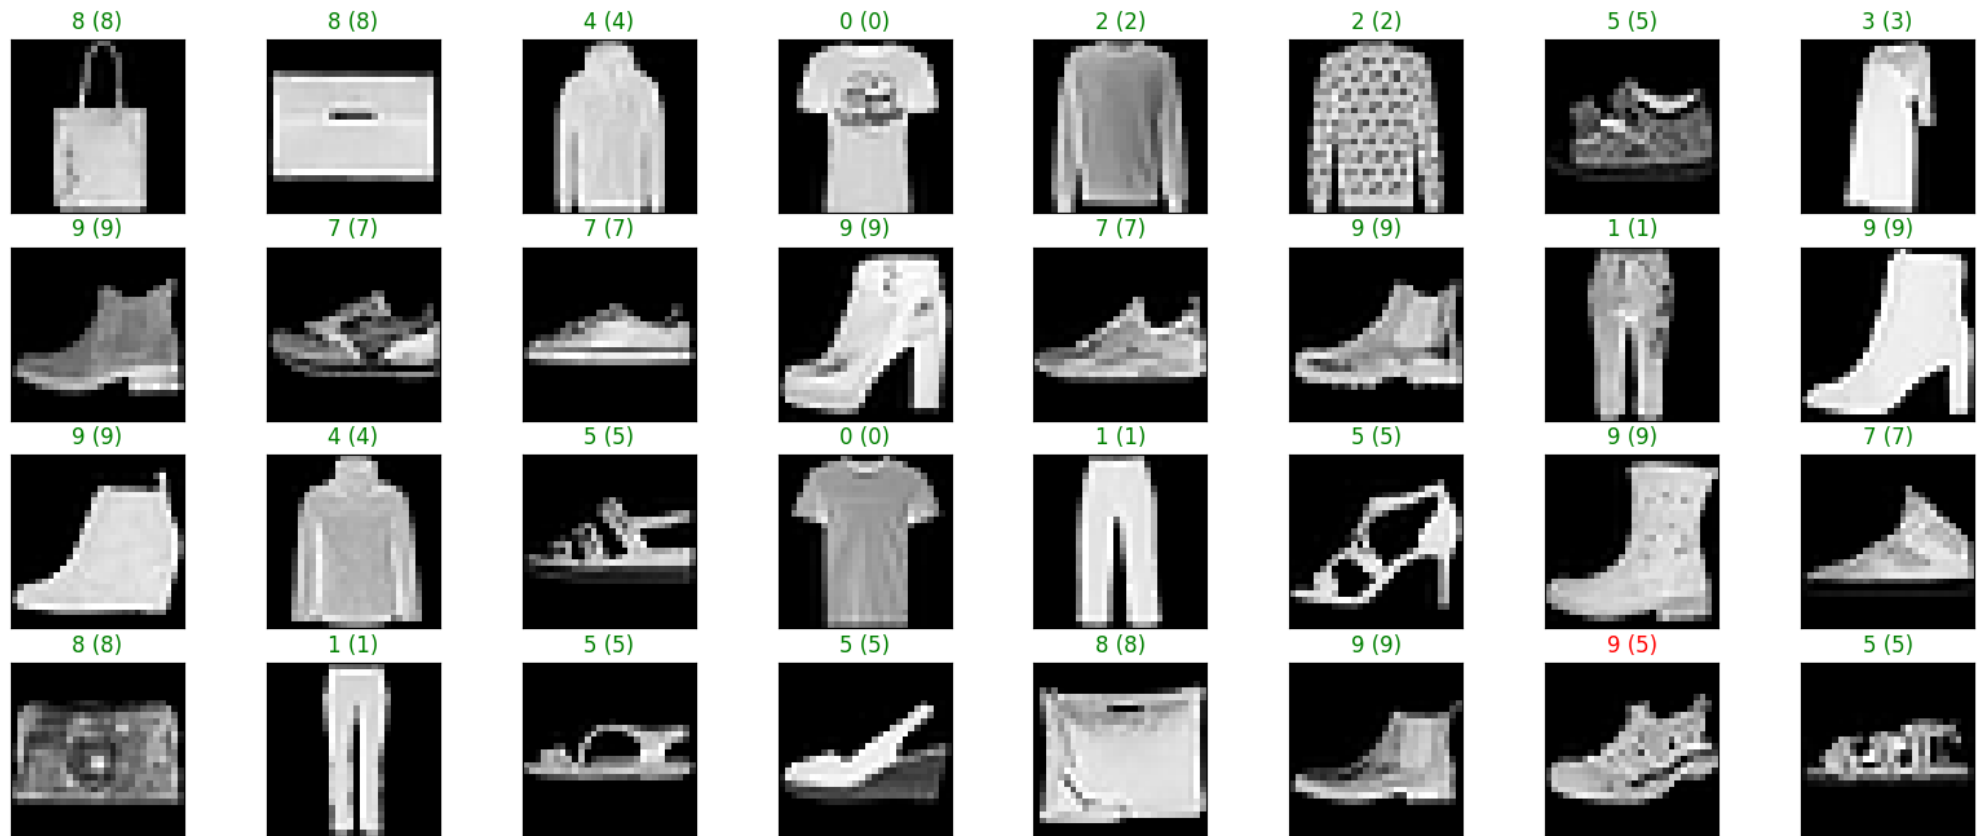
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

313/313 [=====] - 1s 3ms/step

convnet wyniki:



✓ 2.3 Konwolucja model v2

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=4, kernel_size=2, padding='same', activation='relu', input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=2, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 28, 28, 4)	20
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 4)	0
conv2d_3 (Conv2D)	(None, 14, 14, 2)	34
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 2)	0
conv2d_4 (Conv2D)	(None, 7, 7, 2)	18
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 2)	0
flatten_10 (Flatten)	(None, 18)	0
dense_27 (Dense)	(None, 128)	2432
dense_28 (Dense)	(None, 64)	8256
dense_29 (Dense)	(None, 10)	650

```
=====
Total params: 11410 (44.57 KB)
Trainable params: 11410 (44.57 KB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 4s 5ms/step - loss: 1.1637 - accuracy: 0.5790 - val_loss: 0.7959 - val_accuracy: 0.6807
Epoch 2/25
422/422 [=====] - 2s 5ms/step - loss: 0.7650 - accuracy: 0.7042 - val_loss: 0.6997 - val_accuracy: 0.7293
Epoch 3/25
422/422 [=====] - 2s 5ms/step - loss: 0.6898 - accuracy: 0.7345 - val_loss: 0.6478 - val_accuracy: 0.7452
Epoch 4/25
422/422 [=====] - 3s 6ms/step - loss: 0.6401 - accuracy: 0.7561 - val_loss: 0.5995 - val_accuracy: 0.7725
Epoch 5/25
422/422 [=====] - 2s 5ms/step - loss: 0.6076 - accuracy: 0.7702 - val_loss: 0.5994 - val_accuracy: 0.7697
Epoch 6/25
422/422 [=====] - 2s 5ms/step - loss: 0.5841 - accuracy: 0.7803 - val_loss: 0.5755 - val_accuracy: 0.7847
Epoch 7/25
422/422 [=====] - 2s 5ms/step - loss: 0.5668 - accuracy: 0.7889 - val_loss: 0.5513 - val_accuracy: 0.7888
Epoch 8/25
422/422 [=====] - 2s 5ms/step - loss: 0.5510 - accuracy: 0.7934 - val_loss: 0.5310 - val_accuracy: 0.8002
Epoch 9/25
422/422 [=====] - 2s 5ms/step - loss: 0.5404 - accuracy: 0.7985 - val_loss: 0.5261 - val_accuracy: 0.7992
Epoch 10/25
422/422 [=====] - 2s 6ms/step - loss: 0.5260 - accuracy: 0.8053 - val_loss: 0.5091 - val_accuracy: 0.8112
Epoch 11/25
422/422 [=====] - 2s 5ms/step - loss: 0.5172 - accuracy: 0.8079 - val_loss: 0.5029 - val_accuracy: 0.8143
Epoch 12/25
422/422 [=====] - 2s 5ms/step - loss: 0.5076 - accuracy: 0.8123 - val_loss: 0.5075 - val_accuracy: 0.8095
Epoch 13/25
422/422 [=====] - 2s 5ms/step - loss: 0.4981 - accuracy: 0.8141 - val_loss: 0.5007 - val_accuracy: 0.8112
Epoch 14/25
422/422 [=====] - 2s 5ms/step - loss: 0.4893 - accuracy: 0.8189 - val_loss: 0.4824 - val_accuracy: 0.8185
Epoch 15/25
422/422 [=====] - 2s 5ms/step - loss: 0.4806 - accuracy: 0.8219 - val_loss: 0.4747 - val_accuracy: 0.8210
Epoch 16/25
422/422 [=====] - 3s 6ms/step - loss: 0.4753 - accuracy: 0.8247 - val_loss: 0.4730 - val_accuracy: 0.8222
Epoch 17/25
422/422 [=====] - 2s 5ms/step - loss: 0.4711 - accuracy: 0.8226 - val_loss: 0.4662 - val_accuracy: 0.8267
Epoch 18/25
422/422 [=====] - 2s 5ms/step - loss: 0.4608 - accuracy: 0.8279 - val_loss: 0.4752 - val_accuracy: 0.8212
Epoch 19/25
422/422 [=====] - 2s 5ms/step - loss: 0.4599 - accuracy: 0.8277 - val_loss: 0.4673 - val_accuracy: 0.8235
Epoch 20/25
422/422 [=====] - 2s 5ms/step - loss: 0.4538 - accuracy: 0.8313 - val_loss: 0.4557 - val_accuracy: 0.8322
Epoch 21/25
422/422 [=====] - 2s 5ms/step - loss: 0.4494 - accuracy: 0.8326 - val_loss: 0.4492 - val_accuracy: 0.8372
```



```
Epoch 22/25
422/422 [=====] - 3s 6ms/step - loss: 0.4456 - accuracy: 0.8330 - val_loss: 0.4507 - val_accuracy: 0.8318
Epoch 23/25
422/422 [=====] - 2s 5ms/step - loss: 0.4417 - accuracy: 0.8341 - val_loss: 0.4515 - val_accuracy: 0.8302
Epoch 24/25
422/422 [=====] - 2s 5ms/step - loss: 0.4376 - accuracy: 0.8367 - val_loss: 0.4476 - val_accuracy: 0.8372
Epoch 25/25
422/422 [=====] - 2s 5ms/step - loss: 0.4344 - accuracy: 0.8360 - val_loss: 0.4490 - val_accuracy: 0.8320
```



✓ Ewaluacja modelu

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ', score[1])
```

Precyzja: 0.8307999968528748

Wykresy precyzji i błędu

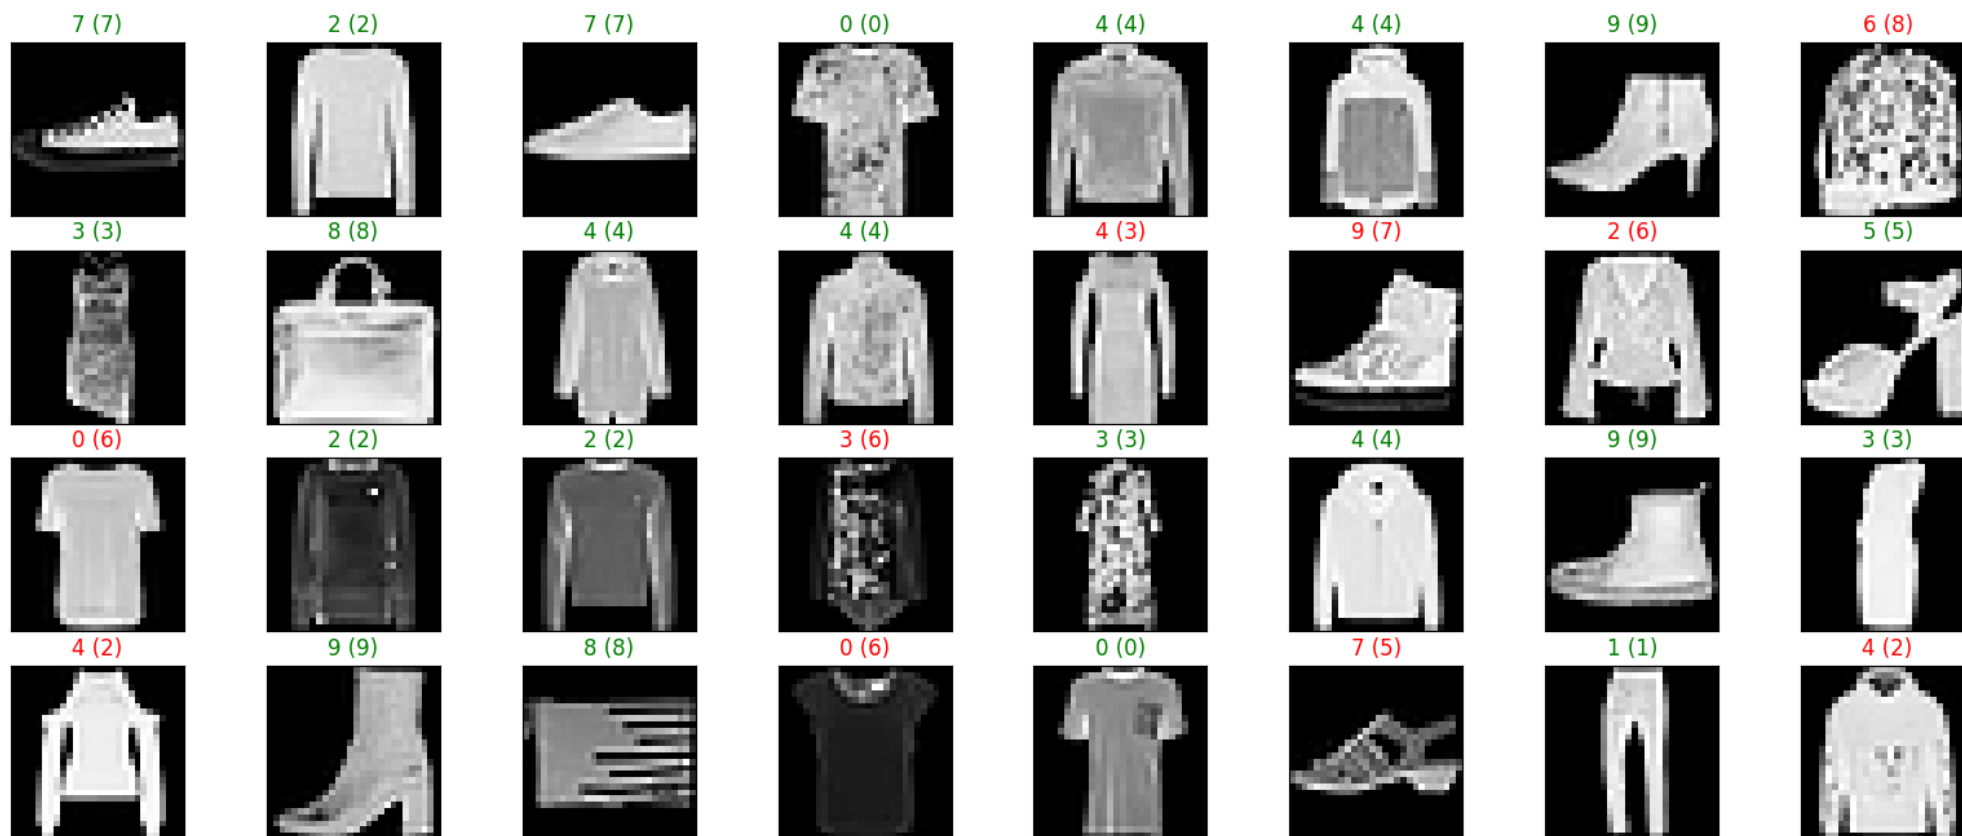
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

```
313/313 [=====] - 1s 2ms/step
```

convnet wyniki:



2.3 Konwolucja model v3

W dotychczasowych przykładach przed warstwą gęstą dodawaliśmy warstwę płaską. Tutaj będzie podobnie, ale przed warstwą płaską dodamy warstwy konwolucyjne i maxpool.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2, padding='same', activation='relu', input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2, padding='same', activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_6 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_6 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_7 (Conv2D)	(None, 7, 7, 16)	2064
max_pooling2d_7 (MaxPooling2D)	(None, 3, 3, 16)	0
flatten_11 (Flatten)	(None, 144)	0
dense_30 (Dense)	(None, 128)	18560
dense_31 (Dense)	(None, 64)	8256
dense_32 (Dense)	(None, 10)	650

=====

Total params: 38074 (148.73 KB)
 Trainable params: 38074 (148.73 KB)

Non-trainable params: 0 (0.00 Byte)

```
history = model.fit(x_train, y_train, batch_size=128, epochs=25, validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 5s 7ms/step - loss: 0.7195 - accuracy: 0.7354 - val_loss: 0.4585 - val_accuracy: 0.8332
Epoch 2/25
422/422 [=====] - 3s 7ms/step - loss: 0.4166 - accuracy: 0.8487 - val_loss: 0.4034 - val_accuracy: 0.8548
Epoch 3/25
422/422 [=====] - 3s 6ms/step - loss: 0.3568 - accuracy: 0.8701 - val_loss: 0.3831 - val_accuracy: 0.8572
Epoch 4/25
422/422 [=====] - 2s 6ms/step - loss: 0.3261 - accuracy: 0.8821 - val_loss: 0.3415 - val_accuracy: 0.8787
Epoch 5/25
422/422 [=====] - 2s 6ms/step - loss: 0.3085 - accuracy: 0.8870 - val_loss: 0.3166 - val_accuracy: 0.8847
Epoch 6/25
422/422 [=====] - 2s 6ms/step - loss: 0.2871 - accuracy: 0.8949 - val_loss: 0.2970 - val_accuracy: 0.8910
Epoch 7/25
422/422 [=====] - 3s 7ms/step - loss: 0.2742 - accuracy: 0.8987 - val_loss: 0.2798 - val_accuracy: 0.9002
Epoch 8/25
422/422 [=====] - 3s 6ms/step - loss: 0.2630 - accuracy: 0.9030 - val_loss: 0.2810 - val_accuracy: 0.8962
Epoch 9/25
422/422 [=====] - 2s 6ms/step - loss: 0.2528 - accuracy: 0.9063 - val_loss: 0.2711 - val_accuracy: 0.8985
Epoch 10/25
422/422 [=====] - 2s 6ms/step - loss: 0.2439 - accuracy: 0.9103 - val_loss: 0.2829 - val_accuracy: 0.8985
Epoch 11/25
422/422 [=====] - 2s 6ms/step - loss: 0.2354 - accuracy: 0.9130 - val_loss: 0.2660 - val_accuracy: 0.9013
Epoch 12/25
422/422 [=====] - 3s 7ms/step - loss: 0.2300 - accuracy: 0.9159 - val_loss: 0.2536 - val_accuracy: 0.9092
Epoch 13/25
422/422 [=====] - 3s 6ms/step - loss: 0.2216 - accuracy: 0.9183 - val_loss: 0.2626 - val_accuracy: 0.9035
Epoch 14/25
422/422 [=====] - 2s 6ms/step - loss: 0.2163 - accuracy: 0.9197 - val_loss: 0.2623 - val_accuracy: 0.9083
Epoch 15/25
422/422 [=====] - 2s 6ms/step - loss: 0.2102 - accuracy: 0.9226 - val_loss: 0.2535 - val_accuracy: 0.9065
Epoch 16/25
422/422 [=====] - 3s 6ms/step - loss: 0.2032 - accuracy: 0.9249 - val_loss: 0.2731 - val_accuracy: 0.9025
Epoch 17/25
422/422 [=====] - 3s 8ms/step - loss: 0.1989 - accuracy: 0.9259 - val_loss: 0.2437 - val_accuracy: 0.9140
Epoch 18/25
422/422 [=====] - 2s 6ms/step - loss: 0.1932 - accuracy: 0.9280 - val_loss: 0.2496 - val_accuracy: 0.9115
Epoch 19/25
422/422 [=====] - 2s 6ms/step - loss: 0.1908 - accuracy: 0.9287 - val_loss: 0.2471 - val_accuracy: 0.9080
Epoch 20/25
422/422 [=====] - 2s 6ms/step - loss: 0.1855 - accuracy: 0.9312 - val_loss: 0.2432 - val_accuracy: 0.9100
Epoch 21/25
422/422 [=====] - 2s 6ms/step - loss: 0.1796 - accuracy: 0.9328 - val_loss: 0.2621 - val_accuracy: 0.9055
Epoch 22/25
422/422 [=====] - 3s 8ms/step - loss: 0.1772 - accuracy: 0.9352 - val_loss: 0.2577 - val_accuracy: 0.9092
Epoch 23/25
422/422 [=====] - 2s 6ms/step - loss: 0.1702 - accuracy: 0.9368 - val_loss: 0.2515 - val_accuracy: 0.9123
```

Epoch 24/25

422/422 [=====] - 3s 6ms/step - loss: 0.1665 - accuracy: 0.9373 - val_loss: 0.2510 - val_accuracy: 0.9117

Epoch 25/25

422/422 [=====] - 2s 6ms/step - loss: 0.1659 - accuracy: 0.9387 - val_loss: 0.2455 - val_accuracy: 0.9143



✓ Ewaluacja modelu

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Precyzja: ', score[1])
```

Precyzja: 0.9099000096321106

Wykresy precyzji i błędu

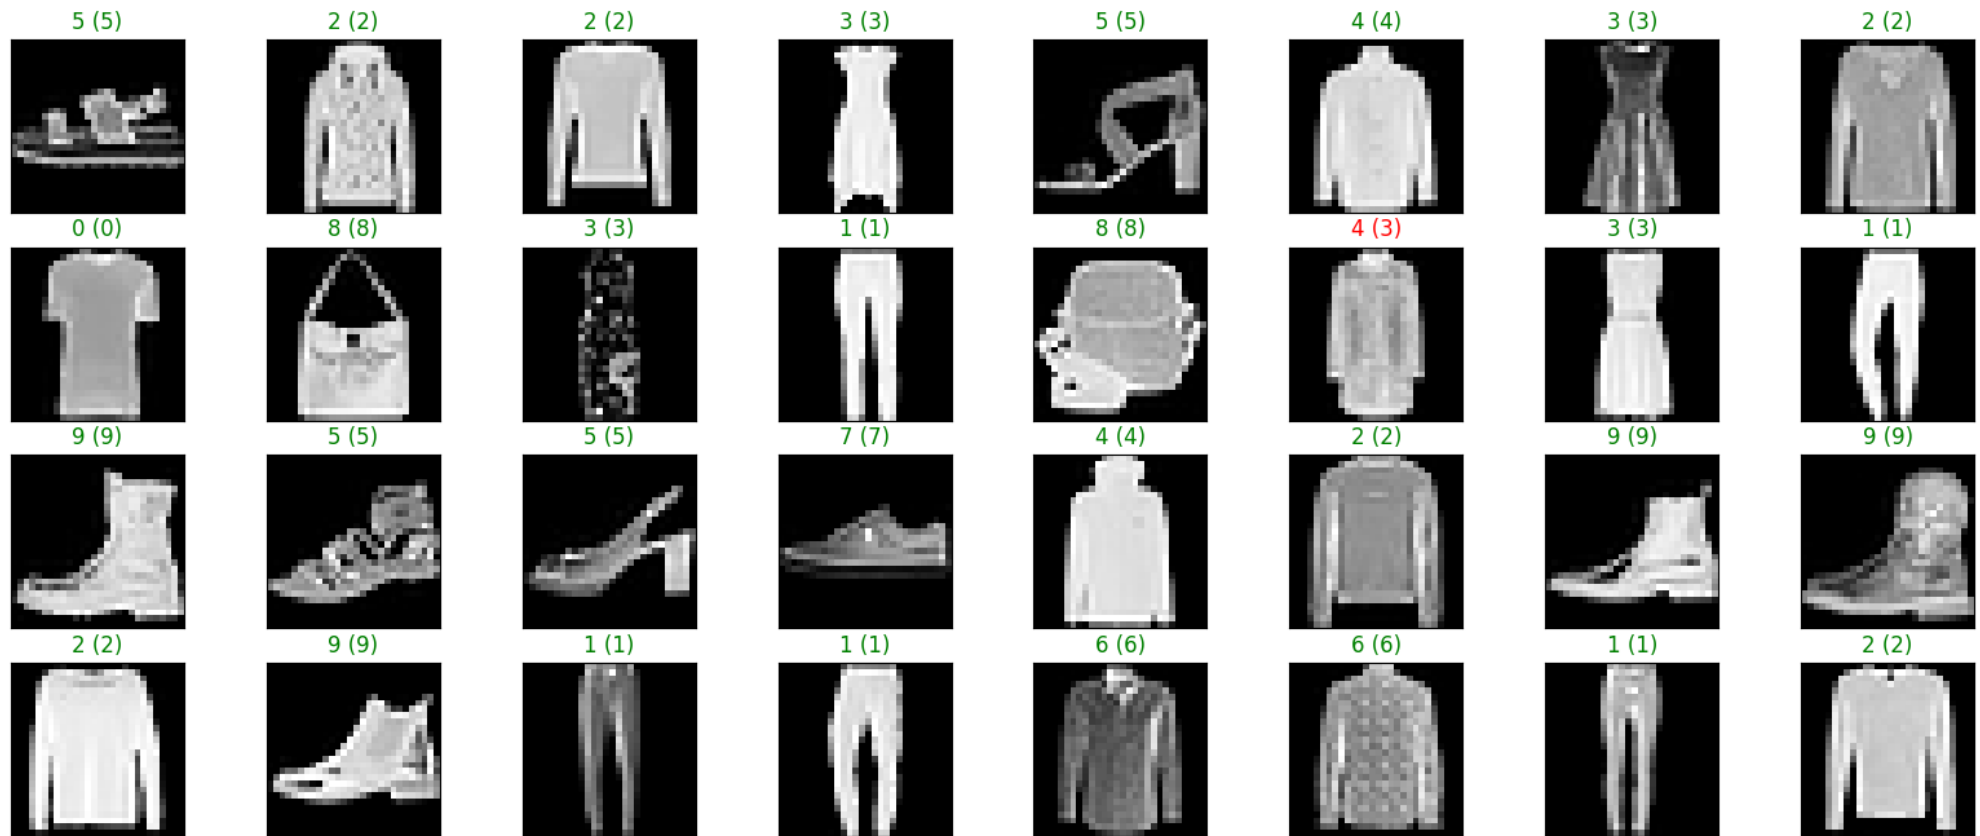
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

```
313/313 [=====] - 1s 2ms/step
```

convnet wyniki:



3 Regularyzacja

Nasz model ma obecnie dużo stopni swobody (ma DUŻO parametrów i dlatego może dopasować się do niemal każdej funkcji, jeśli tylko będziemy trenować wystarczająco długo). Oznacza to, że nasza sieć jest również podatna na przeuczenie.

W tej sekcji dodajmy warstwy dropout pomiędzy głównymi warstwami naszej sieci, aby uniknąć przeuczenia.


```
model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2,padding='same',activation='relu',input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2,padding='same',activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2,padding='same',activation='relu'))

model.add(tf.keras.layers.Flatten())

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout (Dropout)	(None, 14, 14, 64)	0
conv2d_9 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_9 (MaxPooling2D)	(None, 7, 7, 32)	0
dropout_1 (Dropout)	(None, 7, 7, 32)	0
conv2d_10 (Conv2D)	(None, 7, 7, 16)	2064
flatten_12 (Flatten)	(None, 784)	0
dropout_2 (Dropout)	(None, 784)	0

dense_33 (Dense)	(None, 128)	100480
dense_34 (Dense)	(None, 64)	8256
dense_35 (Dense)	(None, 10)	650

```
=====
Total params: 119994 (468.73 KB)
Trainable params: 119994 (468.73 KB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
model.fit(x_train,
          y_train,
          batch_size=128,
          epochs=25,
          validation_data=(x_valid, y_valid))
```

```
Epoch 1/25
422/422 [=====] - 7s 9ms/step - loss: 0.8686 - accuracy: 0.6681 - val_loss: 0.5237 - val_accuracy: 0.8052
Epoch 2/25
422/422 [=====] - 3s 8ms/step - loss: 0.5599 - accuracy: 0.7899 - val_loss: 0.4507 - val_accuracy: 0.8330
Epoch 3/25
422/422 [=====] - 3s 8ms/step - loss: 0.5045 - accuracy: 0.8107 - val_loss: 0.3926 - val_accuracy: 0.8550
Epoch 4/25
422/422 [=====] - 3s 7ms/step - loss: 0.4702 - accuracy: 0.8233 - val_loss: 0.3652 - val_accuracy: 0.8642
Epoch 5/25
422/422 [=====] - 3s 7ms/step - loss: 0.4425 - accuracy: 0.8353 - val_loss: 0.3641 - val_accuracy: 0.8647
Epoch 6/25
422/422 [=====] - 3s 8ms/step - loss: 0.4225 - accuracy: 0.8417 - val_loss: 0.3312 - val_accuracy: 0.8783
Epoch 7/25
422/422 [=====] - 3s 8ms/step - loss: 0.4103 - accuracy: 0.8459 - val_loss: 0.3324 - val_accuracy: 0.8770
Epoch 8/25
422/422 [=====] - 3s 7ms/step - loss: 0.3966 - accuracy: 0.8532 - val_loss: 0.3089 - val_accuracy: 0.8878
Epoch 9/25
422/422 [=====] - 3s 7ms/step - loss: 0.3856 - accuracy: 0.8555 - val_loss: 0.3170 - val_accuracy: 0.8802
Epoch 10/25
422/422 [=====] - 3s 8ms/step - loss: 0.3761 - accuracy: 0.8600 - val_loss: 0.2988 - val_accuracy: 0.8873
Epoch 11/25
422/422 [=====] - 3s 8ms/step - loss: 0.3657 - accuracy: 0.8637 - val_loss: 0.2918 - val_accuracy: 0.8910
Epoch 12/25
422/422 [=====] - 3s 7ms/step - loss: 0.3620 - accuracy: 0.8656 - val_loss: 0.2881 - val_accuracy: 0.8932
Epoch 13/25
422/422 [=====] - 3s 7ms/step - loss: 0.3504 - accuracy: 0.8682 - val_loss: 0.2857 - val_accuracy: 0.8942
Epoch 14/25
422/422 [=====] - 3s 8ms/step - loss: 0.3457 - accuracy: 0.8706 - val_loss: 0.2854 - val_accuracy: 0.8912
Epoch 15/25
422/422 [=====] - 3s 8ms/step - loss: 0.3398 - accuracy: 0.8734 - val_loss: 0.2701 - val_accuracy: 0.8997
Epoch 16/25
422/422 [=====] - 3s 7ms/step - loss: 0.3339 - accuracy: 0.8742 - val_loss: 0.2892 - val_accuracy: 0.8913
```

```

Epoch 17/25
422/422 [=====] - 3s 7ms/step - loss: 0.3320 - accuracy: 0.8759 - val_loss: 0.2703 - val_accuracy: 0.8988
Epoch 18/25
422/422 [=====] - 3s 8ms/step - loss: 0.3276 - accuracy: 0.8782 - val_loss: 0.2721 - val_accuracy: 0.9003
Epoch 19/25
422/422 [=====] - 3s 8ms/step - loss: 0.3185 - accuracy: 0.8808 - val_loss: 0.2638 - val_accuracy: 0.9025
Epoch 20/25
422/422 [=====] - 3s 7ms/step - loss: 0.3200 - accuracy: 0.8796 - val_loss: 0.2626 - val_accuracy: 0.9017
Epoch 21/25
422/422 [=====] - 3s 7ms/step - loss: 0.3187 - accuracy: 0.8794 - val_loss: 0.2589 - val_accuracy: 0.9028
Epoch 22/25
422/422 [=====] - 3s 8ms/step - loss: 0.3165 - accuracy: 0.8826 - val_loss: 0.2549 - val_accuracy: 0.9025
Epoch 23/25
422/422 [=====] - 3s 8ms/step - loss: 0.3134 - accuracy: 0.8820 - val_loss: 0.2575 - val_accuracy: 0.9025
Epoch 24/25
422/422 [=====] - 3s 7ms/step - loss: 0.3082 - accuracy: 0.8850 - val_loss: 0.2508 - val_accuracy: 0.9058
Epoch 25/25
422/422 [=====] - 3s 7ms/step - loss: 0.3082 - accuracy: 0.8843 - val_loss: 0.2527 - val_accuracy: 0.9033
<keras.src.callbacks.History at 0x7f529797fd90>

```

✓ Evaluate model:

```

test_score = model.evaluate(x_test, y_test, verbose=0)
train_score = model.evaluate(x_train, y_train, verbose=0)



print('Train accuracy: ',train_score[1], ' Test accuracy: ',test_score[1])

Train accuracy:  0.9166481494903564  Test accuracy:  0.9003000259399414

```

Wykresy precyzji i błędu

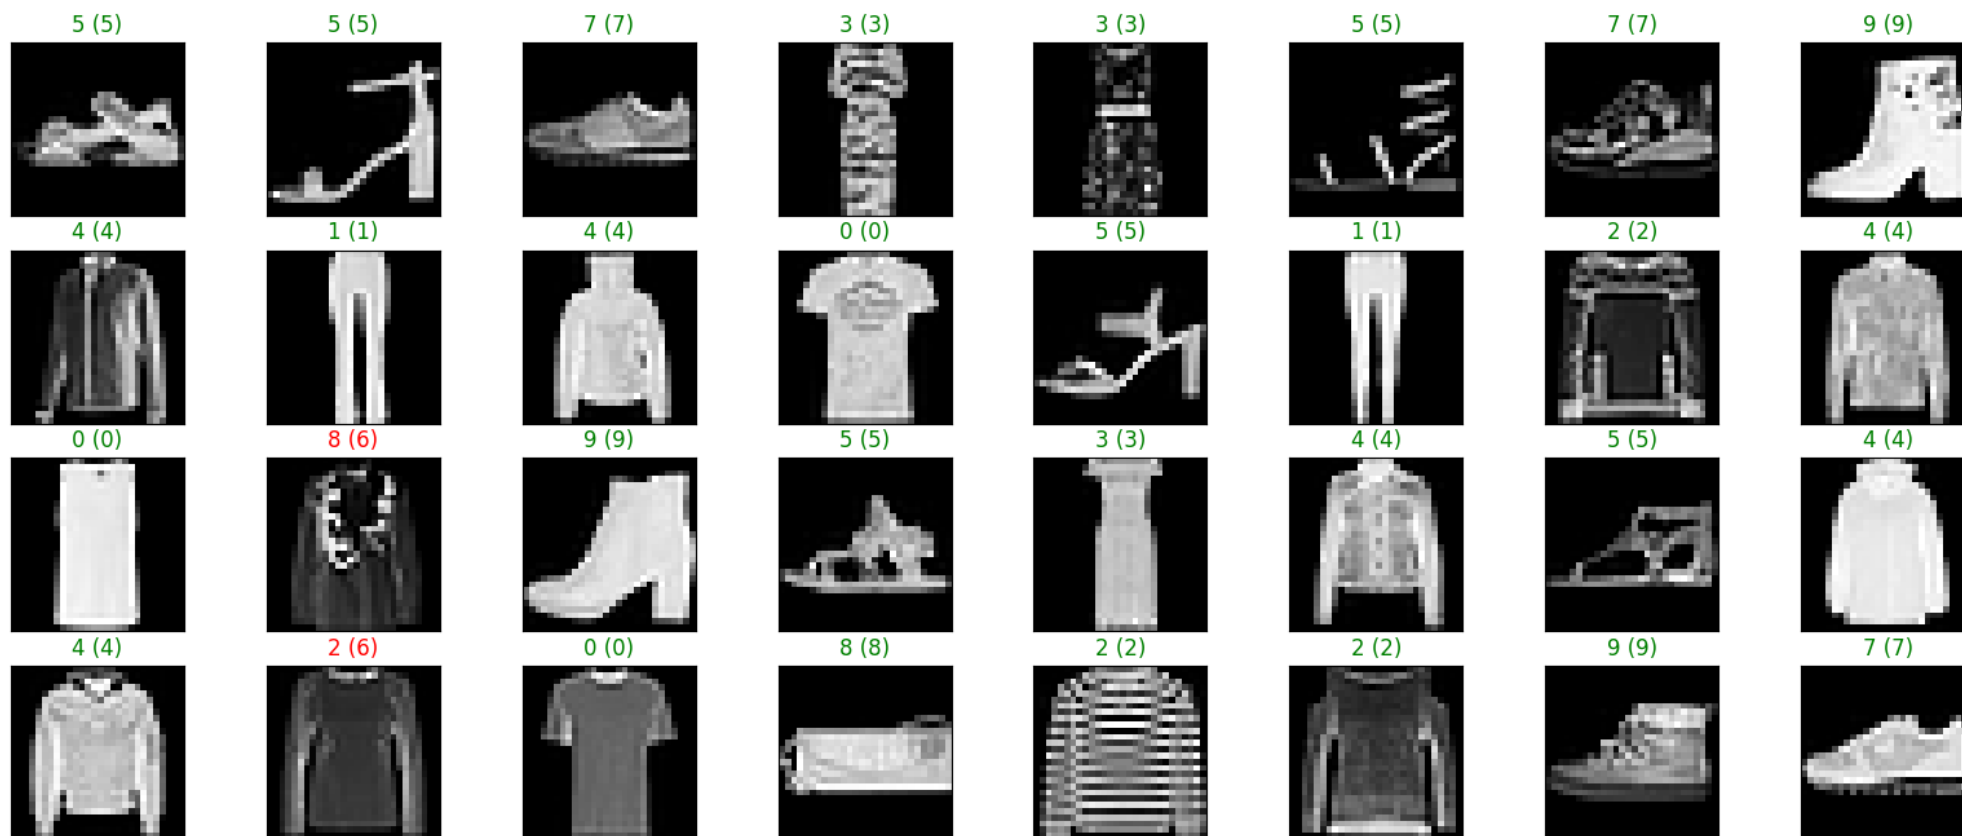
```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
visualize_model_predictions(model, x_test, y_test, "convnet")
```

```
313/313 [=====] - 1s 2ms/step
```

convnet wyniki:



3 Regularyzacja v2

Nasz model ma obecnie dużo stopni swobody (ma DUŻO parametrów i dlatego może dopasować się do niemal każdej funkcji, jeśli tylko będziemy trenować wystarczająco długo). Oznacza to, że nasza sieć jest również podatna na przeuczenie.

W tej sekcji dodajmy warstwy dropout pomiędzy warstwami naszej sieci, aby uniknąć przeuczenia.

```

model = tf.keras.Sequential()

model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=2,padding='same',activation='relu',input_shape=(28,28,1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=2,padding='same',activation='relu'))
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=2,padding='same',activation='relu'))

model.add(tf.keras.layers.Flatten())

model.add(tf.keras.layers.Dropout(0.5))

model.add(tf.keras.layers.Dense(128, activation='relu'))
model.add(tf.keras.layers.Dropout(0.1))
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10,activation="softmax"))

model.summary()
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_10 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_3 (Dropout)	(None, 14, 14, 64)	0
conv2d_12 (Conv2D)	(None, 14, 14, 32)	8224